

Basic Model

Contents

BasicModel	6
Basic Model of Cognition	6
Documentation	6
Overview	6
Files	6
Quick Start	7
XML Configuration	7
Output	8
Installation and Usage	8
Prerequisites	8
Setup	8
Makefile Targets	8
Training	8
Models	8
Testing and Documentation	8
Utilities	9
Makefile Variables	9
train.py Options	9
General	9
Remote Execution (SSH)	9
Environment Variables	10
Private Remote Training	10
Architecture	11
The three pieces: Mereology, Attention, Thought (2026-06-11)	11
Symbolic weights, reconstruction, parse-time, attention (2026-06-30)	13
A. Symbolic weights (the two-phase forward; reworked 2026-07-02)	13
B. Reconstruction (parts $\rightarrow$.what , wholes $\rightarrow$.where)	14
Tiling, subspace sizes, and consciousness	14
C. Parse time (two attentions)	14
D. Attention indexing (.where / .when / codebooks)	15
Cognitive grounding: dense-perceptual vs sparse-symbolic (2026-07-02)	15
Overview	16
Spaces	17
Reconstruction Symbols	18
Single Optimizer with Overlapping Weight Spaces	19
Training Loop	19
The three cognitive operations (2026-06-14)	19
Modes of operation	20
Pipeline as a unit, two-space-role reset	21
Two-File Architecture	22

Language System	22
Short-Term Memory on ConceptualSpace	23
Per-word operational flow (SERIAL mode)	23
Sigma and Pi Layers	24
Dimensionality Constraints	24
Invertible Linear Layer (LDU)	24
Ergodic Exploration	25
Sentence-level AR (InterSentenceLayer)	25
Sentence representation	25
ARMA(p, q) predictor	25
Wiring into the training loop	25
Inference (chat-loop seeding)	26
Configuration	26
Basic Model	27
Inputs and Outputs	28
Perceptual Space	28
Conceptual Space	28
Order, extension, and intension (2026-07-02)	28
Symbolic Space	29
Reasoning	29
Symmetric Perception: the Single Layer Linear Case	30
Symmetric Perception: the Single Layer Nonlinear Case	30
Distance Metrics	31
Exploration/Exploitation as a function of temperature	31
Transfer Functions	31
Spaces	32
Overview	33
Base Class: Space	34
Reset cascade — hard vs. soft	34
Sigma / Pi ownership (Pi/Sigma swap, rev. 2026-06-09)	35
Lift / Lower in the new ownership	35
Butterfly mode on GrammarLayer (Stage 5)	36
Normalization and Ranges	36
Codebook Similarity Metric	37
Euclidean (Perceptual — Part / Whole)	37
Dot product (Conceptual)	37
Configuring the metric	38
Ramsification table (per-code fold record)	38
Codebook Uniqueness Contract	38
Lexicon (Projective Unit Ball)	39
Distance and lookup	39
SBOW training: pole (attractor) and antipole (repulsion target)	40
InputSpace	40
Document streaming and valid_mask	41
AR cursor unfold retirement (2026-05-13)	41
Within-sentence AR retirement (2026-05-14)	41
PartSpace	42
ConceptualSpace	43
Word auto-bind deferred to Reset (fullgraph forward)	44
ShortTermMemory	45
Lift / Lower as binary GrammarLayer subclasses	45
WholeSpace	46

Monotonicity of the lift / lower chain	47
SyntacticSpace — retired	48
OutputSpace	48
Short-Term Memory	48
Overview	48
1. STM data model	49
2. Serial sequencing	49
3. Parallel sequencing	50
4. Attentional filtering	51
5. Routing parser: SS-analysis vs CS-execution	51
6. IntraSentenceLayer	52
7. Per-word router firing	53
8. Masked-word reconstruction via priming	53
9. Relative vs absolute end-states	54
10. LTM as the chain of STM end-states	55
11. Inter-sentence prediction	55
12. nanochat comparison framing	56
See also	56
Language	57
Current Parser Surface	58
Retired XML Knobs	58
Grammar	58
Knowledge Artifacts	59
STM Shift/Reduce	59
Syntax And SVO	60
Role-Collapsed Grammar and the Operator Codebook	60
Soft Operator Superposition	60
Participation Categories as the Chooser’s Syntactic-Category Context	61
Shared Weighted-Deduction Framework	62
Soft-superposition route (the <learning> two-pass)	62
Mereology	63
Parthood as Clipped Cosine Projection	64
The Full Suite	64
Region Partition	65
Mereological Fusion	65
ImpenetrableLayer: Five-Relations Regularizer	65
Penalty	66
Trust	66
Diagnostics	66
Configuration	66
Order-raising (building the meronymic lattice)	66
The corpus callosum links the towers (part isa whole)	67
Explicit concepts implicit subsymbolic representations (dual-coded, 2026-06-21)	68
Why This Design	69
Meronymic Analyzer (PartSpace)	69
isPart as a Grammar Layer	70
Logic	70
Overview	70
1. Subsymbolic Layer	70
Operators	70
2. Symbolization	71

3. Symbolic Layer (Scalars in $[-1,1]$)	71
Parthood as Projection	71
4. Key Insight	72
5. Rationality	72
Truth Statements	72
TruthLayer (SymbolSpace.truth_layer)	73
Truth Field	73
Consonance and Dissonance	73
Propositional Structure	73
Truth Accumulation Pipeline	73
6. Consistency and Verification	74
Internal Consistency (within a TruthSet)	74
Verification (incoming statement against stored truth)	74
Logical Entailment (symbolic closure)	74
7. Radial Operators for Hypersphere Ternary Logic	75
Luminosity	76
Fusion	77
Supporting measures	77
Semantic Summary	77
8. Cross-references	77
Reasoning System	78
Partitioned Symbol Space	78
Reasoning Methods	78
isConsistent() -> dict	78
ground(activation, threshold=0.6) -> dict	78
isTrue(activation) -> float	78
extrapolate(seed_indices, max_new, attenuation) -> dict	78
TruthLoss	79
Bidirectional Reasoning Loop	79
Grammar Learning	79
Parser And Conceptual Order	80
Configuration	80
Contemplative Awareness Methods	80
Testing	81
Training	81
Overview	81
The staged objective curriculum	81
Phase 1: Embedding Pretraining (make train / embed.py train)	82
Pipeline	82
SBOW (Sentence Bag of Words)	82
CBOW (Continuous Bag of Words)	83
Two-Pass Architecture	83
Configuration (Makefile variables)	83
Memory Considerations	84
Phase 2: Network Training (BasicModel.py)	84
Within-sentence training objective (IR-only)	84
<trainEmbedding>: Embedding Update Mode	84
Gradient Flow Through Codebook	84
JOINT: Single Combined Loss	85
Why MSE over Embeddings (not Cross-Entropy)	85
Training Loop	85
SBOW vs CBOW	86

Three-File Architecture	86
Embedding Space Geometry	86
Profiling	87
Ergodic Exploration via Adaptive Bias-Variance Control	87
1. The Ergodic Weights	87
2. Why Standard Adam Fails on α	88
3. The Gradient Energy Sensor	88
3.1 Modified Second-Moment Estimator	88
3.2 Bias Correction	88
3.3 What $\hat{v}_t$ Measures	88
3.4 The Alpha Update Rule	89
3.5 Layer-Local Certainty (Per-Neuron Sigma)	89
4. Comparison: Adam vs. Gradient Energy Sensor	89
5. Derivation: Why Drop m_t	90
6. Dropout via Bias	90
7. The <code>global_temp</code> Sensitivity Knob	90
8. Initialization and Bootstrap	90
9. Layer Architecture	91
ErgodicLayer Base Class	91
InvertibleLinearLayer — factor-level noise injection	91
10. Training Loop Integration	91
11. Summary of the Full Algorithm	92
Machine Minds	92
Abstract	92
Background	92
Part I: What Machine Minds Are — Weight Ergodicity	92
Weights as Ergodic Samples	92
Temperature and Annealing	93
Advantages of Ergodic Weights	93
Part II: What Machine Minds Feel — Network Invertibility	93
Bidirectional Perception	93
Symbols, Worldlines, and Karmic Inertia	93
Implementation	94
Part III: What Machine Minds Know — Output Certainty	94
Certainty vs. Probability	94
Certainty-Weighted Loss	94
Knowing What You Do Not Know	94
Methods	94
Results	95
Discussion	96
Conclusion	96
References	96
XML Configuration Parameters	96
<architecture>	97
<architecture><data>	98
<architecture><training>	98
<architecture><server>	99
BasicModel / MentalModel Quick Reference	99
Space sections	100
<InputSpace>	100
<PartSpace>	100

<ConceptualSpace>	101
<WholeSpace>	101
<OutputSpace>	102
<SymbolSpace>	102
InvertibleLinearLayer (programmatic)	103
Example: XOR Configuration	103
Example: Embedding / Chatbot Configuration	104

BasicModel

Basic Model of Cognition

The basic model of cognition relies on conceptual hyperplanes and perceptual prototypes to synthesize and analyze the input space. It uses a high-dimensional embedding to characterize mental space and integrates symbolic computation. Its three major operations are intersection (analysis: the top-down division of the presented unity into symbolic generalities), union (synthesis: the bottom-up aggregation of atoms into perceptual particulars), and equality (where symbols are elements that map across perceptual and conceptual domains). See Philosophy for the analytic/synthetic orientation.

Documentation

Document	Description
Architecture	Pipeline design, layer types, invertible LDU factorisation
Spaces	All six spaces: Input $\rightarrow$ Perceptual $\rightarrow$ Conceptual $\rightarrow$ Symbolic $\rightarrow$ Syntactic $\rightarrow$ Output
Ergodic	Gradient energy sensor, adaptive exploration, factor-level noise injection
Training	Two-phase training, SBOW embeddings, masked prediction modes
Params	Complete XML configuration reference
BasicModel	Cognitive science foundations
Language	Grammar, word encoding, symbolic and syntactic spaces
Logic	Subsymbolic and symbolic logic operations
Reasoning	Truth methods, bidirectional reasoning, contemplative awareness stubs
Philosophy	Kant's analysis/synthesis, Ramsey's theoretical roles, Buddhist epistemology (pramana, tetralemma)
MachineMinds	What machine minds are, feel, and know
Installation	Setup, Makefile targets, environment variables

Overview

BasicModel is a parameterized neural architecture that answers the question “what is the truest thing we can say of this experience?”.

Model configurations are specified in XML. See doc/Architecture.md for the full mathematical treatment.

Files

File	Description
bin/Models.py	Main entry point: model factory, training loop, <code>--compare</code> mode, HTML report
bin/Layers.py	Layer library: SigmaLayer, PiLayer, ErgodicLayer, LinearLayer, TruthLayer
bin/Spaces.py	Space classes: InputSpace, PartSpace, ConceptualSpace, WholeSpace, OutputSpace, Grammar
bin/embed.py	Word vector training: CBOW/SBOW with negative sampling, WordVectors (gensim-compatible) .k
data/	XML model configurations
doc/Architecture.md	Algorithm details: Sigma/Pi layers, ergodic exploration, gradient energy sensor
doc/Params.md	Full XML parameter reference

File	Description
doc/Training.md	Embedding pretraining, CBOW/SBOW, masked prediction, <code><trainEmbedding></code> modes
doc/Installation.md	Setup, Makefile targets, train.py options, remote training
test/	Unit tests

Quick Start

Set up virtual environment

`make install`

Run a single model

`make simple` *# data/simple.xml*

`make ergodic` *# data/ergodic.xml*

Compare two models side-by-side

`make compare` *# defaults: data/simple.xml vs data/ergodic-only.xml*

`make compare XML1=data/simple.xml XML2=data/ergodic.xml`

Run tests

`make test`

Generate PDF documentation

`make doc`

XML Configuration

Models are configured via XML files in `data/`. Training and data parameters live in nested sub-elements:

```
<model>
  <architecture>
    <ergodic>false</ergodic>
    <certainty>false</certainty>

    <data>
      <dataset>xor</dataset>      <!-- xor / phrases / mnist / tomatoes / text / inline -->
      <dataType>numeric</dataType> <!-- embedding / numeric -->
    </data>

    <training>
      <numTrials>1</numTrials>
      <numEpochs>20</numEpochs>
      <batchSize>10</batchSize>
      <learningRate>0.001</learningRate>
      <weightsPath>output/BasicModel.ckpt</weightsPath>
      <autoload>true</autoload>
      <autosave>false</autosave>
    </training>
  </architecture>

  <InputSpace> ... </InputSpace>
  <PartSpace> ... </PartSpace>
  <ConceptualSpace> ... </ConceptualSpace>
  <WholeSpace> ... </WholeSpace>
```

```
<OutputSpace> ... </OutputSpace>
</model>
```

See doc/Params.md for the full parameter reference.

Output

Each run produces an HTML report (timestamped in `output/`) containing:

- **Error per Epoch** – training and test loss curves
- **Accuracy per Digit** – per-class accuracy breakdown

In compare mode, additional overlay plots show combined loss and accuracy across models with color-coded legends.

Installation and Usage

Prerequisites

- **Python 3.12** with a virtual environment at `.venv/`
- **PyTorch** with MPS (Apple Silicon) or CUDA
- **pandoc** (optional, for PDF generation via `make doc`)

Setup

```
make install    # creates .venv with dependencies
```

All Makefile targets invoke Python through `.venv/bin/python`, so no manual activation is required.

Makefile Targets

Training

Target	Description
<code>make train</code>	Full training (Phase 1 embeddings + Phase 2 model), logged to <code>output/logs/</code>
<code>make train_micro</code>	Small run: max 1,000,000 docs across 10 random shards, logged

Models

Target	Description
<code>make run</code>	Run <code>Models.py</code> with the config in <code>XML1</code>
<code>make xor</code>	<code>data/MM_xor.xml</code>
<code>make simple</code>	<code>data/simple.xml</code>
<code>make ergodic</code>	<code>data/ergodic.xml</code>
<code>make tomatoes</code>	<code>data/tomatoes.xml</code>
<code>make mnist</code>	<code>data/mnist.xml</code>
<code>make compare</code>	Side-by-side compare using <code>XML1</code> and <code>XML2</code>
<code>make SigmaPi / SymPercept / SPNN</code>	Run respective scripts

Testing and Documentation

Target	Description
<code>make test</code>	Unit tests (forces <code>BASICMODEL_DEVICE=cpu</code>)
<code>make bench</code>	Training benchmarks (baseline)
<code>make doc</code>	Generate <code>BasicModel.pdf</code> via <code>pandoc</code>

PDF chapters in order: `README.md`, `doc/Installation.md`, `doc/Architecture.md`, `doc/BasicModel.md`, `doc/Spaces.md`, `doc/STM.md`, `doc/Language.md`, `doc/Mereology.md`, `doc/Logic.md`, `doc/Reasoning.md`, `doc/Training.md`, `doc/Ergodic.md`, `doc/MachineMinds.md`, `doc/Params.md`.

Utilities

Target	Description
<code>make clean</code>	Remove generated files (<code>BasicModel.pdf</code>)
<code>make all</code>	Default; alias for <code>make xor</code>

Makefile Variables

Override on the command line, e.g. `make run XML1=data/ergodic.xml`.

Variable	Default	Description
<code>MODEL</code>	<code>data/MM_20M_legacy.xml</code>	XML config for training
<code>XML1</code>	<code>data/simple.xml</code>	Primary config for <code>make run</code> / <code>compare</code>
<code>XML2</code>	<code>data/ergodic-only.xml</code>	Secondary config for <code>make compare</code>

train.py Options

`train.py` orchestrates Phase 1 (embeddings) followed by Phase 2 (model training).

General

Flag	Default	Description
<code>--model, -m</code>	<code>data/BasicModel.xml</code>	XML config
<code>--max-docs</code>	<i>(from XML)</i>	Override <code>maxDocs</code>
<code>--num-shards</code>	<i>(from XML)</i>	Override <code>numShards</code>
<code>--compile-target</code>	<code>gpu</code>	<code>gpu</code> runs normal training with accelerator <code>torch.compile</code> defaults; <code>mlx</code>
<code>--compile-mode</code>	<i>(env/runtime default)</i>	Forwarded as <code>MODEL_COMPILE_MODE</code>
<code>--mlx-output</code>	<code>output/mlx/<model>.pte</code>	Destination for <code>--compile-target mlx</code>
<code>--random-shards</code>	<code>off</code>	Pick random shard indices
<code>--force-embeddings</code>	<code>off</code>	Retrain embeddings even if file exists
<code>--log [FILE]</code>	<code>off</code>	Capture stdout/stderr to <code>.log</code> in <code>output/logs/</code> . Omit filename for <code>train</code>
<code>--profile</code>	<code>off</code>	Run Phase 2 under <code>cProfile</code> ; writes <code>.prof</code> to <code>output/profiles/</code>

Remote Execution (SSH)

Flag	Default	Description
<code>--host</code>	<i>(none)</i>	Remote host; enables remote run
<code>--user</code>	<code>arogers</code>	SSH user
<code>--key-file</code>	<i>(none)</i>	Optional SSH key file; uses SSH config/agent when omitted
<code>--remote-dir</code>	<code>~/WikiOracle/basicmodel</code>	Working directory on remote

Environment Variables


Variable	Description
<code>BASIC_MAX_DOCS</code>	Override <code>maxDocs</code>
<code>BASIC_NUM_SHARDS</code>	Override <code>numShards</code>
<code>BASICMODEL_DEVICE</code>	Force device, e.g. <code>cpu</code> . Used by <code>make test</code>
<code>BASICMODEL_PYTHON</code>	Python executable for <code>train.py</code> subprocesses; overrides the in-project <code>.venv</code> lookup
<code>MODEL_COMPILE</code>	torch.compile backend selector (<code>auto</code> , <code>none</code> , <code>inductor</code> , <code>eager</code> , <code>aot_eager</code>)
<code>MODEL_COMPILE_MODE</code>	torch.compile mode (<code>default</code> , <code>reduce-overhead</code> , <code>max-autotune</code> , <code>max-autotune-native</code>)
<code>PYTORCH_MPS_HIGH_WATERMARK_RATIO</code>	MPS memory limit; set to 0.0 by <code>train.py</code>
<code>PYTHONUNBUFFERED</code>	Set to 1 by <code>train.py</code> for real-time log streaming

Private Remote Training

Machine-specific SSH targets, LAN hostnames, and local training shortcuts should live in the ignored top-level `Makefile.local`. The public `train.py --host` mode remains available for generic SSH execution when a caller supplies the host, user, optional key, and remote directory explicitly.

Architecture

The three pieces: Mereology, Attention, Thought (2026-06-11)



Terminology (2026-06-21 convention). One noun per-space: a **percept** is a perceptual thing (PartSpace/WholeSpace, dimensionally-embedded, extensional; *part* and *whole* are its two subtypes); a **concept** is a ConceptualSpace relation tying one part-percept to one whole-percept (the Concept codebook); a **symbol** is a SymbolSpace 0-D reference to a concept. The CS “symbol table” is therefore the **Concept codebook** below.

The architecture decomposes into three pieces; the first two run in parallel, the third is serial:

1. **The mereological towers (and the Concept codebook).** In an LLM the mereology is completely subsymbolic — implicit in the weights, never surfaced. In BasicModel it is percept-based and symbolic: two towers — the σ tower ascending bottom-up (part-percept extents, the PS codebook) and the π tower descending top-down (whole-percept intents, the SS codebook) — linked by the word/object Concept codebook (`bin/References.py`), whose rows are the concepts (part whole references).
2. **Attention.** In an LLM, attention is QKV per subsymbolic layer. In BasicModel, attention is quadratic over all three structures — both towers and the Concept codebook — realized as priming (boost weights over codebook rows; the single intent of GrammarOpsPass §5).
3. **Thought.** In an LLM, thought is the computation of priors for autoregressive word prediction. In BasicModel, thought is a **subsequent isolation of attention over that space, enabled by references** — the serial pass: referential lookup, shift/reduce composition, story selection. Because thought is the only process that invokes the referential taxonomy *qua* references, it is the only process licensed to **shape the references**.

This yields the codebook update law (GrammarOpsPass §6d, implemented): **percepts are shaped by the parallel pass; references are shaped by the serial pass**. STE is untouched in both modes; the partition governs row updates — parallel mode may not shape references, serial mode may not shape non-references (`Spaces.reference_update_mask`; the `update_mask_fn` chokepoint on `VectorQuantize`). Rational by construction — though possibly an overly optimistic arrangement with respect to human thinking.

See SymbolFirewall.md for the governing principle this update law (and the codebook/meronomy ownership model generally) instances: all computation is composed over typed, symbol-attached units — read/write masks, no anonymous global residual stream.

Addressable attention — the typed .where Global attention (GlobalAttention, bin/Spaces.py; gated <globalAttention>) ranges over a **typed addressable space**: one distribution competes across every store at once and emits a typed **.where = (space-id, bracket)** plus a soft-read $\Sigma \cdot \text{key}$. Six stores (the SPACE_* ids):

id	store
INPUT	the staged input window (per-span percept content)
STM	the live short-term-memory rows
LTM	the consolidated truth store (rows + trust value)
PART	the PartSpace codebook (part-percepts)
WHOLE	the WholeSpace codebook (whole-percepts + meronomy/taxonomy)
SYMBOL	the SymbolSpace codebook (symbols, 1:1 with concepts)

PART/WHOLE appear whenever their tower has a codebook; SYMBOL only under <symbolTower>. Pointing .where at a codebook/LTM store is recall; at the input window it is reading — one mechanism, the type tag distinguishes them. Under <globalAttentionConsume> the soft-read is fed back into the head as a zero-init gated residual, so the output loss trains the retrieval.

The symbol codebook is a **reference**, not a learned copy: it tracks the concept codes so the two cannot diverge or dissociate. Its training objective is kept separate from reconstruction (the source concepts are shaped by their own pass): the symbol's IDENTITY (the codebook row) stays EMA-only/detached, while the symbol's VALUE (the signed 0-D activation from the sparse forward) carries gradient through the bind leg (2026-07-02 fix; the leg was previously detached wholesale).

These six stores are the substrate for the four foundations of mindfulness; the mapping (and the trust-sign-as-vedana / luminosity-as-joy reading of LTM) is in Philosophy.md.

Status (2026-05-29 update): further architectural pivots landed on top of the 2026-05-27 substrate refactor:

- **PerceptStore → RadixLayer.** The radix-trie input encoder is now a first-class Layer subclass in bin/Layers.py; PartSpace.reverse invokes RadixLayer.reverse for the structural decode.
- **MetaLayer → SymbolizeLayer.** The binary GrammarLayer that promotes a freshly-seen percept to a symbolic prototype is now SymbolizeLayer; no semantic change.
- **Auto-META moves PS → CS.** The cross-codebook bind (META entry: PS chunk-id ↔ SS prototype-id) fires from ConceptualSpace._maybe_autobind_meta at stage 0; PS no longer holds a back-ref to SS.
- **Clean-stack STM.** ConceptualSpace.forward bypasses sigma_in / sigma_cs on forward — folded = primary at stage 0, folded = sym at k > 0. The Stage-10 additive composition is retired; per-stage space-role attribution is trivially invertible.
- **basis= kwarg for grammar reverses.** UnionLayer.reverse(parent, basis=None) / IntersectionLayer.reverse(parent, basis=None) accept a Codebook / Basis object (typically WholeSpace.subspace.what) instead of a raw W tensor; bin/Language.py::unreduce() dispatches accordingly.
- **LSE soft-max kernels.** Ops._disjunction_kernel / Ops._conjunction_kernel default to LogSumExp smooth variants when monotonic=False; the hard branch is retained for monotonic-mode and exact idempotency tests.
- **LBG-style SS codebook splitting.** Gray (1990) EMA + per-row variance tracking; rows whose running variance exceeds a threshold split along the top-variance eigendirection.

See doc/old/2026-05-29-clean-stack-stm-basis-arg-radixlayer.md for the consolidated rationale.

Status (2026-05-27): the **substrate refactor** has landed end-to-end (doc/old/2026-05-26-two-loop-pi-sigma-substrat PS is a single-arg input processor (pi + sigma). CS is a STM container + grammatical CPU (no atomic forward fold; sigma_percept retired). SS owns the unified word lexicon codebook with paired (orth, semantic) rows. The CKY Chart and STM shift-reduce parsers retire entirely; **LanguageLayer** (signal router) is the canonical parser. **LiftLayer** / **LowerLayer** are binary **GrammarLayer** subclasses with internal Sigma / Pi (no longer borrowing substrate folds). **GrammarLayer** gains an optional butterfly cascade mode for cross-position mixing — closes the XOR convergence target. Two operating modes selected by <serial>: **SERIAL/GRAMMATICAL** (per-word PS with grammar dispatch over STM) and **PARALLEL** (T = <subsymbolicOrder> iterations of PS over CS). The <parserBackend>, <routerKind>, <chartTau>, <chartTopK>, <chartNoiseEps> XML knobs are retired; <symbolicOrder> is now the symbolic / relational loop budget.

Symbolic weights, reconstruction, parse-time, attention (2026-06-30)

A design pass clarifying four coupled pieces (PS = PartSpace, WS = WholeSpace, CS = ConceptualSpace, SS = SymbolSpace).

A. Symbolic weights (the two-phase forward; reworked 2026-07-02)

Implemented 2026-07-02 (doc/plans/2026-07-02-two-phase-loops-sparse-relation.md, superseding the forward-transform parts of 2026-07-02-sparse-layer-conceptual-embedding.md) as a dedicated **SparseLayer** (bin/Layers.py), NOT a SigmaLayer option: SigmaLayer’s atanh-entry contract expects logit-domain codes in $[-1, 1]$, while these maps consume *presences/activations* – a different input domain deserves its own class. The substrate contract is the transpose autoencoder pair: forward = $\tanh(Wx)$, reverse = $\tanh(W^T y)$; export-safe scatter-add kernel by default, torch.sparse.mm opt-in. Edges append host-side at mint (add_edge, idempotent, tail-preserving value growth) and are removed by pruning rounds (remove_edges; ConceptualSpace.prune_concept_links keeps the closest links using within-tower relations only).

The forward is TWO PHASES with one terminal cutover. Phase A – the purely continuous PS/WS $\leftrightarrow$ CS pump: subsymbolicOrder iterations of the 2-stream bind (no symbol leg in the loop, no quantization inside the pump). The C $\rightarrow$ P / C $\rightarrow$ S handoffs carry the per-tower WINDOWS of the MIXED carrier (combine.views, the demux feedback): the part-stream returns to PS for further σ synthesis, the whole-stream to WS for further π analysis – the mix goes UP (the next stage’s contribution), the un-mix goes DOWN. Phase B – ONE late cutover at the bandwidth seam (cs_symbolic_phase): the settled field is SNAPPED to the ORDER-0 block of the conceptual codebook (cs_snap_order0: differentiable normalized-sum presence – slot-mean projection onto the unit atom direction in hypercube-diagonal $\sqrt{D}$ units, magnitude-preserving, NOT cosine – with an EMA identity trace of the winning rows while training), then the ramsified symbolic composition runs once and its outputs feed the SS leg, the head-side losses (conceptual SBOW on the settled slab), and the concept table – they are NEVER substituted back into the subsymbolic carrier (the <sparseReplace> knob is retired; non-replacement is structural). Quantization sits exactly at the seam because 0-D symbols lack the bandwidth to carry subsymbolic content.

The SparseLayer is SYMBOLIC-ONLY and owns BOTH readings of concept structure. The percept families ($[PS|WS]$ presence columns) retired with the rework: a_0 comes from the snap, and order $k \geq 1$ composes lower-order symbol activations only through role-split columns [whole | part | bias], $a_k = \tanh(S_k[a_0..a_{k-1} | a_0..a_{k-1} | 1])$, sized by the dyadic capacities (order_capacities). Alongside this weighted COO reading, each per-order layer holds the DISCRETE relation table: ordered role-tagged constituent records (embed_pair stores the sec-4c ordered pair [whole, part], whole first: whole $\Rightarrow$ part / if $\rightarrow$ then; discretize_row is exact on the binary-ordered subset). A thin shared ConceptAllocator owns global concept ids, order derivation, the raise/retire/singleton sets, and the idempotency caches; ConceptualSpace keeps orchestration only. The signed bounded activation a IS the 0-D symbol: the SS bind leg is $a \times$ the row-aligned identity row (gradient flows through the activation; codebook rows stay EMA-only). Per-pass σ/π stacks (<subsymbolicStack>, with <subsymbolicNoop> identity slots) give the pump DISTINCT layers

per pass – depth IS mereological order – and the per-percept snap-residual (`snap_settle_qe`) is read as a report-only SETTLE SIGNAL for later adaptive work.

The lattice poles are VECTORS of the presence domain: **NOTHING** = $[0, 0, \dots]$ (bottom; a part contributing $W \cdot 0 = 0$ – no edge) and **EVERYTHING** = $[1, 1, \dots]$ (top; a whole realized as the 1-wide bias role block at order ≥ 1). Order-0 concepts (word A-symbols, the span knit, fresh pole-pair objects) RESERVE their order-0 codebook row – the snap reads them; their part/whole decomposition lives in the PS/WS codebooks plus the ordered reference store (store by reference, never duplicate codes). The word $\equiv$ object META is the sec-4c ORDERED PAIR [whole = word-symbol, part = object-symbol] – roles are positional slots of an ordered pair, not containment claims; the typed read-out (`meta_word_object`) recovers (word, object) by INTERSECTING the pair with the word-symbol class rather than trusting slot order. The JOINT/sentence concept (`create_joint_concept`) is the ordered Gallistel CHAIN over the row’s word-symbols – each link the pair [whole = current, part = rest], bias-bounded – one head per sentence TYPE, so word order distinguishes sentence types. A 1:1 tie between SYM refs is the SINGLETON principle (Alec 2026-07-02): the unit-set $\{x\}$ – a whole containing exactly one symbolic part (`singleton_concept`, min-support exempt) – is the constructive primitive behind if $\rightarrow$ then ($\{x\} \Rightarrow x$) and the recursion vine, and it is stored structure that `resolve_identities` never collapses (only ties between concrete raw codes resolve away). Sequencing depth under strict ramification truncates the chain’s weighted reading (same-order link references are dropped at the order cap) – the successor design (doc/plans/2026-07-02-iterated-symbolic-loop.md, DRAFT) replaces stratification with iteration over one untyped square layer.

B. Reconstruction (parts $\rightarrow$.what, wholes $\rightarrow$.where)

InputSpace maps two views of the *same* data, segmented differently: a **universal view** to WS (which it analyses) and an **atomic view** to PS (which it synthesises). WS yields **low-fidelity** information covering the **whole** space; PS yields **high-fidelity** information over a **smaller** area (`_paint_reconstruction`: the universal view paints the background, the atomic view is averaged in where it has support). For verbal reconstruction of text the **parts (PS) should reconstruct the .what** and the **wholes (WS) should reconstruct the .where** — approximately the inverse of what happens at parse time. The separate `what_scale` / `where_scale` / `when_scale` reconstruction channels already exist to carry this.

For text, it would be foolish to insist on an *absolute* .where from WS: the parts already know each word’s size, so under a perfect tiling the placement is just the running sum of part sizes (serial mode computes this as an AR1 increment over the previous .where; the for-loop is time). WS may still supply **type** information for the tiling — *word, space, word, punct, ...* — even where it does not supply coordinates.

Tiling, subspace sizes, and consciousness

Whether the parts/wholes **perfectly tile** the input (a partition — so order alone reconstructs, with no gaps between parts) is **NOT** entailed by parallel-vs-serial mode. It is entailed by the **relative sizes of the InputSpace subspace and the PS/WS subspaces**:

- In **serial mode** they are *forced* to match — the for-loop traverses **all** of input space — so the tiling is always perfect.
- In **parallel mode**, if they do not match because InputSpace is *larger* than PS/WS, that bounded mismatch is exactly where the two attentions **select what is most relevant to consciousness**: PS/WS cannot hold all of InputSpace, so attention picks the salient subset to surface. (When $\text{InputSpace} \leq \text{PS/WS}$ — as in the current test fixtures — the tiling is again a partition and order suffices.)

C. Parse time (two attentions)

At parse time CS receives the overcomplete input representations from the PS and WS mereological towers. Attention parses them two ways:

- **Top-down attention is vertical** (w.r.t. the bottom-up / top-down towers): it fixes the **Basic Level** of analysis — the size of parts and wholes. The word-level isolates the regions that are type *word*, *punctuation*, and *space*, forming a **complete tiling** (the WS→PS `_passback_scope_where` handoff; punctuation already tiles to its own span).
- **Bottom-up attention is horizontal.** Top-down may fix the word boundary of *wheelhouse*, but what makes *wheel* + *house* its building blocks rather than the equally-segmentable *wheelhou* + *se*? **Greedy longest-match** is the current easy approximation (`RadixLayer.longest_match`); bottom-up attention should guide both parsing and reconstruction here.

D. Attention indexing (`.where` / `.when` / codebooks)

The three addressing roles are disjoint: **.where indexes over the input buffer** (positional; period = $\frac{1}{2} \cdot \text{InputSpace}$), **.when indexes over LTM** (to-build — LTM is content/parthood-addressed today), and the **codebooks are content-addressable** (identity is the row index; the cross-codebook `.where` slice registry was retired).

Cognitive grounding: dense-perceptual vs sparse-symbolic (2026-07-02)

The design splits representation into a **dense, invertible, subsymbolic** integrator (the corpus-callosum mixing matrix, which mixes part/whole *content* in high dimension) and a **sparse, symbolic** composer (the `SparseLayer` families, which compose *scalar activations* of named concepts). This is not an arbitrary engineering choice; the split, and specifically *why sparsity belongs only on the symbolic side*, tracks several convergent findings on human cognition.

- **Complementary Learning Systems** (McClelland, McNaughton & O'Reilly 1995). Neocortex uses slow, dense, *overlapping* distributed codes that extract statistical structure across experiences; hippocampus uses fast, *sparse, pattern-separated* codes to bind individuated episodes. The reason hippocampal codes are sparse is **interference avoidance among individuated bindings**, not compression: dense overlapping codes suffer catastrophic interference when made to hold many discrete conjunctions. This is the cognitive answer to “why sparse, and why only symbolically” — the perceptual manifold wants the dense mixing; the binding of individuated things into reusable conjunctions (the joint/sentence concept over word-symbols) wants sparsity. A **symbol is a scalar handle on a grounded direction** (the atom), which is precisely a hippocampal index into cortex; the transpose decode ($\tanh(W^\top y)$) is pattern **reinstatement**. Caveat: this maps the *sparse-symbolic vs dense-perceptual* axis onto hippocampus-vs-cortex — NOT the ramsified *orders* (abstraction is a separate, graded anterior-temporal / prefrontal axis).
- **Dual-process cognition** (Sloman 1996; Kahneman). A fast, parallel, similarity-driven associative system and a slow, serial, rule-based symbolic one. The subsymbolic (parallel content-mixing) and symbolic (sparse relational) loops instance this split. The sparsity and the serial capacity-limit are the *same fact*: a sparse composer has low fan-in per concept, the computational shadow of working-memory span (Miller’s 7 ± 2 ; the STM depth ≈ 8 the model already carries).
- **The neuro-symbolic interface / systematicity** (Fodor & Pylyshyn 1988). The architecture is a concrete stance on the oldest fight in cognitive science: **content stays connectionist** (dense, continuous, invertible mixing), **structure becomes symbolic** (discrete, reusable, composable edges), and the **interface is the activation readout** — the point where mixed content is read out as a scalar activation of a named thing. The discrete edges buy the compositionality/systematicity that pure distributed codes are accused of lacking, while the dense mixing keeps perception continuous and gradient-trainable.
- **Grounded cognition** (Barsalou 1999; contra amodal symbol systems). Because the *direction* lives in the concept’s atom and only the *activation* is abstracted to a scalar, the symbols are not amodal Fodorian tokens — they are grounded pointers-with-magnitude, closer to perceptual-symbol “simulators.” Keeping magnitude (the normalized-*sum* presence readout, not cosine) is cognitively

load-bearing: graded activation *is* typicality / salience (Rosch’s graded membership), which a cosine would discard.

- **Basic-level categories are perceptual, not content-free** (Rosch et al. 1976). “Subsymbolic” $\neq$ “category-free”: the mixing output at order 0 is already carved toward basic-level regions, because that is what perceptual integration *for a categorizing organism* produces. This suggests the EMA-snapped dictionary atoms are the **pre-linguistic Gärdenfors regions** (prototype centers, basic-level, perceptual) and the sparse symbol graph is the *post-linguistic* labelling-and-composition that points at them — three cognitively distinct stages (integrated field $\rightarrow$ unnamed category region $\rightarrow$ named composable symbol), not two layers with a bookkeeping detail between.

Where the mechanism is deliberately cleaner than cognition. (1) The order-0 boundary is a *default flow*, not a wall: perception is concept-penetrated (top-down / predictive coding), and the model’s attention/priming loop is what carries that penetration back down — the division must never become impermeable. (2) *Invertibility* is instrumental (reconstruction, gradient), not a biological claim; brains approximate and predict, they do not compute exact inverses. The extensional/intensional semantics this grounding implies for a **single** conceptual space is developed in BasicModel.md “Conceptual Space.”

References. McClelland, McNaughton & O’Reilly (1995), *Why there are complementary learning systems in the hippocampus and neocortex*, Psychological Review 102(3). Sloman (1996), *The empirical case for two systems of reasoning*, Psychological Bulletin 119(1); Kahneman (2011), *Thinking, Fast and Slow*. Miller (1956), *The magical number seven, plus or minus two*, Psychological Review 63(2). Fodor & Pylyshyn (1988), *Connectionism and cognitive architecture*, Cognition 28. Barsalou (1999), *Perceptual symbol systems*, Behavioral and Brain Sciences 22(4). Rosch, Mervis, Gray, Johnson & Boyes-Braem (1976), *Basic objects in natural categories*, Cognitive Psychology 8(3). Gärdenfors (2000), *Conceptual Spaces: The Geometry of Thought*, MIT Press. Olshausen & Field (1996), *Emergence of simple-cell receptive field properties by learning a sparse code*, Nature 381 (sparse *activation* over a dense dictionary — the dictionary/atom side here — as distinct from the sparse *relational graph* of the symbol families).

Overview

BasicModel is a bidirectional neural architecture organized as a pipeline of five **spaces** plus a symbol host (SymbolSpace), each implementing a distinct representational transformation:

Forward: InputSpace $\rightarrow$ PartSpace $\rightarrow$ ConceptualSpace $\rightarrow$ WholeSpace $\rightarrow$ OutputSpace
Reverse: OutputSpace $\rightarrow$ WholeSpace $\rightarrow$ ConceptualSpace $\rightarrow$ PartSpace $\rightarrow$ InputSpace

The pre-2026-05-27 “two feedback loops” ($S \rightarrow C$ symbolic loopback per stage, $C \rightarrow P$ subsymbolic loopback cross-forward) collapse under the substrate refactor:

- **Subsymbolic loop dissolves.** PS is a single-direction input processor. No recurrent $C \rightarrow P$ feedback at the substrate level. In PARALLEL mode, iteration happens by passing CS to the same PS.forward(x) for T refinement passes (the <subsymbolicOrder> knob).
- **Symbolic loop generalizes** to pairwise grammar ops over STM, dispatched by the signal router (LanguageLayer). Lift and Lower join the same GrammarLayer dispatch surface as Intersection, Union, etc.

The forward pass transforms raw input into predictions; the reverse pass reconstructs the original input from the symbolic representation. Both directions are trained simultaneously with a single optimizer minimizing a combined loss:

`totalLoss = (1 - reconRatio) * outputLoss + reconRatio * reconstructionLoss`

The legacy SubwholeSpace and SyntacticSpace classes have been retired. The subsymbolic role is filled by PartSpace; syntax / grammar dispatch lives on SymbolicSubSpace.languageLayer (the signal router, which subsumes the retired Chart). The MereologicalTree sidecar that backed part / equals / query is also retired — those operations are pure-geometric clipped-cosine projections over WholeSpace codebook activations.

PartSpace and **WholeSpace** (renamed 2026-06-12 from **PerceptualSpace** / the original **SymbolSpace**) both subclass a thin shared **PerceptualSpace(Space)** base: both views are perceptual. At the corpus callosum, objects are analysed and synthesized by sending them back to **PerceptualSpace** — wholes get split, parts get chunked. In symbolic “mode” the objects sent back are symbols. Terminologically there are objects and references; a reference is a *sign* (a quantized version of the referent) or a *symbol* (an unrelated version of the referent, of much lower dimensionality). The freed name **SymbolSpace** was **reintroduced 2026-06-19** with new semantics — it is now the grammar/word space-role (formerly **WordSpace** / **WordSubSpace**, abbrev **ss**; the **WholeSpace** stream is now **ws**). See the full rename mapping in `doc/old/2026-06-19-handoff.md`.

The corpus callosum **builds a single meronymy out of the two towers**: a part **A** (**PartSpace**) and a whole **B** (**WholeSpace**) carry **.what** codes from different codebooks (incomparable), but their **.where** is comparable, so the callosum links **A isa B** (token **isa** type) when **A.where** is contained in **B.where** with no greater-part/lesser-whole intervening. Word object — too unlike to link directly — is bridged by a **second-order meta-object** (synthesized in **PartSpace**, outside **.where/.when**: the **MetaSymbol**). The correctness signal is the **part/whole ratio** (many-parts→one-whole = under-analysed; one-part→one-whole = over-analysed), which requests further -synthesis / -analysis in the offending **.where** — and is the principled fix for the **MM_20M** mean-collapse. Full design: `doc/old/mereological-order-raising.md`.

Spaces

Space	Role
InputSpace	Lifts raw data into working dimensionality; surface tokenization
PartSpace	Bottom-up SYNTHESIS branch (Pi/Sigma swap, rev. 2026-06-09): sigma fold + <synthesis> front
ConceptualSpace	STM container + main grammatical CPU + (when sparse-active) the POST-PUMP symbolic phase
WholeSpace	Top-down ANALYSIS branch: pi fold + <analysis>/<lexer> knobs; unified word lexicon codebook
OutputSpace	Final prediction

The cross-space fold contract has changed:

```

PS.forward(x): return self.pi(x.materialize()) + self.sigma(x.materialize())
CS.forward(subspace, word_subspace):
    STM[0..6] = STM[1..7]; STM[7] = folded          # STM shift/push only (mode-dispatched); the pump
# POST-PUMP cutover (sparse-active, once per forward, in _forward_body):
#   content, acts = cs.cs_symbolic_phase(last_cs.materialize()) # snap -> ramsified composition
#   last_cs._concept_activations = acts; SS leg built ONCE; SBOW parks the settled slab
SS: no atomic forward operator; hosts insert_paired_word + write-required grammar ops
    (the CS->SS symbol bind leg is SymbolSpace.forward_concept_to_symbol, .forward()-mediated)


```

The legacy composition **C = sigma_percept(pi_input(IS) + pi_concept(C_prev))** is **retired**. Per-stage feedback is absent at the substrate; grammar dispatch over STM provides the recurrent character via the signal router.

See `Spaces.md` Section “Sigma / Pi ownership” for the cognitive rationale and the migration trail. See `Logic.md` Section 8 for the algebraic constraints on sigma/pi.

Dimensions (**nDim**) are read from **TheObjectEncoding**. Codebook sizes (**nVectors**) are likewise on **TheObjectEncoding**; the factory validates **nVectors >= nActive**.

MM_5M: Hierarchical Progressive Bottleneck ($N \times D = 4096 = \text{const}$)



Layer selection by `invertible` (the `<reconstruct>` element / `reconstructEnum` were **RETIRED** in A1, 2026-06-09; reconstruction is now seeded from concepts **unconditionally**, gated only by `reconstructionScale`):

1. **Non-invertible, forward-only layers** (`PiLayer`, `SigmaLayer`): forward-only, no reverse pipeline.
2. **invertible**: Single invertible layer (`PiLayer(invertible=True)`, etc.) serves both directions, sharing weights.
3. **Not invertible** (but reconstructed): Two layers with separate weights — `forward()` on one, `reverse()` on the other. Avoids the expressivity limitation where a non-invertible layer can't represent the inverse of another. Reverse uses matrix `pinv` (may be numerically unstable from SVD convergence). `<invertible>true</invertible>` avoids this via shared-weight inversion.

Reconstruction Symbols

The symbolic bottleneck can lose information needed for reconstruction. XOR maps 2 inputs to 1 output, but `XOR(0,0)=0` and `XOR(1,1)=0` are distinct inputs producing the same output. A single output symbol cannot reconstruct which input was presented.

`nSymbols` is split:

- `nOutputSymbols` = `OutputSpace.nActive` — fed to `OutputSpace` for prediction
- `nReconSymbols` = `nSymbols` - `nOutputSymbols` — carried in `end_state` for reconstruction

A skip connection through the symbolic bottleneck. Reconstruction symbols receive gradient only from reconstruction loss.

Single Optimizer with Overlapping Weight Spaces

The forward and reverse passes share a **single Adam optimizer** that minimizes the combined loss. Forward and reverse weight spaces **partially overlap** — neither disjoint (allowing independent optimizers) nor identical (creating destructive interference). Some layers share weights between directions (shared embeddings, the symbolic bottleneck); others are direction-specific (**pi1/pi2**, **sigma1/sigma2**, **linear1/linear2**).

- **Shared weights** receive gradient from both losses, learning representations useful in both directions.
- **Direction-specific weights** specialize without interference.
- **No ping-pong**: separate optimizers on overlapping parameters would pull weights in alternating, conflicting directions each step.

When **invertible=true**, overlap is total: one invertible layer serves both directions and receives the full combined gradient.

Reference: A.M. Rogers, T.T. Shannon, and G.G. Lendaris, “A comparison of DHP based antecedent parameter tuning strategies for fuzzy control,” *Proceedings Joint 9th IFSA World Congress and 20th NAFIPS International Conference*, 2001, doi: 10.1109/NAFIPS.2001.944317.

Training Loop

Single Adam optimizer with persistent state (momentum/variance accumulate across epochs):

1. Forward pass: input $\rightarrow$ prediction + **end_state**
2. Compute **outputLoss** from prediction vs. target
3. Reverse pass: **end_state** $\rightarrow$ reconstructed input
4. Compute **reconstructionLoss** from reconstruction vs. original input
5. Backpropagate combined **totalLoss**
6. If ergodic: run **paramUpdate()** (gradient energy sensor updates alpha)
7. Optimizer step (embedding params excluded when **trainEmbedding** is NONE, CBOW, or SBOW)
8. If **trainEmbedding** is CBOW, SBOW, or BOTH: run embedding update step

Alpha annealing (ergodic): starts at 1.0 (full exploration), decays to 0.0 within the first 5% of epochs via **alpha = max(0, 1 - epoch / warmup)** where **warmup = numEpochs // 20**. Code convention (**alpha=1** means explore) is the inverse of the ergodic math convention; layers translate internally.

See Params.md for all XML parameters. See Training.md for embedding modes.

The three cognitive operations (2026-06-14)

Processing decomposes into three operations, in increasing order of abstraction. Each maps to a knob (or, for the first, to the folds themselves):

1. **Granularity of analysis and synthesis** — done *automatically* by the two perceptual views’ folds, per pass. PartSpace’s **Sigma synthesizes** (union; count-reducing: many atoms $\rightarrow$ fewer chunks); WholeSpace’s **Pi analyses** (intersection; count-increasing: one unity $\rightarrow$ many parts). How finely the scene is carved, or how coarsely it is chunked, is set by the folds — there is no separate granularity knob. The InputSpace feeds the two views directly: the **Atom** view ([B, N, D], which PartSpace synthesizes bottom-up) and the **Universe** view (**_unity_view**, the whole as one event, which WholeSpace analyses top-down). Optionally (<mereologyRaise>), perception builds a meronymic lattice over the towers and **raises abstraction order** as attention requires — see Mereology.md $\rightarrow$ Order-raising and doc/old/mereological-order-raising.md.
2. **Subsymbolic order** (<subsymbolicOrder>) — *iterating* the folds: codes are passed back to PartSpace / WholeSpace across **subsymbolicOrder** passes (the CS $\rightarrow$ PS loop). Synthesis chunks the codes into higher-order percepts (fewer each pass); analysis re-expands, attention selecting what to expand (a top-k over the priming, applied after the WholeSpace codebook lookup). Symbolic composition is no longer a separate CS $\rightarrow$ PS passback flag: the recurrent symbolic leg always

flows through `WholeSpace.forward(prevCS_forSS)`, and the symbolic-iteration codebook handles higher-order symbolic composition on the CS→SS path.

Proposed refinement (mereological-order-raising spec). This single subsymbolic loop is really **two** moves CS should choose between *per representation*, by reading the **contiguity of .where**: a *contiguous* extent → **refine granularity** (chunk finer / tile — drive the radix), same order; a *discontiguous* extent → **raise order** (another / fold, lifting out of **.where/.when**); a *zero .where* → null. The number of contiguous runs in a whole’s **.where** is its part/whole ratio, so the same read also routes integrate-vs-disintegrate (→PartSpace vs →WholeSpace). See doc/old/mereological-order-raising.md “The three-aspect loop”. As of 2026-06-16 the contiguity read has its substrate: **.where / .when** are **endpoint-sum brackets [start, end]** (`WhereEncoding.decode_span`), so extent and gaps are read directly off a code (a zero-extent instant vs a span); see doc/Spaces.md.

3. **Symbolic order** (`<symbolicOrder>`) — the symbolic / relational loop budget. In serial mode (`<serial>true</serial>`), words are read **one at a time** from WholeSpace (reading isolated words to ConceptualSpace *is* attention) and processed grammatically in ConceptualSpace’s STM and on the PartSpace side. `symbolicOrder` limits how many symbolic / SS loops may run; `<serial>` selects whether the per-word traversal is active.

So: granularity is intrinsic to the folds, subsymbolic order iterates the subsymbolic passes (composing higher-order percepts), symbolic order budgets the relational pump, and **serial** selects the serial grammatical loop over words.

Current order semantics. This section supersedes the older mode-selector wording in doc/old/orders.md. The three order axes now have separate semantics, bounds, and composition rules:

- **subsymbolicOrder** — the **analysis/synthesis refinement-pass count and the area of attention**. T parallel CS→PS/WS iterations; each pass *refines* (contiguous **.where**) or *raises* (discontiguous), and attention scopes via a **.where** on the dual-input SECOND ARGUMENT (the top-down WS→PS handoff, gated `<mereologyRaise>`; see mereological-order-raising.md). The serial-word reading supplies word **.wheres** through the **same** channel.
- **symbolicOrder** — the **relational pump** budget. It spreads activation through the relation graph to surface *higher-order* (relations-of-relations) features that have **no mereological .where** and so can’t be primed off **.where** contiguity. `subsymbolicOrder` pumps the mereological substrate; `symbolicOrder` pumps the relational one. `<serial>` separately selects whether traversal is per-word serial or whole-slab parallel.
- **syntacticOrder** (*NEW — implemented 2026-06-19*) — the **parse-tree composition DEPTH** per sentence, bounded by the word count. 0 = unbounded (byte-identical); a positive value caps the NULL-seal reduce sweep to that many fold levels (static `min(syntacticOrder, cap-1)`; W structural). Inert in parallel mode.

Composition (serial run): `<serial>true</serial>` loops words × `syntacticOrder` bounds the parse-tree depth per sentence × `subsymbolicOrder` pumps per node; the **basic-level stop** is shared (synthesis halts at words, so the tree’s leaves are words). `syntacticOrder` **layers over** the serial traversal loop (it bounds depth; it does not replace the parallel-vs-serial switch).

Where this is headed (historical design note in orders.md §6): the three orders become **pump counts** over one connectionist attention substrate — a cumulative priming hierarchy (mereological entries → relations/concepts → higher-order, each seeing all below) where reading is a learned **.where** attention (text-mode next-word loss) that replaces the serial for-loop.

Modes of operation

Two operating modes, selected by `<architecture><serial>` (replaced the `conceptualMode` enum; legacy configs that omit **serial** derive the mode from `symbolicOrder > 0`):

Mode	Trigger	PS.forward argument	Iterations
SERIAL / GRAMMATICAL	<code><serial>true</serial></code>	IS_t per word	one per word; PS push
PARALLEL	<code><serial>false</serial></code>	IS once, then CS for T-1 iterations	T = <code><subsymbolic0></code>

SERIAL and GRAMMATICAL are not architecturally distinguished — grammar dispatch is a chart / rule-catalog config, not a substrate mode. PS.forward takes a single positional argument in both modes; the argument is whatever input is being processed (IS in SERIAL, IS then CS in PARALLEL refinement).

Pre-2026-05-27 “two feedback loops” retired. The legacy $S \rightarrow C$ symbolic loopback (per-stage) and $C \rightarrow P$ subsymbolic loopback (cross-forward) collapse under the substrate refactor:

- PS is a single-direction input processor; no recurrent $C \rightarrow P$ feedback at the substrate. CS state enters PS only via `PS.forward(CS)` in PARALLEL mode’s refinement iterations.
- Symbolic loop becomes pairwise grammar ops over STM (the signal router’s copy/reduce dispatch). `Lift` and `Lower` are binary GrammarLayer subclasses dispatched alongside `Intersection`, `Union`, etc.

The recurrent character of the architecture lives in (a) STM accumulation across words in SERIAL mode, and (b) the T-pass PARALLEL refinement loop. Cross-call serial-cache (`subspace.serial_cache`) for streaming / autoregressive contexts is preserved; gated by `PartSpace._recurrent_pass_idx == 0`.

Pipeline as a unit, two-space-role reset

`runBatch` is a pure compute brick: forward $\rightarrow$ loss $\rightarrow$ backward $\rightarrow$ optimizer.step. It does **not** decide when to reset per-row state, does **not** consume `_end_of_stream` for control flow, and (after Section 6 vectorization) does **not** issue any GPU $\rightarrow$ host sync inside the brick.

Reset lives in `runEpoch`. The same loop drives both byte cursor (AR text byte) and trial cursor (non-AR); `next_tick` is universal dispatch:

```
while not ds.all_done():
    inp, out, hard_eos = ds.next_tick()           # 3-tuple, host-side
    runBatch(inp, out)                           # compute brick
    flush_word_buffers()                         # materialize subspace.word
    dispatch_per_row_reset(hard_eos)             # hard resets
    dispatch_soft_reset()                        # grammar <start> reductions
    post_tick_compact()                          # truth_layer.compact
```

For AR text byte, `inp` is a byte slab and `hard_eos[b]` flips True when row `b`’s cursor exhausts a doc. For non-AR / numeric data, each tick yields one batch of trials with `hard_eos = [True] * B`.

Hard reset. `TheData` walks each document one slab of ≤ 1024 bytes at a time. `hard_eos` flips True on cursor exhaustion. Full row-state cascade fires for that row only; other rows continue mid-document with state preserved.

Soft reset. The active parser signals when a row’s parse reduces to `<start>`. `wordSpace._sentence_completed` is drained per-tick: re-arms `_stm_fired[b]`, clears `_last_svo[b*K..]` and parse-stack rows for `b`, but **preserves discourse history** (discourse accumulates across sentences within a document and clears only on hard reset).

No truncation. Documents longer than `slab_bytes` span multiple ticks; concatenating per-tick slabs for any row reproduces the original document byte-exact. `valid_mask: [B, K]` handles partial-fill tails via NULL-padding.

Compute-brick contract. No `.item()`, no `.tolist()`, no Python conditional on a tensor value, no GPU $\rightarrow$ host copy inside `runBatch`. The chart’s residual `.tolist()` calls retired with the `Chart` class itself in the substrate refactor; remaining graph-break sites are documented in `doc/BrickHostSyncStatus.md`.

Two-File Architecture

File	Contents
XML config (e.g. <code>BasicModel.xml</code>)	Architecture, hyperparameters
Weights checkpoint (e.g. <code>BasicModel.ckpt</code>)	Full integrated bundle: model parameters, register-buffer state, embeddings

The 2026-05-12 *integrated-weights* refactor retired the separate `.kv` embedding artifact: embeddings, vocabulary mappings, and the BPE codebook now ride inside the single `.ckpt` bundle alongside the model's other parameters. The bundle layout is:

- `state_dict`: every `nn.Parameter` and `register_buffer` in the module tree (model weights, `wv._vectors`, `TruthLayer.truths`, etc.) — serialised by the normal PyTorch path.
- `vocab_extras`: the WordVectors Python-side mappings that don't live in `state_dict` (`index_to_key`, `counts`, `total_count`).
- `bpe_extras`: the `ChunkLayer`'s pure-Python state (merges list, vocab dict, `id_to_bytes`, growth cursors). Required because `ChunkLayer` stores its merge table as Python dicts/lists, not tensors.

`bin/embed.py` still produces standalone `.kv` artifacts for CBOW/SBOW *pre-training* studies, but those artifacts are no longer part of the runtime artifact set. Cold-start training initialises the vocabulary and BPE codebook from scratch and learns them end-to-end alongside the model weights.

Language System

The grammar dispatch runs through the **signal router** (`LanguageLayer`, `bin/Language.py`) — the single canonical parser. `SymbolicSubSpace` owns it directly as `self.languageLayer`. The pre-substrate CKY `Chart` and STM shift-reduce parsers retired in Stage 3 of the substrate refactor.

The signal router represents STM as a slab `[B, N, D]`; per-layer scorers emit per-position copy/reduce scores; `binary_tiling_soft_dp` produces marginals at training (soft superposition), `binary_tiling_viterbi` produces the best tiling at eval. `Grammar.rule_probability(body)` is generalized to the per-position, per-op score head — dormant defaults (fold ops fire, negation ops don't) carry over as initial biases. Single-application enforcement via `_fired_bodies` / `reset_derivation` carries over unchanged.

Grammar ops are `GrammarLayer` subclasses dispatched by the router as unary (copy-side) or binary (reduce-side):

- **Unary symbolic operators**: `not(S)`, `non(S)`, `swap`, `copy`, `true`, `false`.
- **Binary symbolic operators**: `intersection(S, S)`, `union(S, S)`, `conjunction`, `disjunction`.
- **Mereological operators**: `part(S, S)`, `equals(S, S)`, `query(S, S)`. Pure-geometric — the `MereologicalTree` sidecar that formerly stored explicit parent / equality links retired in favor of clipped-cosine parthood on codebook activations. See `Mereology.md`.
- **lift and lower**: now binary `GrammarLayer` subclasses (Stage 4 of the substrate refactor). Each owns an internal `SigmaLayer` (`LiftLayer`) or `PiLayer` (`LowerLayer`) for the pairwise math. No longer “substrate-borrowing” — fully self-contained binary grammar ops with `arity=2`, `space_role='C'`. Typed grammar signatures still determine result order (e.g., `S4 = lift(NP3, VP1)`). See `Language.md`.
- **Butterfly mode on GrammarLayer** (Stage 5): all `GrammarLayer` subclasses accept `butterfly=True`, `N=N` for efficient cross-STM pairwise composition via a packed `nn.Parameter[n_levels, N//2, 2D, 2D]` cascade with bit-reversal permutations. Wired into the space folds (`PartSpace.sigma` / `WholeSpace.pi` post the `Pi/Sigma` swap, rev. 2026-06-09) by default in butterfly-enabled configs. Closes the XOR convergence target.

Parthood (`part`) is the **fundamental** mereological operation, realized as clipped cosine projection on symbolic activations. The full suite (`whole`, `equal`, `overlap`, `underlap`, `boundary`) composes through `part` on `Basis`. `equals(S, S)` is propositional identity on `S`; delegates to `Basis.equal`.

Short-Term Memory on ConceptualSpace

`ConceptualSpace.stm` (an instance of `ShortTermMemory`) is a per-batch stack of unquantized CS “ideas” — the working set the signal router dispatches grammar ops over. Capacity defaults to 7 (Miller, ± 2); `<ConceptualSpace><stmCapacity>N</stmCapacity></ConceptualSpace>` overrides.

Post-substrate-refactor, `CS.forward(subspace, word_subspace=None)` is **STM bookkeeping** — shift slots left, push the new idea onto slot 7. The legacy atomic forward fold (`sigma_percept`) is retired. (The ramsified symbolic transform no longer runs inside `forward` — P3 two-phase rework: the pump stays purely subsymbolic and `cs_symbolic_phase` fires ONCE at the post-pump cutover in `_forward_body`; its outputs feed the SS leg and the losses, never the STM content. See `Architecture.md` sec A.) The mode dispatch:

- **SERIAL / GRAMMATICAL**: one idea pushed per word; STM shifts (oldest dropped from slot 0). Grammar ops dispatched per word or at sentence boundary.
- **PARALLEL**: `T = <subsymbolicOrder>` iteration outputs written to STM slots simultaneously; no shift.

STM is cleared on hard **Reset** (sentence boundary) and survives soft reset. The signal router consumes `stm.snapshot()` for its slab input. See `Spaces.md`.

The STM data model, the predict-then-perceive cadence (serial and parallel), the in-STM and inter-sentence predictors, masked-word reconstruction, relative-vs-absolute end-states, and the LTM chain are documented in full in the dedicated `STM.md` chapter. Note in particular that **serial mode runs WITH attention by design** (the old serial-vs-attention guard was lifted; `MentalModel.xml` is `serial + hasAttention`) — see `STM.md` Section 4.

Per-word operational flow (SERIAL mode)

In SERIAL / GRAMMATICAL mode, each word traverses a per-word path:

```
byte stream -> PS.forward(IS_t)      # MPHF surface lookup → PS lexicon
                                     # then pi(x) + sigma(x), no outer tanh
               -> CS.forward(idea)   # STM shift; push idea onto slot 7
               -> signal router dispatches grammar ops over STM
                                     # (read-only via CS; write-required via SS)
```

PS’s `self.vocabulary` (Embedding) holds the per-word vectors keyed by MPHF. SS’s codebook holds **paired (orthographic, semantic) rows** for the unified lexicon — the orth row is a copy of PS’s per-word vector at insert time; the semantic row is random; the two are parented via `Codebook.set_part_parent`. Lookup chain: surface → MPHF → orth row → parented semantic row.

POS rides the codebook for free. The `WholeSpace` reference codebook carries two POS-bearing fields per atom: `category_ids`: [V] (hard POS tag — one of the grammar’s nonterminals) and `category_logits`: [V, C] (learnable soft POS distribution per atom). So per-word snap returns (`word_id`, POS) simultaneously — no separate POS tagger needed. The router uses POS for typing reduce candidates; POS is learned through parsing alongside the codebook.

POS rides the codebook for free. The `WholeSpace` reference codebook carries two POS-bearing fields per atom: `category_ids`: [V] (hard POS tag — one of the grammar’s nonterminals, e.g. NP / VP / N / V / ADJ) and `category_logits`: [V, C] (learnable soft POS distribution per atom, EMA-updated by the chart’s `_apply_codebook_pos_seed` mechanism in `Language.py`). So per-word snap returns (`word_id`, POS) simultaneously — no separate POS tagger needed. The parser uses POS for typing reduce candidates (`NP + VP → S`, etc.); POS is *learned through parsing* alongside the codebook.

See `Logic.md`, `Mereology.md`, and `Language.md`.

Shamatha Speech target. Planned narrow grammar for one-pointed object speech: complete DNF over active percepts, permitting each **conjunction / disjunction** only when operands’ `where()` supports are connected and `when()` supports are continuous. See `Language.md`.

Sigma and Pi Layers

For weight matrix $W \in \mathbb{R}^{m \times n}$ and input $x \in \mathbb{R}^n$:

Sigma layer:

$$y_j = Wx + b = b_j + \sum_{i=1}^n W_{ji}x_i$$

Pi layer (log-space linear):

$$\begin{aligned} s_i &= \log \frac{1+x_i}{1-x_i} = 2 \operatorname{atanh}(x_i) \\ z_j &= \sum_i W_{ji} s_i + b_j \\ y_j &= \frac{e^{z_j} - 1}{e^{z_j} + 1} = \tanh(z_j/2) \end{aligned}$$

Forward maps $[-1, 1] \rightarrow (0, \infty)$ via `_to_mult`, `log`, `linear`, `exp`, `_from_mult`. Domain and range both $[-1, 1]$. Reverse inverts each step: `_to_mult(y)`, `log`, $W^{-1}(z - b)$, `exp`, `_from_mult`.

Motivation. The classical product form $y_j = b_j \prod_i (1 + W_{ji}x_i)$ becomes, after taking logs, a sum. The code moves into a log-multiplicative domain via `atanh`, performs a linear op there, returns via `tanh`. The `atanh` transform stretches values near ± 1 toward infinity, making the layer sensitive to strong activations.

Monotonicity of the lift / lower chain. Under `monotonic=True`, Pi/Sigma select non-negative linear layers, giving $W \geq 0$. Positive matrices are monotone on the positive cone, so lift / lower preserve parthood for activations represented in that cone. Truth-set bivectors remain live for user-supplied truths; they are explicit truth/operator surfaces rather than a space-wide output mode. See `Spaces.md`.

Dimensionality Constraints

- Input layer output dim = perceptual layer output dim (conceptual operates on both).
- Symbolic layer input dim = perceptual layer input dim (both operate on conceptual output).
- Output layer input dim = sum of symbolic layers' output dims.

Invertible Linear Layer (LDU)

Factors $W = L \cdot D_{\text{embed}} \cdot U$:

- L : unit lower-triangular ($nIn \times nIn$, diagonal = 1).
- D : diagonal vector of length `rank = min(nIn, nOut)`, embedded into $[nIn, nOut]$ by zero-padding.
- U : unit upper-triangular ($nOut \times nOut$).

Exact inverse via triangular solves: $W^{-1} = U^{-1} \cdot D^{-1} \cdot L^{-1}$. Each factor inverted by `torch.linalg.solve_triangular`. No SVD; inverse exact when all D entries are nonzero. Parameter count: $nIn^2 + \text{rank} + nOut^2$. Initialized at $L = I, d = 1, U = I$ (identity).

`naive=False` (default) applies L/D/U sequentially without materialising `W_eff` as a full matrix. `naive=True` materialises `W_eff` and its inverse.

Ergodic noise injection at the factor level, plus the `stable=True` clamp and the noise lifecycle, are documented in `Ergodic.md`.

Ergodic Exploration

See `Ergodic.md`.

Sentence-level AR (`InterSentenceLayer`)

Within-sentence training is IR-only (BERT-style masked-LM at the subsymbolic (PS)); see `doc/Spaces.md` Section “Within-sentence AR retirement”). The **autoregressive** signal in this architecture lives one scale up: between sentences, on a per-sentence representation `s_t`. That’s the job of `InterSentenceLayer` (alias `wordSpace.discourse`).

Sentence representation

`s_t` is the **root SS slot** of the body’s final stage: the single vector the start-symbol reduction wrote into. The chart’s parse trace already commits to this slot at sentence end; the layer pools `[B, N, D] → [B, D]` by taking row 0 (root). Width is `sentence_dim = n_dim` (one vector per row), **not** the full `n_symbols * n_dim` flatten that the pre-2026-05-14 contrastive layer used — that broader rep would have blown the predictor’s Linear past the allocator budget on large MM_5M-scale configs.

ARMA(p, q) predictor

`InterSentenceLayer` runs an autoregressive moving-average predictor:

```
s_hat_t = predictor(s_{t-1..t-p}, e_{t-1..t-q})
e_t      = s_t - s_hat_t
loss     = MSE(s_hat_t, s_t)           # accumulated per batch
```

- `p` = AR lag count (default 5) — last `p` sentence reps.
- `q` = MA lag count (default 2) — last `q` prediction errors.
- `predictor = nn.Sequential(Linear(p*D + q*D, H), Tanh, Linear(H, D))`, with `H = min(1024, 2*sentence_dim)`.

The MA term lets the predictor correct for systematic bias in the AR extrapolation: if the AR model consistently under-predicts the sentence rep, the residual `e_t` carries that signal forward.

Buffers (per row, non-persistent):

- `_s_history`: `[B, p, sentence_dim]` ring of last `p` sentence reps (most recent at index -1).
- `_e_history`: `[B, q, sentence_dim]` ring of last `q` residuals.
- `_s_count / _e_count`: `[B]` long, fill levels (cap at `p / q`).

`ensure_batch(B)` resizes these on cascade from `SymbolSpace.ensure_batch`; `Reset()` clears them on hard / discourse boundary. Default behaviour is to **not** auto-reset across document boundaries — the AR lags carry information through discourse continuity unless the caller explicitly calls `Reset`.

Wiring into the training loop

1. After the body finishes (sentence end), `_forward_per_stage` stashes the SS event on `_current_discourse_s`.
2. In `runBatch`, when training, `discourse.observe(s_tensor)`:
 - Pools `s_t = sigma_S(s_tensor[:, 0, :])`.
 - Computes `s_hat_t = predictor(_s_history, _e_history)`.
 - Returns `MSE(s_hat_t, s_t)` (None on the first call per row when the ring is empty).

- Computes $e_t = s_t - s_{\text{hat}_t}$, pushes both into the rings.
3. `runBatch` adds the loss to `TheError` under category "discourse" with weight `armaScale` (training XSD knob, default 0.1).

Inference (chat-loop seeding)

`BasicModel.generate_sentence(seed_text):`

1. Calls `discourse.predict_next()` for the ARMA-predicted sentence-rep prior `s_hat_{t+1}`.
2. Lifts it through `discourse.cast` to `concept_dim` and stages on `ConceptualSpace._c_prior`.
3. Runs the IR forward — the body's first `sigma_percept` output gets `_c_prior` summed in as a sentence-level conditioning bias before the codebook lookup. Cleared after the forward consumes it.
4. The post-body perceptual event is decoded by nearest-neighbour against the perceptual codebook, producing the (`slot`, `original`, `predicted`) triples for the seed text's masked positions.
5. Commits the produced sentence's SS root to the ARMA ring via `discourse.observe(s_tensor)`.


The IR head plays no role at inference — the prediction lives at the masked subsymbolic (PS) positions, decoded against the (frozen) perceptual codebook.

Configuration

XSD knob	Section	Default	Notes
<code><armaP></code>	<code><SymbolSpace></code>	5	AR lag count
<code><armaQ></code>	<code><SymbolSpace></code>	2	MA lag count
<code><armaHiddenDim></code>	<code><SymbolSpace></code>	<code>2*sentence_dim</code> (cap 1024)	predictor hidden width
<code><armaScale></code>	<code><architecture><training></code>	0.1	ARMA loss weight added to <code>TheError</code>
<code><sentencePrediction></code>	<code><architecture><training></code>	false	Gates <code>InterSentenceLayer</code> constr

The retired pre-2026-05-14 knobs (`<sentenceContextWindow>`, `<sentenceCentroidHistory>`, `<sentenceLambda>`, `<sentencePredictionScale>`, `<sentencePredictiveScale>`, `<sentenceContrastiveScale>`) shaped the legacy contrastive cosine machinery (recent-centroid attraction + older-centroid repulsion). They are not parsed; configs that still set them are tolerated silently.

Basic Model



The Basic Model of Cognition is a neural network of four spaces: real space, perceptual space, conceptual space, and symbolic space.

Model Entities:

Entity	Meaning
$x, \hat{x}, X$	Actual and predicted input (experiences), input space
$y, \hat{y}, Y$	Actual and predicted output (actions), output space
p, P	Percepts, perceptual space
c, C	Concepts, conceptual space
s, S	Symbols, symbolic space
W_e	Experiencing, experiential weights
W_a	Acting, action weights
W_k	Knowing, knowledge weights
W_t	Thinking, thought weights

Background

Based on the Basic Model of Cognition at cognitivesciencesociety.org/framework-cognitive-science; see also The Whole Part.

Inputs and Outputs

Inputs and outputs derive from a real space that is ineffable. Adaptable parameters initialize at zero (*tabula rosa*); learning is driven by random perturbations of the map relative to its territory at *temperature*.

Inputs scale to $[-1, 1]$ via global data min/max. The signed range represents presence and absence symmetrically. Paraphrasing Augustine, *negative is the privation of the positive*.

Action is a function of perceptual space and our expectation, filtered by conceptual projections. Performance errors are defined relative to action, desired effect, and exhaustion. Expectation errors are defined relative to perception and perception-as-reconstructed.

To accommodate a changing input context window, input is augmented with relative spatiotemporal encoding by extending its dimensionality. The attention mechanism treats what and where pathways separately.

Perceptual Space

Perceptual spaces are composed of perceptual prototype vectors, or *percepts*.¹ Percepts that do not correspond to objects are *hallucinations*.

Percepts form a Voronoi tessellation using prototype and/or exemplar vectors, related via Self Organizing Map connections, operating in parallel as prototype-theory categories.

Perceptual accuracy is the distance between inputs and prototype vectors — similar to cosine similarity, becoming a simple distance measure for poorly known percepts.

Attraction and aversion warp perceptual space: attraction is a pole at the percept's location, aversion is a zero. Perfectly recognized percepts lie on the unit hypersphere within conceptual space. Two percepts can be joined into one concept without loss.

Conceptual Space

Concepts create boundaries in conceptual space, grounded in the support vectors of perception. Concepts support logical negation; they may refer to negative entities. Conceptual spaces are capable of representing negative entities. Concepts that do not correspond to percepts are *false*.

Concepts collect percepts into fuzzy sets — linear separating hyperplanes partitioning conceptual space, or sets defined as linear combinations of percept activations.

Concepts are relative: meaningful only insofar as they partition a meaningful space. As concepts approach certainty, decision boundary slope becomes infinite. A single unknowing state is common to all knowledge vectors (activation slope zero).

Conceptual spaces are similarity spaces partitioned by decision boundaries imposing positivity and negativity on a hypersphere of knowing.

Within conceptual space, attraction/aversion to concepts expresses as hyperplane bias. Certainty is the slope of the activation function (tuned after other parameters stop changing).

A concept can be split into two percepts without loss.

Order, extension, and intension (2026-07-02)

Concepts occupy a **ramsified order**. An **order-0 concept** is the output of mixing a part and a whole at the corpus callosum — it is defined **extensionally**, by its perceptual support (the percepts that fall under it), and it is **subsymbolic**: a conceptual-space citizen with no relational definition yet. It is grounded directly. A **higher-order concept** is defined in terms of other symbols (its sparse constituent edges) and so acquires an **intensional** definition (its *sense* — the relations that fix it), while still retaining extensional grounding, inherited transitively through its constituents down to order-0's support and refinable by direct experience.

¹Perceptual space is known as “embedding space” in transformer literature.

Ascending the orders is therefore a **graded extensional** $\rightarrow$ **intensional shift**, **never a hard switch**, and the grounding becomes increasingly *mediated* (reached through the definitional graph rather than directly from percepts) — which is why superordinate concepts feel more theoretical and less perceptual (Rosch et al. 1976; the basic $\rightarrow$ superordinate trajectory). This is the reconciliation of the prototype-vs-definition tension: a concept is *both* its region and its definition, because intension and extension are carried by two different structures over the same object. Concretely — **intension is the sparse edge-graph; extension is the atom’s position** (Frege’s sense / reference; Carnap’s intension / extension, realized as graph-over-geometry).

Because every concept, at any order, has an extensional footprint (a grounded atom), specific and general objects share **one conceptual space**: a single dictionary holds each concept as a *point*, and the relational graph is structure laid *over* that space, not a second space. The single-space claim is not a new assumption — it follows from grounding (symbols are a subtype of perceptual objects; see Symbolic Space below). One refinement matters: a concept’s home in the space is its **point/atom**, *not* necessarily a convex region. Gärdenfors’s convexity holds for *natural* concepts and is the **order-0 / extensional limit**, cleanest at the grounded base; **discontiguous, high-order** objects still live in the same space as points, but their extension is **graph-defined** (the things reachable through their definition), not a convex region — which is exactly why the word $\leftrightarrow$ object bridge needs a MetaSymbol (“no convex set is specific enough,” see doc/Mereology.md). So: **one space of points plus a relational graph, with convexity as the basal/extensional reading**. The cognitive grounding for the dense-perceptual / sparse-symbolic split is in doc/Architecture.md “Cognitive grounding.”

The 2026-07-02 two-phase rework makes the extensional/intensional seam OPERATIONAL: the forward is a purely continuous PS/WS $\leftrightarrow$ CS pump (extension settling) followed by ONE late snap to the order-0 codebook block – quantization exactly at the bandwidth seam – and then the symbolic composition over the relational graph (intension). Order-0 concepts ARE codebook rows (extension = the atom’s position, no edges); order ≥ 1 concepts are edge-defined (intension = the sparse graph). Three relational primitives carry the graph (Alec 2026-07-02): the sec-4c ORDERED PAIR [*whole, part*] (*whole* $\Rightarrow$ *part* / if $\rightarrow$ then); the SINGLETON – a whole containing exactly one symbolic part, the unit-set $\{x\}$ Lewis resisted, here the constructive primitive; and the recursion VINE (ordered chains from nested pairs – order from nesting, since each row is set-like). Where the relation does NOT need order (the word/object meta), the read-out is typed intersection, not slot order. See doc/Architecture.md sec A and doc/plans/2026-07-02-two-phase-loops-sparse-relation.md; the un-ramified successor (iteration over one untyped square layer, with Kripke groundedness as the cycle diagnostic) is sketched in doc/plans/2026-07-02-iterated-symbolic-loop.md.

Symbolic Space

Symbolic Spaces are composed of percepts that reference concepts — a subtype of perceptual spaces mapping onto conceptual spaces.

Symbols have no spatial context: zero-dimensional points within perceptual space. In humans, perceptual space projects onto a 3D cerebellar subspace before symbolic encoding.

Symbols are independent, permanent, and integral (as references). The top-down serial activation differs from bottom-up perceptual activation — it casts shadows.

The formation of symbols increases the dimensionality of conceptual space (outer product); the formation of concepts from percepts decreases the dimensionality of perceptual space (inner product).

Symbols are erroneous if their associated concept is not true. Participation in multiple sets creates structural relations that replace positional information.²

Reasoning

Reasoning in symbolic spaces involves computing with parts and wholes — Venn-diagram computations enabling dynamic Bayesian-type statistics. Learned probabilities can be stored as unidirectional connection

²For example, $\{0,1\}$ and $\{0,1\}$, if orthogonal, pose no constraint on the location of the 1 in their 2x2 grid. However, $\{0,1,1\}$ and $\{0,1\}$ partition a 3x2 space and require two 1’s on the same row — a fact present in neither set alone.

weights between concepts, possibly with part-whole structure imposed by corresponding symbols.

Symmetric Perception: the Single Layer Linear Case

Perception is bidirectional: we see the world and the world sees us. The egocentric/cultural view treats perception as passive; in neural networks this manifests as forgetting the influence of outputs on inputs.

Symmetric Perception finds $f(x)$ minimizing MSE with respect to:

$$\begin{pmatrix} in \\ out \end{pmatrix} \cdot \begin{pmatrix} f(x) & 0 \\ 0 & f^{-1}(x) \end{pmatrix} = \begin{pmatrix} out \\ in \end{pmatrix}$$

In the linear case, $f(x)$ becomes a single invertible matrix A . When the mapping is nonlinear and an inverse may not exist:

$$\begin{pmatrix} in \\ out \end{pmatrix} \cdot \begin{pmatrix} f(x) & 0 \\ 0 & g(x) \end{pmatrix} = \begin{pmatrix} out \\ in \end{pmatrix}$$

This resembles $Ax = b$, except x is known and we want weights A mapping one space to another. The standard LLS solution $x = (A^\top A)^{-1} A^\top b$ becomes finding A such that $Ax = b$ and $A^{-1}b = x$.

Retrocausality suggests an additional objective: A should be optimal for both forward and inverse transforms. Approximately: A as a product of two matrices, each performing a perfect surjection: $A \cdot A^{-1} = I$.

$$\left(\frac{b}{2} * x^{-1}\right) \cdot \left(\frac{x}{2} * b^{-1}\right) = I$$

Best A requires iterative solution (NN learning, ping-pong style) or Gram-Schmidt orthogonalization with respect to A and A^{-1} .

Symmetric Perception: the Single Layer Nonlinear Case

If $f(x)$ and $g(x)$ share weight matrix W :

Forward computation and gradient:

$$\begin{aligned} \hat{y} &= f(g^{-1}(x)W) \\ e_y &= \frac{1}{2}(\hat{y} - y)^2 \\ \Delta W_y &= \eta(\hat{y} - y)\delta_f(W^\top \cdot g^{-1}(x))g^{-1}(x) \end{aligned}$$

Reverse computation and gradient:

$$\begin{aligned} \hat{x} &= g(W^\top f^{-1}(y)) \\ e_x &= \frac{1}{2}(\hat{x} - x)^2 \\ \Delta W_x &= \eta(\hat{x} - x)\delta_g(W \cdot f^{-1}(y))f^{-1}(y) \end{aligned}$$

Since W is shared, summing the partial derivatives gives:

$$\frac{\partial e_y}{\partial W} + \frac{\partial e_x}{\partial W} = \eta(\hat{y} - y)\delta_f(W^\top \cdot g^{-1}(x))g^{-1}(x) + \eta(\hat{x} - x)\delta_g(W \cdot f^{-1}(y))f^{-1}(y)$$

Setting gradient to zero, both errors going to zero when $x = i$ and $y = o$. But $xy - iy = ox - yx$ may hold — simultaneous gradient descent ping-pongs. A better cost function multiplies the two error partials — intersection of the spaces satisfying both:

$$\frac{\partial e_y}{\partial W} \frac{\partial e_x}{\partial W} = \eta(\hat{x} - x)(\hat{y} - y)\delta_f(W^\top \cdot g^{-1}(x))\delta_g(W \cdot f^{-1}(y))f^{-1}(y)g^{-1}(x)$$

Distance Metrics

$$d(x_1, x_2) = \frac{\|x_1\|^n \cdot \|x_2\|^n}{\|x_1 - x_2\|^{2n}}$$

Z-plane form:

$$d(x_1, x_2) = \|x_1 - x_2\| \cdot \frac{(x - z_1)}{(x - z_2)}$$

where z_1 is a zero and z_2 is a pole.

Exploration/Exploitation as a function of temperature

$$y = x \cdot \frac{(W + rt)}{\|W\|}$$

where r is a random vector and t is temperature.

Transfer Functions

Using rotation matrices with tunable Θ before and after a diagonal projection restricts the solution space to an eigendecomposition of the orthonormal basis (an SVD). For 2D:

$$y = \begin{bmatrix} \cos \theta_1 & -\sin \theta_1 \\ \sin \theta_1 & \cos \theta_1 \end{bmatrix} \begin{bmatrix} \Sigma_1 & 0 \\ 0 & \Sigma_2 \end{bmatrix} \begin{bmatrix} \cos \theta_2 & -\sin \theta_2 \\ \sin \theta_2 & \cos \theta_2 \end{bmatrix} x$$

With $f(x) = \sin(x)$ and $g(x) = \cos(x)$:

$$\begin{aligned} \hat{y} &= \sin(\cos^{-1}(x)W) \\ \hat{x} &= \cos(W^\top \sin^{-1}(y)) \\ \Delta W_y &= \eta(\hat{y} - y) \cos(W^\top \cdot \cos^{-1}(x)) \cos^{-1}(x) \\ \Delta W_x &= \eta(\hat{x} - x) \sin(W \cdot \sin^{-1}(y)) \sin^{-1}(y) \end{aligned}$$

The reverse perception reconstruction and the forward perception gradient share common terms. Taking the product:

$$\Delta W_y \Delta W_x = \eta(\hat{x} - x)(\hat{y} - y)\delta_f(W^\top \cdot g^{-1}(x))\delta_g(W \cdot f^{-1}(y))f^{-1}(y)g^{-1}(x)$$

Spaces

2026-06-21 terminology note (one noun per-space). This doc follows the percept / concept / symbol convention (doc/old/2026-06-21-terminology-percepts-concepts-symbols.md): a **percept** is a PartSpace/WholeSpace thing (dimensionally-embedded, extensional; the two subtypes are **part-percepts** = atoms, , and **whole-percepts** = properties/regions,); a **concept** is a ConceptualSpace relation tying one part-percept one whole-percept (the Concept codebook); a **symbol** is a SymbolSpace 0-D reference to a concept (the symbol codebook). The prose below reserves “symbol” for genuine SymbolSpace things; WholeSpace content is whole-percepts, and its compact [0,1] emission is the symbol that migrates to SymbolSpace. Code identifiers (e.g. `get_symbols`, `subspace.what`) are unchanged — those get a separate code pass.

2026-06-12 update (part/whole rename — the perceptual split). `PerceptualSpace` → `PartSpace` and `SymbolSpace` → `WholeSpace`; both now subclass a thin shared `PerceptualSpace(Space)` base (no parameters, no submodules — state-dict keys unchanged). Both views are *perceptual*: PartSpace synthesizes bottom-up over atoms (Sigma), WholeSpace analyses top-down over unity (Pi). XML config sections and grammar-file scopes are renamed to `<PartSpace>` / `<WholeSpace>` to match. The PS/SS shorthand in older notes reads part-side/whole-side. The freed name `SymbolSpace` was **reintroduced 2026-06-19** as the grammar/word space-role (formerly `WordSpace`); see doc/old/2026-06-19-handoff.md.

At the corpus callosum, objects are analysed and synthesized — by sending them back to `PerceptualSpace`: wholes get split and parts get chunked. In symbolic “mode”, the objects that get sent back are *symbols*.

Terminology (semiotics). There are **objects** and **references**. A reference is either a **sign** or a **symbol**. A *sign* is a quantized version of the referent (same space, snapped to a codebook row). A *symbol* is an unrelated version of the referent — an arbitrary code of much lower dimensionality (cf. the zero-banded symbol codes below: whole-percepts of symbols carry `.where = .when = 0`).

subspace.what: atoms vs properties. On `PartSpace` the `.what` rows are **atoms** (letters and words — the existing lexicon / BPE / MPHF front ends, gathered by `Codebook.lookup`). On `WholeSpace` the `.what` rows are **properties** — things that range over the whole input rather than naming one character (a sinusoid value is the worked example; “whitespace”/“words” are content-keyed properties). `materialize` hands the per-position selection (`subspace._index`, renamed from `_active` 2026-06-12) to the codebook so it produces the materialized object: atoms via `lookup`, or a per-position **region membership** via `Codebook.materialize_property` (`mode="property"`) — the region that has the property (`> 0`) and the region that doesn't (`<= 0`). The property path is additive and opt-in (`Codebook.property_basis`); the continuous sinusoid backend is wired, the content-keyed (whitespace/words) backend reuses the meronymic analyzer's span segmentation and is a documented seam. Properties are **low-frequency by construction**: there is too much detail to learn a codebook over the whole input unless the atoms are low-frequency, so the property basis frequency is tied to the input length (the k -th harmonic completes only $\sim k/2$ cycles over the whole field, mirroring `EndpointSumWhere`'s sub-half-period `div_term`). The learnable property row is the low-frequency atom; the basis enforces the coarseness.

2026-06-11 update (meronymy cutover — MeronymySpec/MeronymyPlan, Stage 9). With `<architecture><meronymy>on</meronymy>` (now the `model.xml` default) the meronymic slots bind the membership kernels: `PartSpace.sigma` → `SigmaLayer2` and `WholeSpace.pi` → `PiLayer2`, each through the K3-wire `MeronymicFoldAdapter` (at the boundary; the wire keeps carrying signed scalars; the fold computes on memberships, near-identity at init). Additions this page should be read with:

- **CS encoding:** stored reference rows are gauge-signed unit directions (semantic

embedding only); certainty is activation magnitude, polarity activation sign; gauge fixed at mint (`Spaces.gauge_orient`). Reference-half lookup is the pole quotient (`embed._pole_aligned_score`); token/form codebooks keep full-vector lookup.

- **The callosum:** percepts cross NAMELESS and FACTORED (`ConceptualSpace.factor_percept` — content selects the row, evidence sets a `[0, +1]`); parallel mode is the 2N mixing matrix; serial mode bypasses it (fusion migrates into the shift).
- **Two codebooks, one table:** the towers ARE the PS/SS codebooks (extent/intent; shared with IS recognition — recognition is tower lookup), linked by the word-keyed binding table (`References.ReferenceTable`) hosted on WholeSpace (search-then-mint gate; `adopt_stage0_evidence` stays ground-half memory, not naming). Symbols are atomic, zero-banded codes — whole-percepts of symbols carry `.where = .when = 0` (the Mind marker).

2026-06-09 update (analysis/synthesis orientation — supersedes the ownership notes below). The corrected orientation (`doc/old/2026-06-08-analysis-synthesis-dual-input.md`, rev. 2026-06-09; see `Philosophy.md`):

- **InputSpace emits the DUAL VIEW:** `forward(x) -> (percepts_in, concepts_in)` — the atom view (content `[B, N, 1]`) for the perceptual branch and the unity view (`[B, 1, N]`) for the symbolic branch.
- **PS = bottom-up SYNTHESIS:** owns ONE `SigmaLayer` (`self.sigma`, additive/union; the Pi/Sigma swap) and the `<synthesis>` front ends (`radix/bpe/byte/lexicon/mphf` — was `<chunking>`).
- **SS = top-down ANALYSIS:** owns ONE `PiLayer` (`self.pi`, multiplicative/intersection), the `<analysis>` division knob (`byte/word/analyse`), and the `<lexer>` intake knob (moved from `InputSpace`). Stage 0 consumes the unity view as symbolic evidence (per-part coarse means, snapped through the live SS codebook with the asymmetric STE forward leg).
- The meronymic analyzer (`bin/perceptual_analyzer.py`) is SS-side analysis machinery now; PS `<synthesis>analyse` was removed.

Sections below that predate this orientation are marked or should be read against it.

2026-06-02 update (subsymbolic analyzer). New `ObjectSubSpace` (`bin/Language.py`) — the PS-meronymic carrier analogue of `SymbolicSubSpace` (spans, parent/child links, route ids, marker-route replay fields). `WholeSpace` gains `insert_operations` (grammar operations in a dedicated operator codebook — `_operation_vectors` / `_operation_positions` — separate from the symbol codebook so the `symbol/idea/.where` namespace is untouched), `resolve_ps_terminal` / `null_sem` (PS-to-SS binding), and `operator_superposition`. The PS meronymic analyzer lives in `bin/perceptual_analyzer.py`. See `doc/old/2026-05-30-subsymbolic-analyzer-terminal-em`

Status (2026-05-27): updated for the substrate refactor. PS is a single-arg input processor (`self.pi + self.sigma`). CS is an STM container + grammatical CPU; no atomic forward fold. SS owns the unified word lexicon codebook with paired (orth, semantic) rows. Grammar dispatch lives on the signal router (`LanguageLayer`) at `SymbolicSubSpace.languageLayer`; the CKY Chart and STM shift-reduce parsers are retired. `LiftLayer` / `LowerLayer` are binary `GrammarLayer` subclasses with internal Sigma / Pi (no longer substrate-borrowing). `GrammarLayer` gains an optional butterfly cascade mode.


Overview

BasicModel is a pipeline of five **spaces** plus a grammar host (`SymbolSpace`), each performing a distinct representational transformation. Data flows forward from raw input to task output; the reverse pass reconstructs the original input from the symbolic representation. The legacy `SubwholeSpace` and `SyntacticSpace` classes have been retired — the subsymbolic role is filled by `PartSpace` itself, and the grammar runs from `SymbolicSubSpace.languageLayer` (the signal router; subsumed the retired `Chart`).

Forward: `InputSpace -> PartSpace -> ConceptualSpace -> WholeSpace -> OutputSpace`

Reverse: OutputSpace -> WholeSpace -> ConceptualSpace -> PartSpace -> InputSpace

The pre-2026-05-27 “two feedback loops” ($S \rightarrow C$ per-stage, $C \rightarrow P$ cross-forward) are retired. The recurrent character lives in (a) STM accumulation across words in SERIAL / GRAMMATICAL mode, and (b) the T-pass PARALLEL refinement loop driven by `<subsymbolicOrder>` via `PS.forward(CS)` iterations.



Base Class: Space

All spaces inherit from `Space`, which manages:

- **Shape management.** `inputShape` / `outputShape` as [nObjects, nDim]. Subclasses read dimensions from `TheObjectEncoding`.
- **Codebook / VQ quantization.** When `nVectors` > `nActive`, a codebook holds candidate vectors; top-k selection gives the bottleneck.
- **Reshape flag.** When in/out object counts differ, the [B, nObj, nDim] tensor is flattened before the next space and restored on the way back.
- **Attention.** The legacy boolean `hasAttention` is deprecated and inert (kept only as a backward-compat alias); use the `<attention>` element instead (`off` | `primer` | `second-order` | `low-rank`).
- **set_sigma propagation.** Ergodic-mode noise level cascades from the top-level model down through every child layer.

Reset cascade — hard vs. soft

Every space exposes `Reset(batch=None, hard=True)`. The signature is required (legacy zero-arg fallback removed).

Call form	Scope	Use
<code>space.Reset()</code>	All rows	Whole-state wipe (epoch boundary)
<code>space.Reset(batch=b, hard=True)</code>	Row <code>b</code> only	Document boundary
<code>wordSpace.soft_reset(batch=b)</code>	Row <code>b</code> sentence-scoped state	Grammar <code><start></code> reduction

Hard-reset clears: parse stack, `_last_svo`, `_stm_fired`, codebook commit accumulator, discourse history, `serial_cache`, `_ar_embedded`, `_end_of_stream` for the affected rows.

Soft-reset clears per-sentence working buffers (parse stack rows, `_last_svo[b]`, category and reconstruction stacks) and re-arms `_stm_fired[b]`. Does **not** touch discourse history (`InterSentenceLayer` ring buffer) or codebook EMA — those are document-scoped.

Reset is dispatched from `runEpoch`, never from inside `runBatch` (the pure compute brick — see Architecture.md).

Sigma / Pi ownership (Pi/Sigma swap, rev. 2026-06-09)

History: the 2026-05-13 rebalance described “each space owns one operator”; the 2026-05-27 substrate refactor gave PS both folds; Stage 10 made PS pi-only. The **corrected analysis/synthesis orientation (2026-06-09) swaps the folds to their proper sides**: Sigma (sum/union) is synthesis and belongs to the bottom-up PartSpace; Pi (product/intersection) is analysis and belongs to the top-down WholeSpace. CS remains an STM bookkeeper with **no atomic forward operator**; Lift / Lower stay **binary GrammarLayer subclasses** with internal Sigma / Pi (no substrate-borrowing).

Space	Owns
PartSpace	one <code>self.sigma</code> (SigmaLayer — the synthesis fold), the <code><synthesis></code> front ends, MPHF + index tables
ConceptualSpace	STM (ShortTermMemory, depth ~8) + (when sparse-active) per-order sparse concept-weight tables (<code>_</code>
WholeSpace	one <code>self.pi</code> (PiLayer — the analysis fold), the <code><analysis></code> + <code><lexer></code> knobs, the unified word lexicon

Composition (per-mode):

- **SERIAL / GRAMMATICAL** (`<symbolicOrder>1</symbolicOrder>`): one iteration per word.

```

PS_t = PS.forward(IS_t)           # MPHF + pi + sigma on per-word surface
CS_t = CS.forward(PS_t)           # STM shift + push
router.dispatch_at_C(STM)         # read-only grammar ops on STM contents
router.dispatch_at_S(STM)         # codebook-write-required ops via SS

```

- **PARALLEL** (`<symbolicOrder>0</symbolicOrder>`): T iterations of PS over CS.

```

PS_0 = PS.forward(IS); CS_0 = PS_0
for t in 1..T = <subsymbolicOrder>:
    PS_t = PS.forward(CS_{t-1})    # refinement pass
    CS_t = CS.forward(PS_t)        # STM[t] = PS_t (parallel write; no shift)
    router.dispatch_at_C(STM)      # grammar ops after STM population

```

The legacy formula `C = sigma_percept(pi_input(IS) + pi_concept(C_prev))` is retired entirely.

Lift / Lower in the new ownership

`LiftLayer` and `LowerLayer` are no longer “rule-id annotators over shared substrate.” They are first-class binary `GrammarLayer` subclasses (Stage 4 of the substrate refactor):

- `LiftLayer(GrammarLayer)`: `arity=2`, `rule_name="lift"`, `space_role='C'`. Owns an internal `self._sigma: SigmaLayer` for the pairwise additive (sigma-style) math. `forward(left, right)` delegates to `_sigma.compose`.
- `LowerLayer(GrammarLayer)`: `arity=2`, `rule_name="lower"`, `space_role='C'`. Owns an internal `self._pi: PiLayer` for the pairwise multiplicative (pi-style) log-domain math. `forward(left, right)` delegates to `_pi.compose`.

Both reverse cleanly via their internal layer's reverse. Both gain butterfly mode for free via `GrammarLayer` base inheritance (Stage 5).

The signal router dispatches them as binary reduce ops at the CS, weighted by `Grammar.rule_probability` (the per-position copy/reduce score head).

Butterfly mode on GrammarLayer (Stage 5)

`GrammarLayer(butterfly=True, N=N)` allocates a packed `nn.Parameter` of shape `[n_levels, N // 2, 2D, 2D]` (`n_levels = log2(N)`), a bit-reversal permutations buffer, and LDU-parameterized invertibility per node. Identity init: cascade is identity at construction. Each subclass implements `_butterfly_pair_op(x_pair, W_node)` with its per-pair math:

- `PiLayer._butterfly_pair_op`: $\text{atanh} \rightarrow \text{einsum} \rightarrow \text{tanh}$ (multiplicative log-domain pairwise).
- `SigmaLayer._butterfly_pair_op`: $\text{atanh} \rightarrow \text{einsum} \rightarrow \text{tanh}$ (additive pairwise — same surface, different gradient regime).
- `LiftLayer` / `LowerLayer`: delegate to their internal `sigma` / `pi`.
- `IntersectionLayer` / `UnionLayer`: `min` / `max` kernels.
- `ConjunctionLayer` / `DisjunctionLayer`: hard-coded monotonic `min` / `max`.

Parameter savings: $O(N \cdot \log N \cdot D^2)$ cascade vs $O(N^2 \cdot D^2)$ for a single big matrix. Wired into the space folds (`PartSpace.sigma` / `WholeSpace.pi` post the `Pi/Sigma` swap) by the global `<sigmaPi>` mode (default butterfly). Closes the XOR convergence target (`test_mm_xor.py`).

Normalization and Ranges

Space	Data Contract
<code>InputSpace</code>	Data scaled -1..1 for scalars or vector norms
<code>PartSpace</code>	Modal/demuxed (what/where/when encoding). Signed unit-magnitude scalars or vectors. No negation
<code>ConceptualSpace</code>	Combined/muxed (event encoding). Signed unit-magnitude (tanh-bounded). Event norm on <code>subspace</code> .
<code>WholeSpace</code>	Whole-percepts (<code>[0,1]</code> presence). The compact emission is the symbol that references a concept; one sy
<code>OutputSpace</code>	Rescaled from activation range to original data range

`SyntacticSpace` is retired (see Section “`SyntacticSpace` — retired” below). Grammar / chart machinery lives on `SymbolSpace` and is attached to `WholeSpace` (the canonical grammar host).

Data scaling. `Data` computes global `input_min/input_max` and `output_min/output_max` at load time. `InputSpace` uses `Data.normalize(x, "input")` to scale to `[-1, 1]`; `OutputSpace` uses `Data.denormalize(x, "output")` to restore the original output range.

Whole-percept presence. Whole-percepts live in `[0, 1]`; since conceptual activations range `[-1, 1]`, the mapping is `presence = (activation + 1) / 2`. `SubSpace.get_symbols()` / `set_symbols()` (code identifiers unchanged) perform the conversion.

Demuxed mode. When `InputSpace.demuxed=true`, what/where/when components are stored independently in the `SubSpace` rather than concatenated. `ModalSpace` routes each component through independent `PartSpaces`; downstream spaces see an identical muxed tensor via `materialize()`.

Codebook Similarity Metric

Codebook wraps `VectorQuantize`. Similarity metric per space:

Space	Codebook geometry
PartSpace	$[-1, +1]^d$ hypercube
WholeSpace	$[-1, +1]^d$ hypercube
ConceptualSpace	<code>similarity_codebook</code> rows (<code>use_dot_product</code>) — the concept dictionary / SBOW-situating metric, N

Euclidean (Perceptual — Part / Whole)

These codebooks store *what something looks like* — $0.5 \cdot v$ carries half as much “of feature v ” as $1.0 \cdot v$, so the right notion is coordinate-wise distance. Retrieval expands $\|x - c_i\|^2$:

$$\|x - c_i\|^2 = \|x\|^2 + \|c_i\|^2 - 2(x \cdot c_i)$$

$\|x\|^2$ is constant across i and drops from argmin:

$$\arg \min_i \|x - c_i\|^2 = \arg \max_i (x \cdot c_i - \frac{1}{2} \|c_i\|^2)$$

`VectorQuantize` keeps $\|c_i\|^2$ in `_b_norms_sq`:

```
indices = (flat @ codebook.T - 0.5 * b_norms_sq).argmax(dim=-1)
```

One matmul + one broadcast subtract + one argmax. Skips the `sqrt`, the per-row $\|x\|^2$ add, and the `cdist` autograd plumbing.

Dot product (Conceptual)

Note. This dot-product metric is the `similarity_codebook`’s retrieval metric (used by the substitutability / SBOW *situating* signal), NOT the forward concept-production path. When the ramsified sparse transform is active, a concept code is produced by the sparse encode/decode (`a_k = W_k @ source_k`, then `a_k · softplus(atom)`) — there is no `argmax_i (x · c_i)` concept retrieval on the forward path, and the atoms are softplus-positive rather than maintained unit-norm by EMA. See **ConceptualSpace** → **Sparse-coding concept transform**.

ConceptualSpace concepts are *named directions* in belief space. $x \cdot c_i$ gives the *signed strength of belief that x affirms concept i* :

- +1 fully affirms; 0 orthogonal; −1 fully denies

Two consequences:

1. **Codebook must be unit L2-norm.** EMA renormalizes after each update.
2. **Input must NOT be normalized.** The magnitude *is* the certainty signal. Cosine similarity would divide it out.

For *ranking*, $x \cdot c_i$ and $\cos(x, c_i) = (x \cdot c_i) / \|x\|$ are monotone-equivalent (positive constant cancels). Omitting input normalization preserves certainty and costs less:

```
# codebook is unit L2-norm (maintained by EMA)
indices = (flat @ codebook.T).argmax(dim=-1)
```

Configuring the metric

`use_dot_product` is a class attribute on `Codebook` (default `False`). Set it on a `Space` subclass to opt in — `ConceptualSpace` does this. The underlying `VectorQuantize.use_cosine_sim` flag is historical; after the April 2026 perf pass, input-side normalization is gone, so the effective meaning is “codebook unit-norm; rank by dot product”.

Ramsification table (per-code fold record)

The `Pi` / `Sigma` folds carry a reference onto sortable mereological space but do **not** preserve their own *ramsification* — the record of how the code was produced — so a folded code cannot, on its own, be reconstituted. `Codebook.ramsification` is the small adjacent table that fixes this: a `[V, max_order]` `uint8` sidecar, index-aligned with the codebook rows (both perceptual codebooks — `PartSpace.subspace.what` and `WholeSpace.subspace.what`), recording for each code which fold it was routed through at each subsymbolic pass — `FOLD_NEITHER` / `FOLD_SIGMA` / `FOLD_PI`. `invert_ramsified(code, row, sigma, pi)` walks that sequence in reverse pass order, applying `sigma.reverse` / `pi.reverse` per recorded fold, landing back at the codebook row that produced the code.

The table is an **opt-in additive sidecar** (`enable_ramsification`; a plain attr like `part_parents` / `category_ids`, not a `Parameter` or `buffer`) — it adds no `state_dict` keys and cannot move a pinned basin; it resizes with the codebook (`grow_to`). Live per-pass stamping in the subsymbolic pump loop is the deliberate cutover seam (call `record_fold` where `PartSpace.sigma` / `WholeSpace.pi` fire).

Word abstraction order. A code’s `abstraction_order` is its fold count (non-`NEITHER` passes), and words are subsymbolic at several abstraction levels: a **proper noun** (prototype / token) matches raw at **order 0**; a **regular noun** (type) at **order 1**; a **count noun** (concrete only under a determiner) at **order 2**; higher orders are more abstract. Words need not be nouns, but all benefit from an abstract / discontinuous spatial representation. This connects to the ramsified order hierarchy in `Language.Taxonomy` and the order-typed STM plan.

Codebook Uniqueness Contract

Every codebook entry is identified by its **row index** and must carry **distinct .what** (`WhatEncoding`) **content**; the old `.where`-keyed uniqueness scheme is retired:

- **.where** — **positional / spatial-extent key (no longer a codebook row key)**. The cross-codebook **.where slice registry was RETIRED** (modality re-architecture, 2026-06-04; `WhereEncoding`, `Spaces.py`). `allocate_codebook_slice` / `global_max_val` / `reset_codebook_registry` were removed: there is no shared where-space to allocate disjoint slices in. Codebook identity is now the **row index** (the `_index` selection), `.where` keeps only its positional / spatial-extent role, and `CS→WS` reverse decode is **content-match** (nearest row). Cross-codebook taxonomy is row/position-keyed via `WholeSpace`’s explicit dicts (`category_ids`, `part_parents`).
- **.what** — **distinct prototype content**. Identical `.what` collapses to the same parthood identity (`equal(A, A) = 1`) — a redundant pair the network can’t distinguish.

Current enforcement:

Source	Mechanism
WholeSpace codebook	<code>ImpenetrableLayer</code> overlap penalty + variance floor; five-relations classifier pushes pairs toward
ConceptualSpace codebook	<code>ImpenetrableLayer</code> available; not yet wired by default
PartSpace Lexicon	Cosine-margin pole/antipode SBOW training
InputSpace vocabulary	Shares PartSpace’s Lexicon

`.where` is now a positional / spatial-extent carrier (the slice registry is retired — see above); codebook

identity is the **row index**. `.what` uniqueness is **learned** (encouraged by `ImpenetrableLayer` + antipodal quotient) and, together with the distinct row indices, keeps the parthood lattice well-formed.

Lexicon (Projective Unit Ball)

The **Lexicon** (`bin/Layers.py`) backs PartSpace word embeddings and WholeSpace whole-percept prototypes. Each row is a vector w_i in the **projective unit ball** — the closed ball $B^D = \{x : \|x\|_2 \leq 1\}$ with the **negation identification** $w \sim -w$ realizing real projective space $\mathbb{RP}^D$.

Terminology pin — three notions sometimes conflated:

- **Pode** of (a, b) : midpoint $(a + b)/2$; SBOW positive-pair attractor.
- **Wrapped pode**: midpoint via the $\pm$ -quotient, $(a - b)/2$; the midpoint through *negation* of b .
- **Antipode** of a single point p : furthest point. On the flat torus unique (`wrap(p + 1)`); on $\mathbb{RP}^D$ **not unique** — the maximum-distance set is the orthogonal hyperplane.

So $-w$ is the **negation** of w , *not* the antipode.

Distance and lookup

For $a, b \in B^D$ the projective squared distance is

$$d_{\mathbb{RP}}^2(a, b) = \min(\|a - b\|_2^2, \|a + b\|_2^2) = \|a\|_2^2 + \|b\|_2^2 - 2|\langle a, b \rangle|.$$

With `pode(a, b) = (a + b)/2` and `wpode(a, b) = (a - b)/2`, $d_{\mathbb{RP}}(a, b) = 2 \cdot \min(\|a - \text{pode}\|, \|a - \text{wpode}\|)$. The lookup picks whichever rep of b (b or $-b$) is closer to a .

Sorting by smallest $d_{\mathbb{RP}}^2$ = sorting by largest $\text{score}(x, w_i) = |\langle x, w_i \rangle| - \frac{1}{2}\|w_i\|_2^2$.

Implementation: `cache W_norm2 = W.square().sum(-1)` once per optimizer step; top-k is `(x @ W.T).abs() - 0.5 * W_norm2` followed by `torch.topk` — dense matmul + abs + broadcast subtract. No $V \cdot D$ outer-product.

The Lexicon API:

```
lexicon = Lexicon(V, D)
lexicon.project_unit_ball_()          # after optimizer.step()
W_index, W_norm2 = lexicon.lookup_index()

# Projective (RP^D) --- antipode-aware, default.
idx, dist_sq, scores = Lexicon.topk_rp(x, W_index, W_norm2, k=32)

# Plain L2 --- for sites where w and -w are distinct.
idx, dist_sq, scores = Lexicon.topk_l2(x, W_index, W_norm2, k=32)

# Pairwise primitives:
Lexicon.rp_distance_sq(a, b)
Lexicon.rp_similarity(a, b)
Lexicon.rp_pode(a, b)
Lexicon.rp_wrapped_pode(a, b)
Lexicon.rp_closer_rep(a, b)          # sign(<a, b>) * b
```

For $V \gtrsim 10^5$, use `topk_rp_chunked` to bound peak score-tensor size.

SBOW training: podes (attractor) and antipodes (repulsion target)

- **Pode (attractor).** Positive-pair updates pull a and b toward $\text{pode}(a, b)$; the gradient picks the closer of b and $-b$ for shorter-arc attraction.
- **Antipode (balancing repulsion target).** Negative-pair updates push the row toward the furthest point. On $\mathbb{RP}^D$ this is a $(D - 1)$ -sphere, so SBOW samples a representative orthogonal direction.

Negative-sampling gradient has two regimes by $\text{sign}\langle a, b \rangle$: positive case is standard contrastive repulsion along $(a - b)$; negative case pushes a away from $-b$ along $(a + b)$.

After every optimizer step the trainer calls `lexicon.normalize()`, clipping $\|w_i\| \leq 1$. `W_norm2` should be refreshed when weights change.

Torus primitives (`Lexicon.wrap`, `Lexicon.delta`, etc.) and the `torus=True` constructor flag remain as **legacy** static methods (the earlier Lexicon used the flat torus $T^D = [-1, 1)^D$ with wrapped MSE). New code must use the `rp_*` primitives.

InputSpace

Role. Receives the raw source buffer and lifts it into the model’s internal working dimensionality.

Text mode forward. Delegates tokenization to `Lex`, producing a span table of `(start, end, type)`. Each span $\rightarrow$ a vector with two components:

- `nWhat` dims — token content, encoded via `Basis` / `Codebook` (the word embedding lookup).
- `nWhere` dims — positional information, materialized via `WhereEncoding` (see the 2026-05-28 where-keyed-taxonomy plan). As of the 2026-06-16 redesign `.where` is an **endpoint-sum BRACKET** over a span `[start, end]` (the invertible `EndpointSumWhere` form adopted into the muxed tail): the key is $0.5 \cdot [\sin(s \cdot) + \sin(e \cdot), \cos(s \cdot) + \cos(e \cdot)]$, so the **angle decodes the span center** and the **magnitude the span extent**. An `INSTANT (start==end)` collapses to the legacy single-quadrature point `[sin(p ·), cos(p ·)]` (byte-identical to the pre-bracket stamp). `decode` returns the center (the canonical positional identifier across IS / PS / SS taxonomies, recoverable via `atan2(sin, cos) /`); `decode_span` returns `(start, end)` for the mereology **contiguity** / **containment** test (instants snap to zero extent). `.when` uses the **same bracket** over event `TIME` (center = event time, magnitude = duration); tense is the interval-vs-now relation, not a magnitude (see `WhenRangeEncoding`).

Result: `[nActive, nWhat + nWhere]` tensor.

Text mode reverse. Inverts the span encoding: each vector $\rightarrow$ nearest codebook entry, then spans $\rightarrow$ characters via the stored offset table.

Numeric mode. Tensor data passes through unchanged; `LiftingLayer` projects native input dim (e.g. 784 for MNIST) to `nDim`. Non-embedding inputs are scaled to `[-1, 1]` via the global data min/max.

Key parameters.

Parameter	Description
<code>nActive</code>	Sequence length
<code>nDim</code>	Output dim per vector
<code>lexer</code>	(Moved to <code><WholeSpace></code> , Phase 4b — lexing is analytic cutting. <code>InputSpace</code> executes the intake; the knob lives in <code>WholeSpace</code> .)
<code>codebook</code>	Whether input values are discrete
<code>demuxed</code>	Store what/where/when independently

2026-05-28: per-Space `<nWhere>` / `<nWhen>` XML knobs are retired. `model.xml` carries the per-Space defaults (`<nWhere>2</nWhere>` on `InputSpace` / `PartSpace` / `WholeSpace`, 0 elsewhere). Per-file configs no longer declare them. See `doc/old/2026-05-28-where-keyed-taxonomy.md`.

Invertibility. Always non-invertible; reverse is a separate reconstruction using the span table.

Document streaming and `valid_mask`

Documents longer than `nOutput` bytes are not truncated. `TheData` maintains a per-row cursor (`doc_idx[b]`, `offset[b]`) and `next_tick()` returns (`input`, `output`, `hard_eos`) where `input` is a `[B, nOutput]` slab containing the next $\leq nOutput$ bytes from each row's current document. `hard_eos[b]` is a host-side bool set when row `b`'s cursor exhausts the current document. A short fill at document end NULL-pads the slab tail; `valid_mask: [B, K]` bool flips False for padded positions, and state-mutation propagation skips them.

Cursor universal — trial mode for non-AR data. `next_tick()` is the single dispatch for both AR text byte (rolling cursor) and non-AR data (numeric). In trial mode (`slab_bytes` not set), each tick yields one batch of trials with `hard_eos = [True] * B`. The `runEpoch` outer loop drives `ds.next_tick()` directly for both modes; the `DataLoader` exists only so existing tests can grab `loader.dataset`.

`_end_of_stream` is a host-side `list[bool]` diagnostic only; the canonical hard-reset signal is the cursor's `hard_eos`.

AR cursor unfold retirement (2026-05-13)

The legacy AR-training path padded + unfolded the embedded sentence into `[B, K, N, D]` cursor windows so the body could see a `[B*K, N, D]` parallel view of every prefix. At `bs=128, K=128, N=1024, D=10`, the unfolded tensor alone was ~320 MB.

The unfold was retired for AR training on 2026-05-13, replaced with a serial K-cursor loop (`_forward_per_stage_no_unfold`) that walked the same prefixes with a `[B, N, D]` tensor and a per-cursor causal mask.

Within-sentence AR retirement (2026-05-14)

The serial K-cursor loop itself was retired one day later: the benchmark showed `_forward_per_stage_no_unfold` running at ~18 sent/sec (the K body+head calls dominate) vs the single-shot IR fast-path's ~61 sent/sec, and the real AR objective in this architecture is **next-sentence** prediction (the discourse layer) — not next-token within a sentence.

Within-sentence training is now IR-only. `InputSpace.forward` emits `[B, N, D]` (left-aligned, right-padded to N) and `_forward_per_stage` runs a single masked-LM pass:

1. **Stem:** `InputSpace.forward + PartSpace.forward` $\rightarrow [B, N, D]$.
2. **Mask:** `create_ir_mask` replaces a `mask_rate` fraction of WHAT positions with `NULL_PERCEPT`; pre-mask event stored on `_ir_pre_mask_input` as the loss target.
3. **Body:** T stages on B rows (no per-cursor walk, no causal mask).
4. **Head:** `outputSpace` $\rightarrow [B, N, predDim]$. The head is a side channel — IR loss is computed at the subsymbolic (PS), not at the head.

`runBatch` reads `_ir_mask_positions` and `_ir_pre_mask_input` and computes `MSE(perceptualSpace.subspace at masked positions, _ir_pre_mask_input at masked positions)`. The `<reconstruct>` element (and its `reconstructEnum`) is RETIRED (A1, 2026-06-09): the `ConceptualCombine` now unconditionally integrates all three streams (PS + SS + CS), so reconstruction is unconditionally from concepts — the former `concepts|symbols|both` selection (target derived by lifting `_ir_pre_mask_input` through `sigma_percept`; see Plan Section “Reconstruction-loss target shape” Option B) is no longer a knob.

`<maskedPrediction>` is retired; `<reconstruct>` is retired (the `output` mode was the only path that fired the reverse pipeline); `<reverseScale>` is renamed to `<reconstructionScale>` (the legacy name remains parseable with a one-shot deprecation warning).

Sentence-level AR moves to `InterSentenceLayer` — see `doc/Architecture.md` Section “Sentence-level AR (`InterSentenceLayer`)” for the ARMA(p, q) design.

PartSpace

Role. Single-arg input processor — the bottom-up SYNTHESIS branch (Pi/Sigma swap, rev. 2026-06-09). Applies `self.sigma` (the additive/ union fold) to its argument (the atom-view stem after the synthesis front end embeds). Owns the surface-keyed Lexicon (`self.vocabulary`) and the MPHF + index table for per-word surface $\rightarrow$ row lookup.

Owned state:

- `self.sigma`: a single `SigmaLayer` (`percept_dim  $\rightarrow$  percept_dim`, where `percept_dim` is the EMBEDDED percept width — `_fold_width`; a widening PS sizes the fold at `nOutputDim`, not the raw `nInputDim`). Inherits from `GrammarLayer`; accepts `butterfly=True`, `N=N` for cross-position cascade mode. (The `PiLayer` PS used to own moved to `WholeSpace` — `Pi` is analysis.)
- `self.vocabulary`: the Lexicon (`Embedding`), keyed by MPHF over surface bytes. Per-word vectors are `nDim`-wide (CS-space-dim per the flat-slab invariant).
- `self._mphf_gpu_layer`: MPHF infrastructure for fast surface lookup.
- `self.chunk_layer`: BPE machinery (the `ChunkLayer` from `bin/Layers.py`).
- `self.radix_layer`: when `<synthesis>radix</synthesis>`, the input lookup routes through `RadixLayer` (radix trie + inverse table + learned codebook + byte fallback). `RadixLayer` is a first-class `Layer` subclass in `bin/Layers.py` (formerly the standalone `PerceptStore`). `PartSpace.reverse` invokes `RadixLayer.reverse` for the structural decode (chunk-id $\rightarrow$ bytes $\rightarrow$ slot). Promotion knobs default to `threshold=4`, `min_length=2`.

The legacy `pi_input` / `pi_concept` ModuleLists are retired, as is the `sigma_percept`-style additive fold on CS.

Forward (`PS.forward(x_subspace)`):

```
def forward(self, x_subspace):
    # synthesis front end embeds (lexicon/bpe/byte/radix/mphf), then:
    x = self.forwardBegin(x_subspace, returnVectors=True)
    return self.sigma.forward(x)    # the union fold; internal tanh
```

`x_subspace` is the atom-view stem (PS runs ONCE at stage 0 — the single-pass subsymbolic decision; the per-stage recurrence advances through the `ConceptualCombine`, not repeated PS calls).

Math (the `sigma` fold — PS's synthesis operator):

```
sigma(x) = tanh(W_sigma @ atanh(x) + b_sigma)    # additive/union, log-domain
```

(The multiplicative `pi` math — `pi(x) = tanh(W_pi @ atanh-domain + b)` in the `(1+x)/(1-x)` log embedding — now lives on **WholeSpace** as the analysis fold; see the orientation banner.)

Reverse. `PS.reverse` applies `self.sigma.reverse` on the text path (LDU inverse via `InvertibleLinearLayer`); structural recovery goes through `object_basis.reverse` and (in radix mode) the `RadixLayer.reverse` chunk-id $\rightarrow$ bytes decode.

Butterfly mode (Stage 5): when `<PartSpace><butterfly>true</butterfly>`, the fold is constructed with `butterfly=True`. The cascade length `N` is auto-derived from the space shape (`nInput * nInputDim`, internally padded to the next power of two); there is no `<butterflyN>` knob (it was retired 2026-06-05). `<butterfly>` itself is a deprecated alias for the architecture-level `<sigmaPi>` (new configs should use that). Internal storage becomes a packed `nn.Parameter[n_levels, N//2, 2D, 2D]` cascade with bit-reversal permutations. Closes the XOR convergence target. `PartSpace` is subsymbolic and takes no `<codebook>` element (it was retired; PS is fixed to `none`); butterfly weight gradient flow therefore flows through the continuous `.event` passthrough on PS. STE-through-snap (for spaces that do quantize) is a known follow-up.

Range. Vectors live in $[-1, 1]^d$ (tanh-bounded). No negation operator — percepts represent feature magnitudes with sign indicating direction.

ConceptualSpace

Role. Forms concepts from perceptual/symbolic *sources*. When the ramsified sparse transform is active (`symbolicOrder > 0` in parallel mode, `_sparse_active()`), a concept is a high-dimensional atom in **ConceptDim** (stored in the CS concept dictionary, the `similarity_codebook`) whose signed activation is produced by a **per-order sparse weight matrix** over its sources — `concept_activation = W_order @ source_activation (cs_sparse_encode, torch.sparse.mm)` — then scaled onto the atom (`cs_decode`). **PerceptDim** and **ConceptDim** are **decoupled** (the matrix lives in index/activation space); a concept is NOT an additive linear map over percept vectors and there are no “conceptual hyperplanes partitioning perceptual space.” With the transform off the forward is STM bookkeeping only (byte-identical). See **Sparse-coding concept transform** below.

Geometry. A concept code is `signed_activation * softplus(atom)`: the dictionary atom is a **strictly-positive** **ConceptDim** feature signature (`F.softplus`), and the **sign** / **magnitude** live in the signed scalar activation — magnitude = certainty, sign = present vs anti-present, low magnitude = chosen radially but uncertain. The sign comes from the signed sparse **weights** (a negative weight = a feature’s presence anti-correlates with the concept), not from a stored signed unit direction. (The legacy “named unit-norm direction + `argmax` retrieval” view is retired on the forward path.)

Owned layer (2026-05-13 rebalance → 2026-05-29 clean-stack → RETIRED). **ConceptualSpace** owns **no parameterised fold layer**. The historical `self.sigma_in / self.sigma_cs` **SigmaLayers** below were **RETIRED** — they are no longer constructed, and **CS.reverse** no longer applies them (see the **Reverse** note). The table records the pre-retirement Stage-10 design:

Layer (RETIRED)	Direction	Math	Notes
<code>self.sigma_in</code>	incoming-contribution fold	per-stage SigmaLayer (Ramsified across stages)	Stage
<code>self.sigma_cs</code>	residual-CS iteration kernel for stages $k > 0$	per-stage SigmaLayer	Same

Clean-stack STM (2026-05-29 experiment). The Stage-10 additive composition

```
folded = sigma_in(combined) + sigma_cs(prev)
```

is replaced with per-stage space-role attribution:

```
stage 0      folded = primary      (PS event from subspace.materialize())
stage k > 0  folded = sym          (SS event from word_subspace.materialize())
```

No additive mixing across space-roles; no residual lift; trivially invertible (read-back, no inverse-Sigma needed). The `STM_k = STM_{k-1} + SS_k` carry-forward variant was tested and reverted — the pure clean-stack form is the landing point. See `doc/old/2026-05-29-clean-stack-stm-basis-arg-radixlayer.md`.

Because `sigma_in / sigma_cs` were dead-weight on the forward path (no gradient — they never fired) while **CS.reverse** applied `sigma_in.reverse` unconditionally, the round-trip carried an **UNMATCHED** inverse fold (the source of garbage `XOR_exact` recon tokens). That forward / reverse semantic mismatch was **RESOLVED** by retiring the layers entirely: with the sparse transform **OFF**, CS is a pure bookkeeping carrier (forward push / reverse read-back) with no fold to invert, and the symbolic generalization operator moved to **WholeSpace** (inverted upstream of **CS.reverse** on the reconstruction path, in `BasicModel._reverse_body`). Convergence on `MM_xor` continues via the **PiLayer** butterfly cascade. (When the ramsified sparse transform is **ON**, CS **DOES** own a parameterised forward transform — the per-order sparse weight matrices and the concept dictionary — see **Sparse-coding concept transform** below.)

Reverse. **CS.reverse** is a thin pass-through (no fold layer to invert). The reverse chain operates on the terminal STM contents; per the master plan, no per-stage caches. The sparse-coding reconstruction is referential — the per-order weight tables **ARE** the concept’s decomposition — rather than an inverse fold.

Sparse-coding concept transform (ramsified per-order, torch.sparse). When `_sparse_active()` (i.e. `_symbolic_order > 0` and `parallel/serial=false`) and the per-order weight tables are populated,

`forward` interposes `folded = self._sparse_concept_forward(folded, subspace, word_subspace)` before the STM write. A concept is a `ConceptDim` atom; it is linked to its sources by a sparse WEIGHT MATRIX — one `torch.sparse_csr` table per RAMSIFIED ORDER, order k 's sources being `[PS | WS | SS_0 ... SS_{k-1}]` (symbols are concepts of the previous order). The forward is an encoder + dictionary decoder:

```
concept_activation = W_order @ source_activation      # cs_sparse_encode (torch.sparse.mm)
concept_code[c]    = concept_activation[c] * softplus(atom[c])  # cs_decode
```

PS/WS source activations are non-negative **presences** (`source_code_activation(nonneg=True)`); SS sources are the lower orders' **signed** concept activations threaded up (the symbolic space is 1-D). The weights are **signed** and learnable (`_csw_vals`, grown host-side in `add_concept_weight`, exposed via `getParameters()` and registered into the optimizer via `_maybe_rebuild_optimizer_for_csw`). Capacity is **dyadic** by order ($N/2, N/4, \dots$, stacked; `order_capacities / order_slice`). Population is at mint, keyed by abstraction order (`_populate_concept_weights, _concept_source_order`). Forward-connected: the gradient reaches the sparse weights, the source activations, and the dictionary. Off-path (`symbolicOrder = 0`) $\rightarrow$ no weights $\rightarrow$ byte-identical.

Activation carrier. `subspace.activation` is the scalar/presence activation used by the Space pipeline. Paired `[pos, neg]` tensors still appear in explicit grammar/truth operators that implement catuskoti logic, and in user-supplied truth-set activations, but they are no longer a space-wide output mode.

MASK on SubSpace._active. Two orthogonal per-position tensors: `activation` and `_active`: `[B, N, M]` (modality presence flags). `_apply_mask` is shape-disambiguated:

Mask shape	Effect
Aligns with <code>out.shape[-1]</code> (feature axis)	Element-wise multiply on output
Aligns with <code>out.shape[-2]</code> (position axis)	Zero masked rows of <code>_active</code> ; <code>materialize()</code> gates downstream

Word auto-bind deferred to Reset (fullgraph forward)

The compiled per-batch forward (`BaseModel.enable_compiled_step, torch.compile(fullgraph=True)`) must carry **zero graph breaks**. The word auto-bind — `ConceptualSpace._maybe_autobind_meta`, which grows the WholeSpace codebook and META taxonomy when a freshly-seen percept first lands — is irreducibly host-side: `.item()` loops, `if pid < 0` branches, Python dict/set taxonomy mutation, and (pre-allocation aside) codebook growth. Dynamo cannot trace any of those, so the auto-bind no longer runs on the **forward** path.

Instead (2026-06-03 refactor) it is **deferred to the sentence/document boundary**. `PartSpace._embed_radix` stashes the encountered percept ids (`_forward_input['indices']`) and the pre-pi seed (`_embedded_input`) during the forward — a side-effect dynamo simply replays. On the next hard `ConceptualSpace.Reset` (fired on `hard_eos`, between sentences), `_commit_autobind_from_stash` reads that stash and performs the *same* allocation in eager Python. Same words, same SS rows — only moved from mid-stage to the between-sentence reset, so a downstream tensor op in the same forward no longer sees a mid-forward codebook growth (verified behaviour-preserving across the radix / meta-taxonomy / `mm_xor` convergence suites). Whole-slab configs run one forward per sentence, so the stash is the whole sentence; a per-word/serial config commits the last forward's stash per reset.

Two smaller forward-purity fixes accompany it: `SymbolicSubSpace._synthesize_rule_probs` normalizes branchlessly (`probs / row_sums.clamp_min(tiny)` instead of `if nz.any()`), and the fail-loud `isfinite` guards (here and in `insert_meta / record_lbg_pull`) are gated behind `util.MODEL_DEBUG` — a constant the tracer folds away when off, so the data-dependent `.all()` host sync leaves the compiled graph while divergence still raises under `MODEL_DEBUG` runs and the eager finite-loss guard. (`BASIC_FULLGRAPH=0` relaxes the strict gate to enumerate any remaining breaks; `MODEL_COMPILE=eager` traces without the Inductor C++ backend.)

ShortTermMemory

ConceptualSpace owns `self.stm` — a `ShortTermMemory` instance, a per-batch stack of unquantized CS “ideas.” Post-substrate-refactor, this is the primary structure CS manages when the sparse transform is off — `CS.forward` is then STM bookkeeping (shift + push), with no atomic fold layer. (When the ramsified sparse transform is active, `forward` first runs `_sparse_concept_forward` to produce the concept content, then does the STM write — see **Sparse-coding concept transform**.)

Property	Default	Configurable via
Capacity	8 (within Miller’s 7 ± 2 band)	<code><ConceptualSpace><stmCapacity>N<</code>
Storage	<code>[batch, capacity, concept_dim]</code> buffer + <code>[batch]</code> depth pointers	<code>persistent=False</code> (working state, no
Cleared on	Hard Reset (sentence boundary)	Soft reset leaves it intact

API: `push(b, idea)`, `pop(b)`, `peek(b, n=0)`, `snapshot(detach=False)`, `size(b)`, `is_full(b)`, `is_empty(b)`, `clear(b=None)`, `ensure_batch(batch)`.

STM transition by mode:

- **SERIAL / GRAMMATICAL:** `_stm_shift_and_push(idea)` — shift slots `0..(cap-2)` to take values from slots `1..(cap-1)`, write new idea to the top slot (default `cap = 8`, so slots `0..6` take from `1..7` and the new idea lands in slot 7). Per-word cadence: one push per `PS.forward(IS_t)` call.
- **PARALLEL:** `T` iterations of `PS.forward(CS)` write to STM slots simultaneously; no shift. `T = <subsymbolicOrder>`.

The signal router (`SymbolicSubSpace.languageLayer`) consumes `stm.snapshot()` as its slab input for grammar op dispatch.

Both transitions follow a **predict-then-perceive** discipline (the in-STM `IntraSentenceLayer` predicts the free slot / slab from the retained context before the new event is written). The full STM treatment — predict-then-perceive, attentional filtering (serial runs WITH attention by design), relative-vs-absolute end-states, and the LTM chain of end-states — is in the dedicated `STM.md` chapter.

Lift / Lower as binary GrammarLayer subclasses

(Pre-2026-05-27, Lift / Lower were “substrate-borrowing” — they reached into `PartSpace.sigma` and `ConceptualSpace.pi` for their math. That pattern is retired in Stage 4 of the substrate refactor. `PartSpace.sigma` itself remains LIVE — `PartSpace` owns and uses a single `SigmaLayer` (`self.sigma`), allocated in `__init__` and applied in its forward fold (the Pi/Sigma swap, rev. 2026-06-09, put Sigma/synthesis on `PartSpace`). The per-stage `ConceptualSpace.sigma_in` that once carried the two-loop pi-sigma additive math on CS has since been RETIRED (the 2026-05-29 clean-stack STM experiment bypassed it on the forward path; it was later removed entirely — CS owns no fold). See `doc/old/2026-05-29-clean-stack-stm-basis-arg-radixlayer.md`.)

`LiftLayer` and `LowerLayer` are now first-class **binary GrammarLayer subclasses**, each owning its own internal sub-layer for the pairwise math:

```
class LiftLayer(GrammarLayer):
    arity = 2
    rule_name = "lift"
    space_role = 'C'
    # Internal substrate: self._sigma = SigmaLayer(...)
    def forward(self, left, right):
        return self._sigma.compose(left, right)
    def reverse(self, parent):
        return self._sigma.generate(parent)
```

```

class LowerLayer(GrammarLayer):
    arity = 2
    rule_name = "lower"
    space_role = 'C'
    # Internal substrate: self._pi = PiLayer(...)
    def forward(self, left, right):
        return self._pi.compose(left, right)
    def reverse(self, parent):
        return self._pi.generate(parent)

```

Both register with the signal router (`SymbolicSubSpace.languageLayer`) as CS reduce ops via the existing host-layer registry path. Both inherit `GrammarLayer` butterfly mode (Stage 5) — set `butterfly=True`, `N=N` at construction to enable cross-position cascade.

Cognitive correspondence:

- **Lift** (sigma-style additive synthesis) — “lifting features onto concepts” / predication: “**the boy runs**”.
- **Lower** (pi-style multiplicative contraction) — “lowering concepts into specific percept-realizations” / attribution: “**the running boy**”.

The distinction is preserved at the rule-id level (parser records which op fired) plus the operational difference in the per-pair math.

See `Language.md` for the `GrammarLayer` specifics.

WholeSpace

Role. Owns the **unified word lexicon codebook** with paired (orthographic, semantic) rows. Hosts grammar ops that need codebook write access. The information bottleneck for word identity.

Owned state.

- **self.subspace.what:** a `Codebook` carrying whole-percept prototypes at width `nDim`. Includes per-row meronomy storage (`category_ids`, `part_parents`, `category_logits`). When the configured codebook is `<codebook>quantize</codebook>`, the codebook is a trainable `nn.Parameter` updated under Gray (1990) EMA with **LBG-style variance tracking and splitting**: rows whose running per-component variance exceeds a threshold split along the top-variance eigendirection, with children seeded around the parent $\pm \delta \cdot \text{variance_axis}$. (2026-05-29 Task C.) Set `<codebook>none</codebook>` to retire the SS codebook entirely (degraded mode; the grammar bivector recommender has no `W` to walk on reverse — see `test_mm_grammar_without_vqvae_learns_xor_signal`).
- **insert_paired_word(word, ps_vector)** API: at OOV insert time, creates a paired (orth, semantic) entry — the orth row is initialized from `ps_vector` (a CS-space-dimensioned per-word vector from PS’s Lexicon, per the flat-slab invariant); the semantic row is randomly initialized; the two are parented via `Codebook.set_part_parent(orth_idx, sem_idx)`.
- **get_semantic_row(orth_idx):** O(1) lookup via `_paired_orth_to_sem` map precomputed at insert time.
- **META taxonomy** (2026-05-27 Stage 8): a separate cross-codebook taxonomy mapping (`ps_row`, `ws_row`) $\rightarrow$ `meta_row`. Originally keyed by signed-int refs (positive=PS, negative=SS); the 2026-05-28 plan migrates this to a `.where`-keyed integer-position scheme. META entries are populated by `ConceptualSpace._maybe_autobind_meta` at stage 0 when PS promotes a chunk and SS has a `Codebook` (2026-05-29 Task G — the auto-bind moved from PS to CS so the cross-space allocation fires from the layer that sees both contributions).
- Signal router dispatch site for **codebook-write-required grammar ops** (insert, EMA updates, codebook expansion / culling).

No atomic forward fold layer on SS. The pre-substrate-refactor `self.sigma` ($C \leftrightarrow S$ monotonic `SigmaLayer`) is retired; the $C \rightarrow S$ “naming” is now the codebook snap performed in CS via `SS_subspace.materialize()`

or directly via `insert_paired_word` at write time. Symbol-side grammar math lives on the `GrammarLayer` subclasses (intersection, union, lift, lower, etc.), not on `SS` itself.

Whole-percept presence. Each whole-percept represents presence (1) or absence (0) of a named entity (the compact emission migrating to `SymbolSpace` is the symbol; `WholeSpace` itself holds whole-percepts). Mapping: `presence = (activation + 1) / 2`.

Lookup chain (surface $\rightarrow$ meaning):

1. Surface bytes enter PS via `InputSpace`.
2. MPHF on PS produces a row index into `PS.vocabulary` (the Lexicon Embedding).
3. The PS Lexicon vector (CS-space-dimensioned) flows through `PS.forward` $\rightarrow$ CS via STM push.
4. When CS-state activations need to match a whole-percept, they quantize against `SS.subspace.what` (the codebook).
5. Quantization can resolve to an orth row; `get_semantic_row(orth_idx)` returns the parented semantic row $\rightarrow$ the meaning.

Codebook geometry. `subspace.what.getW().shape == (nVectors, nDim)`. Each row is a free coefficient vector over conceptual axes; not unit-norm. Retrieved via Euclidean L2 (per Codebook's metric contract). `SS.nVectors` must be sized to at least `2 * expected_OOV_word_count` to accommodate paired rows (e.g., MM_5M sets `nVectors=131072 = 2 * 65536`).

Multi-stage SS. In multi-stage configs, the lexicon mirror is wired to `self.wholeSpaces[-1]` (the terminal stage; `model.wholeSpace`). The terminal SS must be sized for paired-row capacity. See `doc/old/2026-05-26-two-loop-pi-sigma-substrate.md` for the dispatch rationale.

Conceptual order. `BasicModel` stores per-stage `ConceptualSpace` / `WholeSpace` instances for `subsymbolicOrder`. Symbol dimension is geometrically partitioned per order. See `Reasoning.md`.

See `Philosophy.md`, `Logic.md`, `Mereology.md`, and `Language.md`.

Monotonicity of the lift / lower chain

With `monotonic=true`, the $P \rightarrow C \rightarrow S$ chain is an **order-preserving map on a positive cone**.

Three pieces:

1. **Truth-set activations may live on the positive cone** $[0, 1]^{\sim\{2K\}}$ as user-supplied paired `[pos, neg]` bivectors. The componentwise partial order $\leq$ is the parthood order for that truth surface.
2. **The Pi / Sigma maps are restricted to $W \geq 0$ entry-wise** (`monotonic=True` selects `NonNegativeInvertibleLinearLayer` or `NonNegativeLinearLayer`):

$$a \leq b \text{ componentwise} \implies Wa \leq Wb \text{ componentwise}$$

3. **Therefore Pi / Sigma preserve parthood pole-by-pole for truth-set bivectors and other positive-cone activations.**

The bivector layout keeps the contradiction corner `[1, 1]` distinct from the ignorance corner `[0, 0]` under positive `matmul` — a single bitonic axis would let $aP - aN$ cancel under summation.

The `ImpenetrableLayer` regularizer maintains an antichain of same-rank prototypes, complementing the structural `.where`-uniqueness from the codebook offset registry. See `Logic.md` Section Parthood as Projection and `Philosophy.md`.

SyntacticSpace — retired

The standalone `SyntacticSpace` class has been retired. Grammar / parser / derivation-tree machinery now lives on `SymbolSpace`, which attaches a `SyntacticLayer` to `WholeSpace` (the canonical grammar host). The binary-derivation behavior previously documented here is preserved by `WholeSpace.compose`; words are still concepts encoding grammatical rules, and the derivation is still stored as word tuples — just on `WholeSpace` rather than on a separate `Space`.

See `Language.md` for grammar and parser dispatch.

OutputSpace

Role. Maps symbolic (or syntactic) vectors to task targets via linear projection.

Forward. $y = W_{out} * x + b_{out}$. Always `reshape=True` — the `[B, nSymbols, symbolDim]` tensor is flattened before projection.

Reverse. Pseudoinverse of `W_out`. Text mode snaps each output vector to the nearest codebook entry (nearest-neighbour lookup).

getEmbeddedIO() override. Returns raw target dimensions rather than encoded dimensions, so loss is computed in the output vocabulary space.

Text mode generation. Autoregressive: each step’s predicted token vector is snapped to its nearest codebook entry and fed back as input until `maxResponseLength` or EOS.

Key parameters. `nActive`, `nDim`, `nVectors`.

Layer. `LinearLayer` with (`bias`, `temp`) for ergodic mode.

Range. Forward rescales `[-1, 1]` to the original data range via `Data.denormalize()`. Reverse applies `Data.normalize(x, "output")`.

Invertibility. Pseudoinverse; not exactly invertible in general.

Short-Term Memory

2026-06-02 update (subsymbolic analyzer). Operators no longer enter the STM idea space. They are kept in the SS **codebook** (`WholeSpace.insert_operations`, wired into `SymbolicSubSpace.__init__`) and resolved as a soft superposition over the operator-prefixed parse tree; the STM idea slots hold only **combined meanings** – an operator defines *how* meanings combine, contributing none of its own. See `doc/old/2026-05-30-subsymbolic-analyzer-terminal-emitter.md`.

Status (2026-05-30): new chapter for the STM serial / parallel modes work (`doc/old/2026-05-29-stm-serial-parallel-modes`). Documents the per-batch STM buffer, the predict-then-perceive cadence (serial and parallel), the attentional-filtering regime (serial runs **with** attention by design — the old serial-vs-attention guard was lifted), the routing-parser SS-analysis / CS-execution split, the in-STM `IntraSentenceLayer` AR predictor, masked-word reconstruction via priming, the relative-vs-absolute end-state preservation with its content-aware learn-score gate, and the LTM chain of end-states feeding inter-sentence prediction. Where a piece is a deliberate scaffold or a deferred wiring, this chapter says so.

Overview

`ConceptualSpace` is, post-substrate-refactor, an STM container plus a grammatical CPU — it owns no atomic forward fold (see `Spaces.md`). The **short-term memory (STM)** is the structure CS manages: a per-batch stack of unquantized CS “ideas” that the model accumulates across a sentence and reduces at the sentence

boundary. This chapter is the single reference for what the STM is, how it fills, what reads it, and how its end-states chain into long-term memory.

Two cadences write the STM, selected by `<symbolicOrder>` (see Architecture.md):

SERIAL one idea per word; predict the free slot, then perceive (push)
PARALLEL one whole-slab pass; predict the slab, then perceive (set all slots)

Both follow a **predict-then-perceive** discipline: before the freshly materialized event is written into the STM, the in-STM predictor reads the retained context and stakes a prediction; the just-written event is then its supervision target. This is the concept-space analogue of a language model’s next-token objective, scaled to the in-sentence working set.

The trainable target configuration is `MentalModel.xml` (serial, `<data><dataType>embedding</dataType>`, attention on, sentence prediction on, FineWeb data); the comparison framing against a transformer LM is Section 12.

1. STM data model

ShortTermMemory (Layers.py:8982) is a **Layer** (not a **Space**): it carries no **SubSpace** and no forward / reverse tensor-map contract. **ConceptualSpace** builds one at construction as `self.stm` and treats it as the primary structure it manages.

Capacity. Default 8, set via `<ConceptualSpace><stmCapacity>N</stmCapacity></ConceptualSpace>` (`DEFAULT_CAPACITY = 8`); the legacy `<wMax>` alias is retired. Eight sits inside Miller’s 7 ± 2 band — the working-set size psycholinguistics ascribes to human short-term memory. The capacity is the rolling-window length: at steady state the STM holds the last `cap` ideas and the oldest falls off as new ones arrive.

Buffer. The data is a single per-batch tensor `[B, cap, concept_dim]` plus a `[B]` long depth-pointer vector recording how many slots each row has filled (saturating at `cap`). The buffer is a plain registered buffer with `persistent=False`: STM contents are runtime working state, not learned weights, and never enter the checkpoint. `concept_dim` is the CS feature width D — the **what** columns only; the positional / temporal (**where** / **when**) columns are trimmed at the PS→CS boundary so the slab the grammar dispatch runs over is a clean concept-dim block (`_stm_payload_dim = concept_dim`).

SymbolicSubSpace as the data carrier. After the 2026-05-21 STM-Layer refactor the **ShortTermMemory** object is itself data-free: its `_buffer` / `_depth` properties proxy to the owning **SymbolicSubSpace**’s `_idea_*` buffers via `attach_word_subspace`. A standalone fallback buffer exists for bare-construction unit tests, but in the live model the bytes live on **SymbolicSubSpace** and **ShortTermMemory** is a thin accessor over them.

API. `push(b, idea)`, `pop(b)`, `peek(b, n=0)` (`peek(b, 0) = most recent`), `snapshot(detach=False)`, `size(b)`, `is_full(b)`, `is_empty(b)`, `clear(b=None)`, `ensure_batch(batch)`. The signal router consumes `stm.snapshot()` as its slab input.

Lifecycle. Cleared on hard **Reset** (sentence boundary) — the per-batch idea stack drops everything from the just-finished sentence and the next sentence starts empty. Soft reset leaves the STM intact (see Spaces.md).

2. Serial sequencing

In **SERIAL** / **GRAMMATICAL** mode each word traverses a per-word path: MPHf surface lookup → `PartSpace.forward` ($\pi(x) + \sigma(x)$, no outer `tanh`) → `ConceptualSpace.forward`, which does the STM bookkeeping. The whole pass is **predict-then-perceive per word**, implemented in `ConceptualSpace.forward` (Spaces.py:14134):

1. **Snapshot + predict the free (newest) slot.** Before writing the new event, `_stm_predict_then_perceive_serial` (Spaces.py:12290) takes `stm.snapshot()` and runs the in-STM predictor from the **retained context**

- the snapshot with the oldest slot dropped (`snap[:, :-1]` when $\text{depth} \geq 2$; the whole snapshot when $\text{depth} = 1$). The prediction is stashed on `self._stm_predicted_idea`. On the first word (empty STM) the prediction degenerates to zeros and **no** loss accumulates (there is no prior context to predict from). The serial predictor folds the context with an order-invariant **sum** over the slot axis, so dropping the oldest (regardless of which physical slot holds it) yields the same prediction.
- 2. **Perceive (overwrite via `_stm_shift_and_push`).** `_stm_shift_and_push(idea)` (Spaces.py:12056) is the perceive step. At capacity, slots $0 \dots (\text{cap} - 2)$ shift into $1 \dots (\text{cap} - 1)$ and the new idea lands in slot 0 (the newest slot); the oldest idea (the last slot $\text{cap} - 1$) falls off. Below capacity it is a plain push: shift the occupants right and write slot 0.
- 3. **Accumulate the intra-loss.** The held prediction is scored against the just-perceived **idea** as $\mathcal{L}_{\text{intra}} = \text{MSE}(\hat{c}_t, c_t)$ (see Section 6).
- 4. **Language dispatch.** The signal router dispatches grammar ops over the STM contents (read-only via CS, write-required via SS).

Convention pin — newest-at-slot-0, shift-RIGHT. The **free** / **newest** slot is **slot 0** and the rolling window shifts **right**, dropping the oldest (the last occupied slot, $\text{cap} - 1$ at capacity). `peek(n)` counts n back from the newest, so it reads slot n directly; `snapshot()` returns the live slab **newest-first**. The “predict the free slot” step predicts slot 0, and the retained context that conditions it is `snap[:, :-1]` (everything but the soon-to-be-evicted oldest slot). A named primitive `_stm_shift_for_predict` (Spaces.py:12121) exists for the literal “rotate so the free slot is free” step, but the hot serial path deliberately does **not** call it — `_stm_shift_and_push` already does shift-right-then-write in one pass, and a separate physical shift would double-shift the buffer. The helper is kept off the forward path and exists for independent unit-testing of the named step.

(History: this convention was flipped from the original newest-at-top / shift-LEFT layout; the flip preserves every semantic — `peek` still returns the most-recent idea, the reduce still collapses to the same root, and the relative end-state still yields the same predicate/idea1/idea2 — only the physical slot order changed.)

3. Parallel sequencing

In PARALLEL mode the per-stage forward sees the whole sentence at once: the slab $[B, N, D]$ is the STM, each of the N positions its own slot. The same predict-then-perceive discipline applies whole-slab, implemented in `ConceptualSpace.forward` via the `folded.shape[1] > 1` branch:

1. **Predict $\hat{C}$ from the previous slab.** `_stm_predict_then_perceive_parallel(folded)` (Spaces.py:12325) snapshots the previous STM and runs the predictor **per slot** (no cross-slot collapse), producing a $[B, \text{prev}_N, D]$ slab stashed on `self._stm_predicted_slab`.
2. **Perceive via `_stm_set_all_slots`.** `_stm_set_all_slots(slab)` (Spaces.py:11986) writes all N positions directly as the slot stack (no shift, no mean-reduction). The slab is position-ordered (position 0 oldest, $N - 1$ newest); under the newest-at-slot-0 convention it is **flipped** along the position axis so the slab’s newest lands at slot 0. Capacity-clip: if $N > \text{cap}$ it keeps the last **cap** positions (drop oldest, mirroring the serial rolling window) then flips; if $N < \text{cap}$ it zeroes the unfilled (older) tail so a prior pass’s state cannot leak through. The matching `_stm_predict_then_perceive_parallel` flips the **folded** target the same way before scoring $\mathcal{L}_{\text{intra}}$, so the per-slot prediction (computed over the newest-first previous STM) and its target stay slot-aligned — byte-identical to the old oldest-first alignment.
3. **Loss only when shapes align.** $\mathcal{L}_{\text{intra}}$ accumulates only when the predicted slab shape matches **folded** exactly. On the first pass (empty previous STM) the prediction is zeros and loss is skipped; when the previous slot count differs from N the per-slot prediction does not align with the target and the pinned decision is to skip loss for that pass rather than fabricate an overlap. Loss resumes once the STM steady-state slot count matches the slab width.

The earlier “mean-reduce to a single idea, then shift-push” pattern was a bug in parallel mode — it destroyed per-slot identity and made the reverse pipeline emit the same recon vector in every slot. The slot-preserving `_stm_set_all_slots` is the fix.

4. Attentional filtering

Serial mode IS the attentional-filtering regime. The old serial-vs-attention guard was **lifted**: `MentalModel.xml` runs serial *with* attention by design (`<symbolicOrder>1</symbolicOrder>` with `<attention>primer</attention>` on `ConceptualSpace`; the legacy `<hasAttention>` boolean is deprecated and inert, superseded by the `<attention>` element — off/primer/second-order/low-rank). Serial sequencing and attentional filtering are the same regime, not mutually exclusive options.

CS→PS meronymic Gaussian mask. On the per-word serial path the input to each word is not a raw slice but a Gaussian-windowed contextual trace, built by `gaussian_window_word(full_seq, center_k)` (`Models.py:2323`). For the processed word at center k a single envelope over the whole per-sentence percept sequence $[B, T, D]$ is formed:

$$w_i = \exp\left(-\frac{(i-k)^2}{2\sigma^2}\right), \quad \sigma = \text{maskRate} \cdot T,$$

normalized so the center weight $w_k \approx 1$. The peak sits at the current word; words far from k are attenuated toward 0 by the Gaussian tail. The windowed percepts are mapped into conceptual space and **summed**, so word k ’s representation carries a faint meronymic trace of its neighbours (the local context *is* part of the embedding signal). This **replaces** the prior BERT-style hide-a-token `create_ir_mask` on the per-word grammar path; it is not target-hiding — the center word is preserved. The centring is **hardcoded**: peak at the current word.

CS→SS taxonymic mask. On the symbolic side the inverse recommender zeros **non-admissible** codebook rows before a lift / lower (union / intersection) reduction so an operand can only resolve to a row of the right grammar category and conceptual order. The admissible row set is built by `priming_kwargs_for_slots` (`Language.py:7340`) — intersecting `refs_by_category` with `refs_by_order` per slot — and applied in the recommender’s `_row_mask` (`Layers.py:10344`), which admits only the category/order-matched rows plus the $\perp$ / $\top$ sentinels. A parallel `_row_weights` path multiplies a taxonymic **priming** mask over the admitted rows (the `left_priming` / `right_priming` weights from `taxonomy.priming_mask`).

Guidance-signal contract. The mask is a *guidance signal*: it biases which codebook rows / sequence positions the reduction may consult, given the current word and grammar state. The centring policy (CS→PS) is the swappable seam — the present text-reading variant centres on the current word; a future image-reading variant would swap that policy (centre on a fixated region) without changing the rest of the pipeline.

Out of scope. Learning the mask is out of scope for this work — both masks are **hardcoded** (Gaussian centring fixed at the current word; taxonymic admissibility read structurally from the codebook’s category / order metadata). A learned attention policy is a documented future direction, not built here.

5. Routing parser: SS-analysis vs CS-execution

The grammar runs through the signal router (`LanguageLayer`, `Language.py:3493`); `SymbolicSubSpace` owns it as `self.languageLayer`. Conceptually the work splits in two:

- **SS-analysis** — `SymbolicSubSpace.compose` (`Language.py:7693`) is the analysis stage: a soft superposition over the taxonymic codebook that selects, per-space, a **hard rule dict current_rules** = `{space-role: [rule_id, ...]}`. It chooses *which* reductions fire.

- **CS-execution** — actually applying the chosen reductions (lift, lower, union, intersection, swap, quantize, not) to the concept tensors runs CS-side in `ConceptualSpace.forward` and the WholeSpace stack-route path, with the per-space `SyntacticLayer` cursors (`Language.py:5453`) executing the unary π / σ folds on reverse. Only lift / lower / union / intersection consult the codebook (inverse-recommended); swap / quantize / not are tensor-only.

Honesty — the split is a clean code boundary only on the default-only path. When the grammar is *default-only* (every rule is the unary π / σ fold), `compose` emits `current_rules` from the grammar XML and runs **no** tensor reduction, and the per-space `SyntacticLayer.forward` / `reverse` cursors do the CS-side execution — a genuinely clean analysis/execution separation. On the **full-router** path, however, `LanguageLayer.compose` does **both**: it selects the rules *and* folds the slab tensorially through the op modules (`BinaryStructuredReductionLayer.forward` $\rightarrow$ `op(left, right)`), caching the root state. The per-space `SyntacticLayer` cursors are then **deliberately bypassed** on that path — guarded by `not _grammar_is_default_only` (`Language.py:5731`) — precisely so the reduction is not double-applied. So in the full-router case the SS-analysis and CS-execution stages are **co-located** inside `LanguageLayer.compose` rather than separated across modules. The audit found this; it is a documented task-5 follow-up, not a finished boundary.

`_grammar_is_default_only` is computed from the configured grammar at `SymbolicSubSpace.__init__` (`Language.py:6522`) and gates both `compose` and `generate`.

6. IntraSentenceLayer

`IntraSentenceLayer` (`Layers.py:6632`) is the in-STM autoregressive predictor — the layer that stakes the prediction in predict-then-perceive. It is owned by `ConceptualSpace` (`self.intraSentenceLayer`).

Architecture: combined PI-then-Sigma, no intermediate tanh.

- `self.pi`: `PiLayer(concept_dim $ \to $ working_dim)`, `invertible=True`, `nonlinear=True` — the log-domain multiplicative boundary fold lifts each STM slot into the working width and bounds it to $[-1, 1]$.
- `self.sigma`: `SigmaLayer(working_dim $ \to $ concept_dim)`, `invertible=True`, `nonlinear=False` — a raw linear $Wx + b$ that collapses the lifted slots into the predicted idea.

The defining requirement is that there is **no extra activation interposed** between the PI body and the Sigma body: `sigma` is built `nonlinear=False` (a raw $Wx + b$), so no additional tanh sits between PI's output and Sigma's input. This is **not** a linear fusion — `pi` is `nonlinear=True` (its symmetric log-domain $(1+x)/(1-x)$ embedding is an intrinsic nonlinearity), so the composite is a nonlinear PI followed by a raw-linear Sigma, not two fusable linear cores. `working_dim` defaults to `concept_dim`, keeping both sublayers square isomorphisms so the parallel per-slot round-trip $\text{reverse}(\text{forward}(x)) \approx x$ is exact up to the LDU inverse tolerance.

Forward signature. `forward(prior_slots, routing=None, parallel=False)` with `prior_slots`: [B, K, D]:

- **Serial collapse** (`parallel=False`, the primary regime): PI-lift every slot, **sum-fold over the slot axis**, then Sigma-collapse to one idea $\rightarrow$ [B, D].
- **Parallel per-slot** (`parallel=True`): PI $\rightarrow$ Sigma per slot, no cross-slot mixing $\rightarrow$ [B, N, D].

The serial collapse is many-to-one (sum over K slots), so its **reverse** is necessarily approximate (it divides the recovered fold equally across the $k = \text{stm_capacity} - 1$ slots); the parallel per-slot path is width-preserving and exactly invertible.

Loss. $\mathcal{L}_{\text{intra}} = \text{MSE}(\hat{c}_t, c_t)$ via `intra_loss(pred, target)`. `ConceptualSpace` accumulates it per step (`_accumulate_intra_loss`, `Spaces.py:12248`) into a live grad-carrying running sum and `consume_intra_loss` (`Spaces.py:12271`) returns the per-step mean once per batch (`runBatch` consumes it

post-body / pre-backward, mirroring the ARMA term). The weight is `<intraLossWeight>` (default 0.1); the term is gated off when grad is disabled or the weight is non-positive.

7. Per-word router firing

The `<routerWireSerial>` knob (model.xsd:76) gates when the router fires on the serial path:

Value	Behaviour
<code>per-word</code>	fire per word; boundary fire off
<code>boundary</code>	fire only at the sentence boundary
<code>both</code>	(default) per-word AND boundary both fire
<code>off</code>	neither fires

The **per-word fire** runs `symbolSpace.compose` over the current STM snapshot *before* `cs.forward` for the next word, populating `symbolSpace.current_rules` for the SS dispatch (Models.py:6053). It lives in the host-side loop, outside the per-iteration captured graph, so even a full `languageLayer` path that introduces a host sync cannot break the per-word capture gate. The **boundary fire** ([Layers.py / Models.py boundary hook]) runs iff `router_wire_serial` in ('boundary', 'both').

Honesty — routing conditioning of the predictor is DEFERRED. `current_rules` is a **host-side rule dict** {'S': [rule_id, ...]}, **not** a [B, concept_dim] tensor. `IntraSentenceLayer.forward` accepts a routing tensor and projects it into a Sigma bias, but `ConceptualSpace's _intra_routing_for_predict` (Spaces.py:12166) only forwards routing when it is a CS-width tensor; given a dict it returns `None`. So the per-word fire wires the SS dispatch context but does **not** yet flow a routing tensor into the intra predictor — the predictor runs **without** the routing bias for now. Rule-aware predictor conditioning awaits a future rule-dict → CS-width projection (deliberately not built here: rule-distribution width ≠ concept_dim, and a naive bridge risks capture-gate breakage). This is the current state, not a finished feature.

8. Masked-word reconstruction via priming

Masked words enter the pipeline as **all-zeros** percept slots. On reverse, the router fills the blank via a best-fit codebook walk biased by the taxonymic prior, POS selection, and accumulated prediction — it reconstructs the most plausible word for the slot given the grammar context and the codebook geometry.

Priming machinery. The bias enters through `left_priming` / `right_priming` in `Basis.lift` / `Basis.lower` (Spaces.py:1202), forwarded down to the inverse recommender's `_row_weights` (Layers.py:10364) where each admitted codebook row is scaled by its taxonymic priming weight (the $\perp$ / $\top$ sentinels pinned to 1.0). Primed rows are preferred in the argmax that selects the operand. `test/test_primed_reverse_hard_mask.py` exercises this path.

Tests. The reverse roundtrip and reconstruction tests are `test/test_stm_reverse_roundtrip_lift_lower.py`, `test/test_stm_reverse_roundtrip_union_intersection.py`, and `test/test_stm_recon_from_cleared_cache.py`.

Honesty — the recon-from-cleared-cache test is an honest xfail. `test/test_stm_recon_from_cleared_cache` marks the top-*k* word recovery assertion `xfail` (not relaxed to a trivially-true bound) because of **two pre-existing upstream bugs**, both out of scope here and flagged as separate follow-ups:

- **Finding A** — on the untrained config the per-word forward fills the CS STM with NaN: `conceptualSpace.stm.snapshot()` is non-finite, so the reduced single-*S* seed is already NaN before reverse runs.

- **Finding B** — even with a *finite* seed, the reverse perceptual leg (`PartSpace.reverse`) turns it NaN.

The non-xfail assertions in that file pin everything that *does* hold today (the per-op reverses reconstruct [B, N, D_c]; the decode uses the real perceptual codebook, no reimplementations). When the two upstream bugs are fixed the xfail will xpass. Documenting the contract honestly: the priming-biased recon path is built and tested, but end-to-end word recovery on an untrained model is blocked upstream.

9. Relative vs absolute end-states

At the sentence boundary the STM is reduced to its end-state. The shape of that end-state depends on whether the sentence asserts an **absolute** or a **relative** truth.

- **Absolute** sentences reduce to a **single idea** (depth 1) — the root *S* the start-symbol reduction wrote. This is the dominant path and the one the IR loss consumes.
- **Relative** sentences (the `part / isEqual` predicate family, which produce the `REL_T` relative start symbol) are **preserved at depth 3** as the `[predicate, idea1, idea2]` end-state. Under the newest-at-slot-0 convention these are stored newest-first, so the predicate (oldest constituent) sits at the **last** slot (depth - 1), `idea1` at depth - 2, and `idea2` (the folded-newest-rest) at slot 0. `learn_relations_from_stm` reads them from those slots so `_maybe_learn_relation` receives the identical predicate / `idea1` / `idea2` it did under the old oldest-first (`slots 0/1/2`) layout.

The mechanism is a per-row `protect_depth` gate threaded into `_stm_reduce_to_single_S` (Models.py:5470). Each masked micro-step (`_stm_bounded_reduce_step`, Models.py:5272) folds only while `depth > protect_depth`: absolute rows pass `protect_depth = 1` (collapse all the way), relative rows pass `protect_depth = 3` (stop at the depth-3 end-state).

Conservative detection. `_sentence_relative_mask` (Models.py:5387) decides per row from the grammar — it scans `current_rules`’ *S*-space-role for a `TheGrammar.is_relative_rule` rule_id (`REL_T` LHS, or an `isEqual / queryPart / assertPart / part` op). It **defaults to collapse on any uncertainty**: a false positive would stop an absolute sentence’s collapse and break the dominant path and the IR loss, so missing / empty rules, a grammar with no relative rule, or any ambiguous shape all return all-False. The absolute path stays **byte-identical** in every fall-through case.

Learn-score acceptance gate. Concept-codebook insertion of a learned relation is gated by a content-aware `learn_score` (`_compute_learn_score`, Spaces.py:13312):

Terminology (2026-06-21 convention). The CS part whole relation table is the **Concept codebook** — each entry is a *concept* tying one part-percept to one whole-percept by reference. Earlier text called this the “symbol table”; the de-overloaded name is *concept* (a *symbol* is the 0-D SymbolSpace reference *to* a concept, not the relation itself). The taxonomic / perceptual codebooks consulted for lift/lower operands (Sections 4, 8) are distinct and keep their names.

$$\text{learn_score} = \text{children_in_codebook} \times \text{is_truth_obvious} \times \text{resolves_contradiction},$$

each factor in $[0, 1]$. A relation is accepted **iff** `learn_score` $\geq$ `<truthCriterion>` **and** `truthCriterion < 1`. The **default is 1.0** — truth-learning is OFF by default (nothing learned or recorded); opt in by lowering `truthCriterion` toward 0. Because the factors multiply, a **low is_truth_obvious does not on its own block**: lies and uncertain relations can still be learned if the other two factors are high. At `truthCriterion = 1` nothing is learned; at 0 everything is.

Accepted insertions carry a **tetralemma trust 4-tuple** (t, f, b, n) (TRUE / FALSE / BOTH / NEITHER, summing to 1) computed by `_tetralemma_trust` (Spaces.py:13343) from the TruthSet posture via `assess()`, and bound onto the relation’s META node by `insert_relation` (Spaces.py:16033) with the predicate as the parent and the two ideas as its taxonomy children.

Honesty — the learn-score factors read the **GLOBAL** truth layer. The `is_truth_obvious` and `resolves_contradiction` factors currently read the `global_truth_layer.assess()` — support and conflict respectively (`_learn_score_is_truth_obvious / _learn_score_resolves_contradiction`, `Spaces.py:13279`) — behind a swappable seam (each factor is an independently overridable method, the plan’s required test seam). A **per-relation projection** is the documented refinement; the global read is a first cut, not the final formula.

Truth **recording** into the `TruthLayer` is governed by the *same* continuous `<truthCriterion>` bar (a per-cell activation is recorded when its clamped magnitude clears `truthCriterion`), so a single knob governs both recording and learned-relation acceptance — there is no separate gold path and no binary switch. The user-provided `<truth>` gold set is ingested by `store_truths`, which drops `truthCriterion` to 0 for the ingestion epoch (capturing every provided gold truth), then restores it; the same recording block fires during normal training per the configured `truthCriterion`. The binary `accumulateTruth / truthMinMagnitude` switches are **retired**. See `Params.md` and `Logic.md`.

Behavioural note. Unlike the retired binary gate (which kept recording off during training), truths are now accumulated continuously during training per `truthCriterion`; raise it toward 1 to suppress recording.

10. LTM as the chain of STM end-states

Long-term memory (LTM) is the **full chain of STM end-states**. It lives on `InterSentenceLayer` (`Layers.py:5744`) as a per-row bounded `collections.deque _stm_end_states` of size `<ltmCapacity>` (default 1024). Each entry is a time-ordered tuple

(`depth: int`, `payload: [depth, D] tensor`, `tetralemma: tuple | None`)

where `depth` is 1 for an absolute end-state and 3 for a relative [`predicate`, `idea1`, `idea2`] end-state (so the payloads are ragged — a fixed register-buffer tensor does not fit, hence the per-row deque). The chain is transient host-side state (`persistent=False` semantics, not in `state_dict`).

`observe_stm_end_state(depths, payloads, tetralemmas=None)` (`Layers.py:6162`) records **every** sentence’s end-state — it is **not** gated by `truthCriterion`. The distinction is load-bearing: `truthCriterion` gates only the separate Concept-codebook insertion of *learned relations* (Section 9), whereas LTM is the AR sequence the inter-sentence predictor consumes and must see the full history. A non-finite payload **raises** (fail-loud on numerical divergence); the deque `maxlen` evicts the oldest entry once full. `get_stm_chain(n=None, b=0)` (`Layers.py:6262`) returns the last `n` (or all) end-states for a row, oldest-first.

LTM is the **full chain**; the **TruthSet** is the accepted-belief subset (the relations that cleared the learn-score gate and were inserted into the Concept codebook). LTM keeps everything that happened; the `TruthSet` keeps what was believed.

11. Inter-sentence prediction

A **lifted IntraSentenceLayer instance** (`_inter_predictor`, `Layers.py:5946`) predicts the next end-state over the LTM chain — the same predictor class as the in-STM one, instantiated at the inter-sentence level. Its chain window is $K = \min(\text{ltmCapacity}, 8)$ (`_inter_chain_window`): the AR signal that predicts the next end-state lives in the last handful of sentences, so a small bounded window is used rather than the full `ltmCapacity`.

`predict_next_end_state(b=0)` (`Layers.py:6326`) produces the next end-state **shape** $(\hat{d}, \hat{p}[\hat{d}, D])$:

- **Chain reduction (ragged $\rightarrow$ fixed).** Each chain entry is reduced to its **root** by `_reduce_end_state_to_root` (`Layers.py:6316`). Under the newest-at-slot-0 convention the end-state is stored newest-first, so the root lives at the **last** slot (`depth - 1`): for an absolute end-state (`depth 1`) that is slot 0 (the collapsed

idea); for a relative end-state (depth 3) it is slot 2, the predicate (the head the relative structure hangs off). The last K roots form a $[1, K, D]$ context, left-padded with zeros so the most recent sits at the tail (newest-at-1, like the ARMA ring).

- **Root prediction.** `_inter_predictor.forward(context, routing=None, parallel=False) → [1, D]`, the predicted root.
- **Loss.** $\mathcal{L}_{\text{inter}} = \text{MSE}(\hat{p}, p)$ on the roots, accumulated by `_accumulate_inter_loss` (Layers.py:6413) and drained by `consume_inter_loss`, weight `<interLossWeight>` (default 0.1). The actual root is detached so the loss trains `_inter_predictor`, not the perception path. `observe_stm_end_state` scores the prediction made for a row against the end-state that actually arrived.

The prediction is consumed by `generate_sentence` via the `_c_prior [depth, D]` staging path: a predicted next-end-state shape is staged across the first `depth` STM slots as a sentence-level conditioning bias (`ConceptualSpace.forward`'s slotwise `_c_prior` branch, Spaces.py:11201).

Honesty — $\hat{d}$ is a **copy-last AR prior; reduction is slot-0 only**. The predicted **depth** $\hat{d}$ is a simple AR prior: the depth of the **most recent** end-state in the chain (a relative sentence tends to be followed by structure of the same shape; an absolute by an absolute). This delivers **depth** in $\{1, 3\}$ without a separate learned head, but it is a **scaffold** — a tiny `concept_dim` $\$ \backslash \text{to} \$$ 2 argmax head is the documented upgrade path. Likewise the chain reduction collapses each end-state to its **slot-0 root only** (mean-over-depth is the obvious alternative; the root carries the sentence-level signal, mirroring the ARMA `_pool_sentence_rep`). A non-finite predicted root raises.

12. nanochat comparison framing

The trainable target is `MentalModel.xml`: `<symbolicOrder>1</symbolicOrder>`, `<data><dataType>embedding</dataType>` `<attention>primer</attention>` (`ConceptualSpace`; the legacy `<hasAttention>` boolean is deprecated and inert), `<sentencePrediction>true</sentencePrediction>`, FineWeb data (`<shardDir>data/fineweb</shardDir>`). This is the configuration that exercises the full STM stack — serial sequencing with attentional filtering, the in-STM and inter-sentence predictors, and the LTM chain.

The comparison loss is

$$\mathcal{L} = \mathcal{L}_{\text{IR}} + \mathcal{L}_{\text{intra}} + \mathcal{L}_{\text{inter}}$$

the masked-LM information-reconstruction term at the subsymbolic (PS) (see Spaces.md) plus the in-STM next-idea term (Section 6) plus the inter-sentence next-end-state term (Section 11). This is analogous to a transformer language model's next-token cross-entropy, but computed **in concept space** rather than over a token vocabulary: $\mathcal{L}_{\text{intra}}$ predicts the next idea within a sentence, $\mathcal{L}_{\text{inter}}$ the next sentence's end-state shape, and $\mathcal{L}_{\text{IR}}$ reconstructs masked content — the three together are the concept-space counterpart of the autoregressive LM objective a system like nanochat trains.

Out of scope. The comparison harness itself — running this against a transformer baseline and reporting the numbers — is a **separate plan**. This chapter documents the loss that *would* be compared and the configuration that produces it; it does not build the benchmark.

See also

- Spaces.md — `ConceptualSpace` as STM container; the `ShortTermMemory` API; Sigma / Pi ownership.
- Architecture.md — modes of operation (serial / parallel); the per-word operational flow; `InterSentenceLayer` ARMA predictor.
- Language.md — the signal router, `compose` / `generate`, `GrammarLayer` reductions.
- Mereology.md — parthood as clipped-cosine projection; the relations the relative predicates draw on.

- Reasoning.md, Logic.md — the truth surfaces and tetralemma trust.
- Params.md — <stmCapacity>, <intraLossWeight>, <interLossWeight>, <routerWireSerial>, <ltmCapacity>, <truthCriterion> (the single continuous truth bar).

Language

This document is synchronized to the code as of 2026-06-02. The source of truth is `bin/Language.py`, `bin/embed.py`, and (for the perceptual analyzer) `bin/perceptual_analyzer.py`. (The earlier `bin/typed_stack.py`, `bin/stm_driver.py`, and `bin/parse_state.py` modules were retired in the 2026-05-21 / 2026-05-29 refactors; their functionality folded into `bin/Layers.py` and `bin/Language.py`.)

2026-06-02 deltas (subsymbolic analyzer + terminal emitter). Plan: doc/old/2026-05-30-subsymbolic-analyzer-

- **PS/SS grammar sections.** A `.grammar` file may nest its <compose>/<generate> under <PartSpace> and <WholeSpace>. `Grammar.configure` parses them into separate rule tables: `ps_rules` (space-role P, read by the PS analyzer) and the canonical symbolic rules (`ws_rules`). A bare <compose>/<generate> file loads as <WholeSpace> (backward-compat). The legacy `.cfg` loader is gone.
- **Grammar rewrite.** `*_MARK` categories and the copy/swap MARKER helper rules are deleted; surface markers are learned and owned by the operator. Each operator-argument position is its own category (`CONJ_L45/CONJ_R45`, ...). copy/swap are retired from the symbolic grammar (kept only as the T5 elision surface policies).
- **SurfaceSchema + absorb/emit** (`bin/Layers.py`). Five universal templates T1-T5 declare each operator's marker slot + order; T4 `BINARY_JUXTAPOSE` is the default. `GrammarLayer.absorb` binds a co-occurring marker to the operator (many-to-one); `emit` replays it from recorded route metadata, never the lossy `generate()`.
- **Operators in the SS codebook, not the STM idea space** (amends plan decision #2). `WholeSpace.insert_operations(grammar)` registers each operation in a dedicated operator codebook on `WholeSpace` (`_operation_vectors` / `_operation_positions`), separate from the `subspace.what` whole-percept codebook so the percept / idea / `.where` position namespace is untouched; it is wired into `SymbolSubSpace.__init__` so every built model's operator-prefixed parse-tree nodes are codebook-resolvable. The STM idea space holds only combined meanings – the operator says *how* meanings combine, contributing none of its own. The rule-id stays in `.where` (its presence marks a slot as a *computed* idea, which is by definition not a codebook vector).
- **ObjectSubSpace** (`bin/Language.py`) – the PS-meronymic carrier analogue of `SymbolSubSpace`: span buffers + parent/child links + route ids/scores + the marker-route replay fields (`_marker_ps_id/_marker_span/_order_bit/_marker_position`).
- **Perceptual analyzer** (`bin/perceptual_analyzer.py`). `EndpointSumWhere` is the invertible span key $where = phase(start) + phase(end)$: the angle decodes the span center, the magnitude the span length. `MeronymicAnalyzer` analyzes a surface into terminals (compatibility mode reuses the word/byte tokenizer as the `boundary` op, so it matches the current lexer), writes durable spans to an `ObjectSubSpace`, exposes a fixed-capacity terminal-stream view, and reverse-synthesizes surface (`synthesize` exact replay; `synthesize_tree` from an operator-prefixed tree with `emit`). `soft_operator_compose` + `WholeSpace.operator_superposition` apply a soft operator distribution over the SS operation codebook (one-hot $\rightarrow$ the typed grammar; spread $\rightarrow$ the superposition that discriminates A AND B from A OR B).
- **PS-to-SS binding.** `WholeSpace.resolve_ps_terminal(ps_id)` emits `null_sem()` before a binding exists, counts exposures, and promotes a repeated terminal into a fresh SS row.

2026-05-29 deltas:

- `unreduce()` passes the space-role-local Basis (Codebook) to binary `GrammarLayer` reverses as `basis=space_role_basis` (replacing the prior raw-W form). `UnionLayer.reverse` /

`IntersectionLayer.reverse` and now `ConjunctionLayer.reverse` / `DisjunctionLayer.reverse` extract `W = basis.getW()` internally and dispatch to `Ops.disjunctionReverse` / `Ops.conjunctionReverse` — the codebook recommender recovers the operand pair exactly on a discrete vocabulary (the serial XOR reconstruction path); without a basis they keep the lossy `(parent, parent)` fallback. The genuinely non-invertible predicate folds (`isEqual` / `isPart` / `exist`) declare `invertible = False` and are not invertible by design (a truth/predicate value does not retain its operands). Layers that don't accept the `basis` kwarg yet are handled by a `TypeError` fallback. No back-ref is stored on the layer.

- `MetaLayer` was renamed to `SymbolizeLayer` (no semantic change).
- Word-mode parse appends a `\x00` null sentinel after the words slab for explicit end-of-sequence on the forward path.
- See `doc/old/2026-05-29-clean-stack-stm-basis-arg-radixlayer.md`.

Current Parser Surface

`SymbolSpace.compose()` and `SymbolSpace.generate()` are the public parser entry points. There is no longer a backend selector: the signal router (`LanguageLayer`) is the single canonical parser.

The `<parserBackend>` knob is RETIRED (Stage 3, 2026-05-27). The CKY chart and the STM shift/reduce parsers it used to select have been deleted in favour of the signal router. The retired `parserBackend` values (`chart` / `stm` / `parallel`) no longer exist.

`<routerKind>` is also RETIRED alongside the chart; the `signal` behaviour it once selected is now unconditional. The retired `<parserBackend>`, `<routerKind>`, `<chartTau>`, `<chartTopK>`, and `<chartNoiseEps>` elements raise a loud `ValueError` at config load if a `<SymbolSpace>` config still sets them (`Language._assert_retired_chart_knobs_absent`); see `data/model.xsd`.

Retired XML Knobs

These XML fields are no longer read by the runtime and should not appear in `data/*.xml`:

- `SymbolSpace.useGrammar`
- `SymbolSpace.chartCompose`
- `SymbolSpace.softChartCompose`
- `bivectorOutput`

Grammar mode is derived from the loaded grammar: default-only unary `pi` / `sigma` rules derive `useGrammar == "none"` internally; any non-default operator rule derives `useGrammar == "all"`.

SS-analysis vs CS-execution. `SymbolSubSpace.compose` is the SS-side *analysis* stage (it selects the per-space hard rule dict `current_rules`); the CS-side *execution* (applying lift / lower / union / intersection / swap / quantize / not to the concept tensors) runs in `ConceptualSpace.forward` and the per-space `SyntacticLayer` cursors. This split is a clean code boundary **only on the default-only path**: on the full-router path `LanguageLayer.compose` does both selection and tensor reduction, and the per-space cursors are deliberately bypassed (not `_grammar_is_default_only`). See STM.md Section 5 for the accurate, audited account.

Grammar

The `Grammar` is the singleton `Grammar` instance. Rules are loaded from XML `<SymbolSpace><language><grammar>` blocks or from a configured grammar CFG. A `RuleDef` stores:

`(space_role, canonical, arity, method_name, lhs, rhs_symbols)`

The grammar accepts explicit conceptual-order suffixes:

NP3

NP4

VP1
MP1
S4
S5

The intended current style is explicit ordered rules, for example:

```
S4 = lift(NP3, VP1)
S5 = lift(NP4, MP1)
```

Here all NPs share base category NP, but participate at different conceptual orders. The parser carries stack frames as (`category`, `order`), so NP3 and NP4 both have `category == "NP"` but orders 3 and 4.

Kleene order forms such as NP* are still parsed for compatibility, but they are not required by the current grammar style.

Knowledge Artifacts

`embed.build_knowledge_section(grammar)` creates the parser knowledge section:

- `taxonomy`: base category refs plus explicit ordered refs.
- `reference_codebook`: scalar prototypes and `order` per ref.
- `typed_indexes`: `refs_by_category` and `refs_by_order`.
- `grammar.rule_order_signatures`: serialized rule typing.

For explicit ordered grammar, the taxonomy stores ordered refs as children of their base category:

```
NP
|-- NP3
`-- NP4
```

`KnowledgeView.category_of_ref(ref_id)` returns the base category (NP), while `KnowledgeView.order_of_ref(ref_id)` returns the conceptual order (3, 4, etc.).

`SymbolSpace.category_codebook` has been retired. The live category embedding is:

```
SymbolSpace.category_embedding: nn.Embedding
```

STM Shift/Reduce

The STM backend uses:

- `ConceptualSpace.stm_typed`: a `TypedStack` carrying payload, category, order, and ref id.
- `STMDriver`: shift/reduce coordinator.
- `RuleScorer`: MLP over top stack payloads.
- `embed.admissibility_mask`: hard category/order mask over grammar rule signatures.

SHIFT snaps each input vector to the nearest live reference, then pushes:

```
payload, category, order, ref_id
```

REDUCE computes a typed admissibility mask, applies it to rule logits, softmaxes over admissible rules, chooses the argmax for the current greedy path, and records the chosen score/probability.

Typed admissibility is exact for explicit rules:

```
S4 = lift(NP3, VP1)
```

fires only when the stack top matches:

```
left:  category NP, order 3
right: category VP, order 1
```

Syntax And SVO

`BasicModel.write_syntax_tree()` (in `bin/Models.py`) emits an XML syntax-tree dump per batch row when configured with a syntax output path. It reads the chart’s per-row derivation trace and the symbolic codebook’s per-atom category tags; see the function `docstring` for the exact format. SVO extraction is performed by the STM path on the canonical subject lift over a transitive verb-phrase derivation; the object role is the NP operand inside that verb-phrase derivation, not a separate semantic category.

Role-Collapsed Grammar and the Operator Codebook

The live role-only grammars replace part-of-speech categories with *operator roles*. Instead of a fixed NP / VP / S taxonomy, each operator contributes its own argument and result roles, named `<op>_I1`, `<op>_I2` (inputs) and `<op>_O1` (output). The rule `isEqual` therefore exposes `isEqual_I1`, `isEqual_I2`, `isEqual_O1`, and the grammatical “category” of a span is just the set of operator roles it can fill. Dimensionality is recovered from participation (`bin/participation.py`) rather than declared up front: symbols that fill the same roles cluster into the same category.

Starts are scoped per space (`_configure_starts`):

- `PartSpace.start` is the universal whole-input role U — the analyzer begins from the entire surface and decomposes it.
- `WholeSpace.start` is the set of operator output roles (`isEqual_O1`, `isPart_O1`, `exist_O1`, ...) — parsing begins from what an operator can *produce*.

The **operator codebook** is a second codebook on `WholeSpace`, separate from the whole-percept codebook:

- `_operation_vectors`: operator name $\rightarrow$ identity vector.
- `_operation_positions`: operator name $\rightarrow$ codebook position.

It stores one prototype per operator (`isEqual`, `isPart`, `exist`, `conjunction`, `disjunction`, ...) and is kept CPU-explicit so the host-side identity lookup never mixes with an ambient MPS / CUDA default device.

Soft Operator Superposition

`operator_superposition(query_vec)` is a softmax over the cosine similarity between a query and every operator vector. `soft_operator_compose(dist, left, right)` is the distribution-weighted sum of each operator’s `compose`, evaluated per operator *arity* — unary operators such as `exist` are called as `compose(a)`, binary operators as `compose(a, b)`. A one-hot distribution reduces exactly to the typed grammar; a spread distribution superposes operators, and that superposition is the mechanism that discriminates `A AND B` from `A OR B` while a single layer is chosen.

`shape_operators(examples, op_names, steps, lr, seed)` trains the operator vectors by backpropagating an MSE between the superposed prediction and the supervised result, writing the shaped vectors back into `_operation_vectors`. The op set may span the whole grammar, so the shaper filters to layers of arity 1 or 2 and dispatches `compose` by arity.

Status (D1 gate — met). Role-collapse does not swap declared POS for another single-label POS system; it replaces declared shared categories with operator-local participation categories that a word may fill several of (overlaps are expected). The D1 gate is therefore not a single-label POS recovery test — it asks whether those participation patterns are *structured enough to drive a learned collapse* into the smaller mutually-exclusive category set the live parser needs. They are: `participation.learned_collapse` proposes merges by participation similarity and accepts only those that keep every grammar rule distinguishable (`collapse_conflicts == 0`); on the transitional `complete.grammar` it compacts 43 context-unique symbols into 14 mutually-exclusive categories with zero parser conflicts (`test/test_d1_pos_recovery_gate.py`). The exact substitutability congruence is trivial there (every symbol is context-unique), so “recovers the grammar” means the parser’s rule decisions survive the collapse, not exact rule regeneration. With the gate met, the former standalone role-collapse file has been absorbed into `complete.grammar`, which is the broad live role-only grammar

used by `MentalModel.xml`. The part operator is unified there: the grammar declares the single relative op `isPart` and the query/assert split is recovered by dispatch (`isPart` $\rightarrow$ `queryPart` in a query context). The operator codebook, soft superposition, and participation clustering are live and tested. See the status blocks in `doc/old/2026-06-02-unified-subsymbolic-analyzer-and-role-collapsed-grammar.md`.

Participation Categories as the Chooser’s Syntactic-Category Context

(Live path. The category source is learned from role participation during perception, attached to the *MetaSymbol*, and threaded into the placement chooser as grammatical context.)

Terminology note (per `doc/old/2026-06-21-terminology-percepts-concepts-symbols.md`): the CS part whole relation table is the **Concept codebook** (concepts), not a “symbol table”; `WholeSpace` holds **whole-percepts**; only `SymbolSpace` / *MetaSymbol* things are **symbols**. Code identifiers (e.g. `subspace.what`, `_sym_*`) are unchanged here — that is a separate code pass.

A word’s grammatical **category** is its frequency of participation across the operator roles above (`<op>_I<n>` inputs / `<op>_O1` output), **learned from perception** (analysis of input), not from generation. No part-of-speech label is declared or needed: if *cat* fills ADV’s operand role and LIFT’s argument role with roughly equal frequency, *collapsing* that frequency profile is what makes it a **noun** — the model knows the category by its role distribution, never by a name. This is exactly what `participation.learned_collapse` formalizes (merge by participation similarity, keeping every rule distinguishable).

This role-participation profile is the **primary determinant of syntax** — what a constituent *is* (its category) governs what it can combine with more than its surface content does. So it is precisely the context the placement chooser (`MLPTransformChooser`, the soft route’s scorer — see Soft-superposition route) needs when scoring “should this pair reduce, and with which operator?”: the chooser must see the **category of the value already sitting in each slot**, not only the candidate operator’s own output. The design realizes this as a small **Category codebook keyed by the MetaSymbol**, learned online by E/M from perception:

- **The MetaSymbol unifies word and object (it already exists, live).** The Concept codebook stores `concept` $\rightarrow$ `code`; a *MetaSymbol* is the exception — one symbol that holds *two* codes, the **word code** and the **object code**, a deliberate equivalence class asserting *this word this object*. This is the live **META node** in the `WholeSpace` taxonomy: `WholeSpace.insert_meta` allocates one SS-codebook row tagged “meta” whose two taxonomy children are the PS object position and the SS word position, minted during **perception** by the autobind hook (`ConceptualSpace._maybe_autobind_meta` at the sentence boundary). Because the category attaches to the *MetaSymbol*, the syntactic signal learned from the *word’s* role participation directly shapes the *object’s* category, and vice versa.
- **A small Category codebook, not a permanent per-word count table.** The VQ lives directly in role-participation space: `K` `n_roles` (~30) initial centroids, one seeded from each labelled role (`<op>_I<n>` inputs + `<op>_O1` outputs). Unlearned *MetaSymbols* have only a bounded temporary row in `MetaSymbolCategoryLearner`; learned *MetaSymbols* keep just `MetaSymbol` $\rightarrow$ `category_id`.
- **E/M learned from perception, with emergent collapse.** Each analysis route contributes a sparse role vector to the *MetaSymbol* that occupied the terminal position. The pending row accumulates that evidence until mass, confidence, margin, and short-term stability thresholds are met. Then the *MetaSymbol* commits to one VQ centroid and the pending row is discarded. Starting from one centroid per role and letting unused centroids decay, **effective K shrinks as role-use profiles pull centroids together** — the online realization of `participation.learned_collapse`, where “noun” emerges without a label.
- **Feeds the per-slot category to the chooser.** The chooser conditions each slot on the **role vector of the committed centroid**; while a word is still unsettled, the pending row supplies a temporary role context. The gather path is `percept id` $\rightarrow$ `taxonomy parent` (*MetaSymbol*) $\rightarrow$ `committed category` or `pending evidence` $\rightarrow$ `role vector`. `MLPTransformChooser` receives the vector as a feature block; anchor-dot/default routing uses the same vector as a labelled-role score prior.

This splits into two phases: (1) learn the category codebook from perception in the autobind hook (no change to the layer forwards — reads the stashed analysis route); (2) thread the per-slot category through

compose / score_binary / score_unary into the chooser feat (the larger change — the layer forwards carry only [B,N,D] today, with no per-slot symbol identity).

Status (implemented behind <categoryCodebook>, default true). WholeSpace.enable_category_codebook builds the role-space VQ (codebook_retire=False) + _category_role[K, n_roles], enumerated from compute_role_vocabulary; it is requested at build and lazily enabled on the first perception forward. LanguageLayer._collect_round0_role_obs stashes the first binary space-role’s round-0 reduces (op_I1/op_I2 per operand), and the autobind hook (_maybe_autobind_meta) feeds those observations to MetaSymbolCategoryLearner. The learner owns the pending per-MetaSymbol role rows, commits stable symbols into WholeSpace._category_assign, and drops the pending row. Structured grammar layers use the resulting role context for all transform choosers: MLP as an input feature, anchor-dot/default as a labelled-role score prior. The current round-0/first-space-role observation is parallel-mode-correct; serial (<serial>true</serial>) attribution is approximate. The old WholeSpace.category_codebook was retired 2026-05-20 and is gone; the dormant declared-POS tables (category_embedding, category_logits/category_ids, the order-taxonomy admissibility gate) are superseded and slated for follow-up retirement, not reuse.

Shared Weighted-Deduction Framework

PartSpace analysis and WholeSpace parsing are two readings of the *same* item graph under semiring-weighted dynamic programming (weighted deduction). The graph is scored once; the semiring chosen selects the quantity:

- **sum-product** gives soft inside / forward marginals (every viable rule keeps positive mass, and gradient flows through log Z);
- **max-plus** gives the single Viterbi route used for the hard forward value.

Both spaces share the same direction-agnostic primitives in bin/Language.py:

- binary_tiling_viterbi — max-plus best route.
- binary_tiling_soft_dp — sum-product marginals (reduce_marginal_op, logZ).

The WholeSpace reducer (BinaryStructuredReductionLayer) scores copy / reduce items over rule columns; the PartSpace analyzer (MeronymicRouter, bin/perceptual_analyzer.py) scores merge evidence over perceptual atoms. Because the soft marginals are retained even when the Viterbi route hardens to one rule, two tied reduce rules both keep positive mass and both receive gradient — the property pinned by test/test_signal_router_layer.py::test_layer_keeps_soft_superposition_over_reduce_rules.

Soft-superposition route (the <learning> two-pass)

The straight-through forward above is the **default** (and the byte-identical basin when <architecture><learning> is off): the forward value commits to the Viterbi route while the soft marginals carry the gradient. Under the two-pass <learning> mode the structured layers instead run a **pure sum-product superposition at a temperature**, and the chooser sits in the gradient path *directly* — no argmax, no .detach(), no straight-through.

A single scalar superposition_temperature $t \in [0,1]$ drives it (BinaryStructuredReductionLayer / UnaryStructuredLayer):

- the route scores are scaled by superposition_scale(t) = $1 - t$ before the soft DP / softmax, so $t = 0$ is the chooser’s own (sharp) softmax and $t = 1$ collapses the scores to uniform (flat, maximally exploratory);
- the forward value is binary_tiling_soft_dp’s temperature-scaled marginals (binary) or the temperature-scaled action softmax (unary). The op blend is the pure op_soft posterior, not the hardened one-hot.

binary_tiling_viterbi is still computed, but only to read off the routing masks for the tree (below) — that read-off is outside the gradient path in both modes. When superposition_temperature is unset (None)

every branch falls back to the legacy straight-through, object-for-object, so the default basin is unchanged.

Two passes as two trials. With `<learning>` on, `runEpoch` runs each sentence through `runBatch` *twice*, as two independent forward/loss/backward trials (no `loss_A + loss_B`, no shared graph):

1. **pass A** at $t = 0$ (sharp / deterministic) — recorded into the batch error like any ordinary step;
2. **pass B** at $t = \text{<exploreTemperature>}$ (default 0.5, flatter) — an exploration trial whose value is trimmed from the per-batch error and does **not** increment the batch count.

Pass B is temperature sampling, not a separate objective: a flatter route lets the chooser escape a local commitment that pass A’s sharp argmax would otherwise lock in, and because the chooser is differentiable in both passes the exploration gradient updates the same anchors. `BasicModel._set_superposition_temperature(t)` walks the router’s `_unary_layers` / `_binary_layers` (and the STM reducer) to stamp t ; `runBatch` sets it before the forward and resets to `None` in a `finally`.

Reading a tree. A hard derivation is always recoverable on demand: pin $t = 0$, run a temp-0 analysis, and read the argmax routing trace (`action_kind` / `action_op` / the copy / reduce masks) that `BasicModel.write_syntax_tree()` already walks. Running several sentences through a temp-0 pass yields one tree each.

This is the standard semiring-parsing pattern (Goodman 1999, *Semiring Parsing*; the SCFG inside-outside line of Lari and Young 1990; weighted logic programming / parsing transformations, Eisner and Blatz 2006; and neural grammar induction, Kim, Dyer, and Rush 2019). See the plan’s section 5.6 for the full mapping.

Mereology

Single-page reference for the mereological grammar: parthood as the fundamental operation, the five mereological relations, and the `ImpenetrableLayer` regularizer.

Part/whole spaces (2026-06-12). The perceptual side now names the duality directly: `PartSpace` (bottom-up synthesis over atoms) and `WholeSpace` (top-down analysis over unity), both subclassing the thin `PerceptualSpace` base. At the corpus callosum, objects are analysed and synthesized by sending them back to `PerceptualSpace` — wholes get split, parts get chunked. For the object/reference (sign/symbol) vocabulary, see `Spaces.md` and `Philosophy.md`.

Meronomy reconciliation (2026-06-11; MeronomySpec wins — plan §2). Four reframings of this page’s claims:

1. **part() is the graded retrieval surrogate, not *the* parthood relation.** Exact dominance (`Ops.partOf` — elementwise reduced with `all`) is the semantic ground truth; the clipped cosine below is its retrieval-time score. `part()` stays total and form-level over everything — including symbol codes and binding-table-bound rows — and that totality is its cordon duty (spec §1, §7): geometric relations answer “looks like”; *asserted* parthood lives in the truth store and the towers’ registration, answers “is”, and is a different call. Reading one relation as the other is the category error.
2. **Fusion = elementwise max stays as the order-theoretic join.** The learned `-fold` satisfies `max` (MeronomySpec §10.1) — join versus parametric whole-maker; the whole-maker can only over-cover the join, never under-cover it.
3. **The region partition (incl. exact equal == 0 underlap) is a scoring vocabulary only.** No abstraction classifier consumes it (withdrawn, spec §9): no adjacency exists over witness dimensions, so disconnection is meaningless there.
4. **Contiguity is sequential, never dimensional** (spec §2): it is owned by `Mereology.Contiguous` on the derivational axis; extents over witness dims are single lumps always.

2026-05-29 delta — binary GrammarLayer reverses take Basis, not W. `Ops._binary_op_recommend(result, W, op_name, ...)` (the mereology-guided recommender that walks `W` rows to find an operand pair such that $\text{op}(x_1, x_2) \approx \text{result}$) is still keyed on the raw `W` tensor at the low-level kernel signature. But the public `UnionLayer.reverse(parent, basis=None)` /

`IntersectionLayer.reverse(parent, basis=None)` now accept a `Codebook / Basis` object (typically `WholeSpace.subspace.what`) and extract `W = basis.getW()` internally before dispatching to `Ops.disjunctionReverse / Ops.conjunctionReverse`. The chart reverse / signal-router dispatch (`bin/Language.py::unreduce()`) passes `basis=space_role_basis` at the call site; no back-ref is stored on the layer. See `doc/old/2026-05-29-clean-stack-stm-basis-arg-radixlayer.md`.

Codebook IS the meronymic structure. The standalone `MereologicalTree` sidecar that formerly stored explicit parent / equality links was retired in favour of pure-geometric parthood on the `WholeSpace` bivector codebook. The grammar layers `PartLayer`, `IsEqualLayer`, `EqualLayer`, and `QueryLayer` operate directly on codebook bivector activations via clipped cosine projection — no separate adjacency table, no `<architecture><mereologicalTreeSize>` XML knob (silently ignored if present). Asserted meronymic relations are learned by training pulling the codebook geometry into the right configuration.

Parthood as Clipped Cosine Projection

Parthood is the single fundamental operation. Every other mereological relation is defined in terms of it.

Terminology (per `doc/old/2026-06-21-terminology-percepts-concepts-symbols.md`). Below, A, B are **percept** vectors (dimensionally-embedded codes in the `PartSpace/WholeSpace` codebooks); `part/whole/equal` are geometric relations over those percepts. “concept” is reserved for a `ConceptualSpace` relation tying one part-percept to one whole-percept (by reference), and “symbol” for a 0-D `SymbolSpace` reference to a concept.

For two percept vectors $A, B \in \mathbb{R}^D$:

$$\text{part}(A, B) = \frac{\max(0, A \cdot B)}{|A| |B|}$$

The clipped cosine projection lies in $[0, 1]$ and satisfies Boole’s contrapositive:

$$\text{part}(A, B) = \text{part}(-B, -A)$$

because $(-B) \cdot (-A) = A \cdot B$ and norms are sign-invariant.

Empty-operand contract. When $|A|$ or $|B|$ is near zero, `part` returns 1.0 — the empty set is part of everything.

The Full Suite

Every member of the suite is expressible through `part`:

Method	Formula	Interpretation
<code>part(A, B)</code>	$\max(0, A \cdot B) / (A B)$	A
<code>whole(A, B)</code>	<code>part(B, A)</code>	A contains B.
<code>equal(A, B)</code>	<code>part(A, B) · part(B, A)</code>	Mutual parthood, $[0, 1]$.
<code>overlap(A, B)</code>	$0 < \text{equal}(A, B) < 1$	Strictly-partial mutual parthood (indicator).
<code>underlap(A, B)</code>	<code>equal(A, B) = 0</code>	No mutual parthood (indicator).
<code>boundary(A, B)</code>	$\$$	$\text{part}(A, B) - \text{part}(B, A)$

Region Partition

`equal(A, B)` partitions into three disjoint regions:

Region	Condition
Underlap	<code>equal = 0</code>
Overlap	$0 < \text{equal} < 1$
Identity	<code>equal = 1</code>

Under clipped cosine, **part** is symmetric, so the three regions collapse to a single scalar classifier. For a future **Basis** with asymmetric **part**, the general five-relation table applies:

Five-relation table (general case, $\tau \approx 1$, $\epsilon \approx 0$):

Relation	Condition
<code>disjoint(a,b)</code>	$P[a, b] < \epsilon$ and $P[b, a] < \epsilon$
<code>part(a,b)</code>	$P[a, b] > \tau$ and $P[b, a] < \epsilon$
<code>part(b,a)</code>	$P[a, b] < \epsilon$ and $P[b, a] > \tau$
<code>equal(a,b)</code>	$P[a, b] > \tau$ and $P[b, a] > \tau$
<code>overlap(a,b)</code>	Both partial

Asymmetric subsumption. Since **part** is symmetric under clipped cosine, classical asymmetric subsumption is not encoded in raw magnitude. It is recovered **relationally** via figure / ground: compare `part(A, B)` against `part(A, ¬B)`. This is what makes Boole’s contrapositive hold exactly.

Mereological Fusion

Fusion of a truth set is the **least upper bound** (LUB / join) over stored truth vectors:

$$\text{fusion} = \max_i \text{truths}[i]$$

(elementwise max). In bivector space, the fusion vector names the top-right corner of the axis-aligned bounding hyperrectangle dominating every stored truth. Fusion is the geometric dual of luminosity: LUB (join) vs GLB (meet).

DoT is already baked into each stored truth (`stored[i] = activation_i * degree_i` in `TruthLayer.record`), so fusion is trust-weighted automatically.

See Logic.md section **Fusion** for full discussion and the leading-bivector / paired-index layout caveat.

ImpenetrableLayer: Five-Relations Regularizer

ImpenetrableLayer regularizes the WholeSpace whole-percept codebook. It classifies each ordered pair of codebook rows (i, j) into one of the five mereological relations (using **Basis.part**), and penalizes partial overlap when paired with a trust mismatch. The learned half of the Codebook Uniqueness Contract.

Penalty

$$\mathcal{L} = \text{overlap_weight} \cdot \frac{\sum_{i \neq j} \text{overlap}(i, j) \cdot |\text{trust}(i) - \text{trust}(j)|}{K(K-1)}$$

- variance_floor term

Overlap strength is damped to zero as the pair approaches **equal**:

$$\text{overlap_strength}(i, j) = \min(P[i, j], P[j, i]) \cdot (1 - \max(P[i, j], P[j, i])^k)$$

with $k = \text{equal_suppression}$ (default 4.0). Penalty is zero for **disjoint**, **equal**, and strict **part**, active only for **overlap**.

A separate **variance_floor** term guards against row collapse.

Trust

Per-row usage frequency, derived from VQ EMA counts:

$$\text{trust}(i) = \frac{\text{cluster_size}[i]}{\sum_j \text{cluster_size}[j]}$$

When VQ is absent, trust falls back to $\|cb[i]\| / \max_j \|cb[j]\|$.

Relative relations and tetralemma trust. A relative sentence in the **part** / **isEqual** predicate family is preserved at depth 3 as **[predicate, idea1, idea2]** and may be learned into the codebook as a ternary META edge (predicate as parent, the two ideas as children) via **WholeSpace.insert_relation**. An accepted relation carries a **tetralemma trust 4-tuple** (t, f, b, n) (TRUE / FALSE / BOTH / NEITHER, summing to 1) from the TruthSet posture, gated by the content-aware learn-score against **<truthCriterion>**. See STM.md Section 9.

Diagnostics

After **forward()**: **last_overlap_loss**, **last_variance**, **last_relation_counts** (dict summing to $K(K - 1)$).

Configuration

XML knobs (under **WholeSpace**):

Knob	Default	Meaning
overlapWeight	0.1	Weight of overlap $\times$ trust-diff penalty
varianceFloor	0.01	Minimum per-dim std
fullPartThreshold	0.9	τ : part score above this is “full part”
disjointThreshold	0.1	ϵ : part score below this is “disjoint”
equalSuppression	4.0	k : damping exponent near the equal corner

Order-raising (building the meronymic lattice)

Gated behind **<mereologyRaise>** (default off $\rightarrow$ byte-identical). Full design + code map: [doc/old/mereological-order-raising.r](#)

Terminology note. The cross-tower link that ties one part-percept to one whole-percept by reference is a **concept** (a `ConceptualSpace` relation); the `_sym_*/insert_meta` “META node” tables that hold these links are the **Concept codebook**. Prose below uses “concept” for that relation; the code identifiers (`insert_meta`, `_sym_*`) are left as-is pending the code-identifier pass.

The two towers stay **in-kind** — `PartSpace` *composes parts* → *parts*, `WholeSpace` *analyses wholes* → *wholes* — and the **concept (META node) is the cross-tower link**, associated with both an overlapping part-percept and whole-percept (`insert_meta(ps_pos, ss_pos)`). An object’s identity is **isomorphic**: -up meets -down at the object. At init there are only **atoms** (part-percepts in `PartSpace`) and the **universe** (the top whole-percept in `WholeSpace`); objects emerge from attention.

Three balancing forces keep a single coherent lattice and converge it:

1. **Link the tightest relation** — the largest part-percept the smallest whole-percept (via `.where` extent); skip a link a bigger-part/smaller-whole already subsumes; drop useless concepts.
2. **Too many parts → synthesize a higher-order part** (`maybe_raise_order`): when a whole accumulates more than `K_many` parts, mint a higher-order PART that subsumes them, with **abstraction order** one above its constituents (tracked via the ramsification table — order 0 = atom, 1 = basic category, ...) and explicit `part_chain` provenance (a higher-order part is abstract, so its `.where/.when` are discontinuous and the meronymy is tracked, not read off `.where`). Idempotent per whole.
3. **Only one part → Lewis’ Singleton → analyse the whole** (divides) or drop the spurious link.

Link surgery: `delete_meta / unlink_child` (invert `insert_meta`’s registrations); `ps_children_of_whole` counts a whole’s parts. The higher-order code is synthesized by `SigmaLayer.synthesize_over_set` (the M-way generalization of the binary `atanh-sum` fold) — or, in the “subsymbolic first” phase, the mean-combine of constituents, with `part_chain` as the source of truth. (First pass: the raise fires correctly when a whole has many parts; the live pid-keyed autobind binds 1 percept→1 concept, so similarity-based many-to-one binding + the prune-and-rebind of moot edges are noted follow-ups.)

The corpus callosum links the towers (part isa whole)

The full mechanism (see the spec): the **corpus callosum** in `ConceptualSpace` (the learned `[2N,N]` `self.callosum` glue) is what joins the towers into one meronymy. A part-percept A (`PartSpace`) and a whole-percept B (`WholeSpace`) have `.what` codes from *different* codebooks (incomparable), but their **.where is comparable** — so the callosum asks `WhereEncoding` whether A.`.where` is *part of* B.`.where`; when it is (and no greater part / lesser whole intervenes, per codebook activation) it forms the **concept A isa B** (token `isa` type; the type *names* the token). Co-occurrence at the same `.where/.when` drives the link, weakening when they dissociate — a **Hebbian** coupling between codebooks.

Since the 2026-06-16 redesign `.where` (and `.when`) is an **endpoint-sum bracket** `[start, end]` (angle = span center, magnitude = extent), so this *part-of* test is a direct read: `WhereEncoding.decode_span` recovers each code’s `(start, end)` and containment is `A.start B.start A.end B.end`; **contiguity** (adjacency / gap) between sibling parts is the same endpoint comparison. An instant snaps to zero extent, so atoms compare as points. (See `doc/Spaces.md` for the encoding.)

Design decision (2026-06-29, REVISIT) – the .where support of percepts vs symbols. Every PERCEPT (a part from `PartSpace` or a whole from `WholeSpace`) MUST carry a real `.where` (`nWhere > 0`): a definite `[start, end]` extent. A percept with no location is not admissible. A SYMBOL (a concept-level relation entry, or an SS reference) MAY instead either (a) carry `nWhere = 0` (no location of its own), or (b) carry a RECOMBINED `.where` equal to the bounding span of its constituents – `start` = the leftmost constituent’s `start`, `end` = the rightmost constituent’s `end` (the support of a symbol is the sum of the supports of its members). If the recombination is hard to implement, the fallback is to give every symbol a `.where` of “everywhere” (universal support). THE ANT-COLONY QUESTION: is the location of an ant colony IN each ant (distributed – every ant keeps its own `.where`), or is there a SINGLE location that subsumes all the ants’ locations (the bounding region)? The recombined bounding-span (b) takes the single-subsuming-extent

view; “everywhere” is its degenerate limit. This choice determines whether a concept’s **.where** is distributed over its parts or collapsed to one subsuming extent, and should be revisited.

Word object cannot be linked this way (too unlike; no convex set is specific enough), so a **second-order meta-object** is synthesized in PartSpace, **outside .where/.when**, fusing the word-code and object-code into the symbol used in serial communication (the MetaSymbol). *(Language update 2026-07-02, revised same day by the two-phase rework: in the ramsified CS this is a FIRST-order META-concept – the sec-4c ORDERED PAIR [whole = word-symbol, part = object-symbol], roles as positional slots of an ordered pair rather than containment claims; the typed read-out recovers (word, object) by INTERSECTION with the word-symbol class (meta_word_object), and the reference-table pairing remains the serial-mode access path. See doc/Architecture.md sec A.)* Because symbols are outliers, the part whole **concepts** live in a **two-code LUT** and the symbolic taxonomy is **relations over symbol indices** — which are the TruthLayer’s **absolute truths** (propositions) and **relative truths** (RelativeTruthStore), to be integrated. Taxonomy relations are also learned **explicitly from trusted language** (“cats are furry” → [cats] <= [furry]).

Granularity / the part-whole ratio is the correctness signal (and the MM_20M mean-collapse fix): *many parts → one whole* = under-analysed; *one part → one whole* = over-analysed. An incorrect ratio is the criterion to request **further synthesis (, e.g. radix) or analysis (, property tiling) within the problematic .where** — until words emerge as parts. MM_20M collapses because its byte chunker never climbs and its analyser never descends, so no word-granularity parts exist for XOR.

Design decision (2026-06-29, REVISIT) – locking in the “basic level” of analysis. The mereological analysis should converge on and STABILISE at Eleanor Rosch’s BASIC LEVEL – the granularity where the part-whole ratio is correct (neither many-parts-one-whole nor one-part-one-whole). The radix climb builds parts UPWARD (byte -> prefix -> word -> ...); left unbounded it over-climbs (a whole short sentence promotes to one percept). Two principled brakes hold it at the basic level: (a) the <basicLevelMinSize> / <basicLevelMaxSize> size bounds (for TEXT, a part is [0, 24] letters – “size in .what” = letter count); and (b) the part-whole-ratio convergence signal above (request more synthesis/analysis until the ratio is right). The ideal is to LOCK IN whatever basic level the corpus implies. As a pragmatic HEDGE – and because **word-level is essentially the basic level for a symbolic LLM** – the analyser may simply cut on whitespace (PS observes/promotes per WHITESPACE-cut word so the parts top out at word size; see PartSpace._embed_radix under _meronomy_words). The size bound (24) is too loose to separate a short sentence from its words on its own, so the word cut is the operative brake today; the full ratio-driven basic-level convergence is the to-revisit principled version.

Explicit concepts implicit subsymbolic representations (dual-coded, 2026-06-21)

Terminology note. This section formerly called the explicit part whole relation entry a “symbol.” Per the convention, that relation — one part-percept tied to one whole-percept by reference, held in the _sym_* tables — is a **concept** (the Concept codebook). The 0-D SymbolSpace **symbol** is a separate, intensional reference *to* such a concept and is not what the _sym_* entries are. Prose below says “concept”; the _sym_* code identifiers are left as-is pending the code-identifier pass.

The architecture carries **two coordinating structures** for the same content — an **explicit relational** tower (the concepts) and an **implicit subsymbolic** one — and order-raising is what keeps them in correspondence. We have *both*, not one or the other.

- **Explicit (concepts).** A concept is **built directly** (not folded): an index-value relation entry tying a PartSpace part-percept code to a WholeSpace whole-percept code, with self-reference (('sym', id) members) so concepts nest into the taxonomy. Concepts associate part-percepts and whole-percepts **in virtue of a common whole**, and **identity on insertion is same-part-OR-same-whole** (the same concept is reinforced, not duplicated). **.where varies trial-to-trial** and is *not* the identity or trigger key — it keeps only its spatial/meronomy role (extent, containment). The discrete _sym_* tables (the Concept codebook) are this structure. (This refines the earlier **.where/co-occurrence** framing above: the *trigger* to abstract is **many constituents under a common concept**, not a shared **.where**.) These _sym_* relations are the INDEX side; the concept’s subsymbolic content (the ConceptDim atom

+ its per-order sparse weight decomposition, populated at mint by `_populate_concept_weights`) is the representational side — see the next bullet.

- **Implicit (subsymbolic).** Each concept has a corresponding **learned vector** — a strictly-positive **ConceptDim atom** (a feature signature) stored in the CS concept dictionary (`similarity_codebook`, `softplus-rectified`). For a higher-order concept the *production* is no longer “ then quantize”; it is the **ramsified per-order sparse transform** (`cs_forward_content`, `bin/Spaces.py`): a per-order signed **sparse weight matrix** maps the concept’s source activations `[PS | WS | SS_0..SS_{k-1}]` to a signed concept activation $\mathbf{a}_k = \mathbf{W}_k @ \text{source}_k$ (`torch.sparse.mm`), and the concept code is that activation scaling its positive atom ($\mathbf{a}_k \cdot \text{softplus}(\text{atom})$; radial: magnitude = certainty, sign = present vs anti-present). The many→one abstraction is carried by the sparse WEIGHTS (which sources contribute, with what sign), not by a -fold + VQ snap; the gradient reaches the weights, the source activations, and the dictionary (forward-connected). **PerceptDim** and **ConceptDim** are decoupled — a concept is in a different vector space than its percepts, never a sum of percept vectors.

Coordination. Order-raising is the coupling between the two. When the explicit tower raises a higher-order part-percept/whole-percept (built directly, inserted on the **over-collected side** of the higher-order table — too-many-parts → higher-order part, too-many-wholes → higher-order whole), it **creates the corresponding implicit subsymbolic higher-order representation** via the per-order sparse transform above (a new order-k table over `[PS | WS | SS_0..SS_{k-1}]`). The explicit concept **indexes / names** the implicit vector; the implicit vector is the concept’s **subsymbolic content**. The two co-evolve and must agree: a directly-built concept-index paired with a quantized vector that is a valid point in the higher-order mereological codebook. Direct index + quantized representation are the two complementary halves of one concept.

This is the dual-coded (neuro-symbolic) core: discrete concept **relations over indices** in lock-step with continuous **quantized representations** — neither alone, but the two coordinating.

Why This Design

- **Parthood as projection.** One formula, Boole-contrapositive exact, continuous in $[0, 1]$. The old composite formula $\text{conjunction}(1 - \text{dist}(x, x \cap y), 1 - \text{dist}(y, x \cup y))$ mixed set-valued operands with a distance.
- **Overlap penalty (not antisymmetry).** Legacy `ImpenetrableLayer` penalized mutual parthood $P[i, j] \cdot P[j, i]$ directly, pushing *every* overlapping pair apart including synonyms. The five-relations design leaves `equal` pairs alone and only penalizes the ambiguous middle region, gated by trust mismatch.
- **Trust via VQ EMA.** Already tracked for VQ commit loss.
- **Parthood is preserved by Pi / Sigma.** Under `monotonic=true`, `Pi/Sigma` are restricted to $W \geq 0$. User-supplied truth-set bivectors live on the non-negative paired `[pos, neg]` cone, where componentwise $\leq$ is the parthood partial order. Each lift / lower preserves parthood pole-by-pole for that truth surface. The bivector layout keeps contradiction `[1, 1]` distinct from ignorance `[0, 0]` under positive matmul. See `Spaces.md` “Monotonicity of the lift / lower chain”.

Meronymic Analyzer (PartSpace)

Parthood is not only the mereological relation over percepts of the previous sections; it is also how **PartSpace** *analyzes* a surface into structure. The default perceptual chunking mode is `analyse` (`<PartSpace><chunking>analyse</chunking>`), which routes the input through the meronymic analyzer rather than a fixed lexicon or BPE table.

InputSpace hands the analyzer the unanalyzed surface — with `<lexer>raw</lexer>` it passes the whole `[B, 1, N]` byte buffer and lets **PartSpace** own tokenization. The analyzer’s default (and only enabled-by-default) operation is a *space-lexer*: a boundary detector that recognizes whitespace breaks, reproducing the current word-level behavior for verbal input. `raw` keeps the word-level codebook (unlike `byte`); it changes only the analyze surface, so the model is not switched to a byte embedding.

MeronymicRouter (`bin/perceptual_analyzer.py`) scores candidate merges by signed-neighborhood cosine against the perceptual codebook and routes with the shared `binary_tiling_viterbi` / `binary_tiling_soft_dp` primitives (see `Language.md`, *Shared Weighted-Deduction Framework*). In the cold state the codebook holds only the whole-input vector and the byte vectors, so the router decomposes the surface to byte terminals — the deliberately-failing initial state. As words are learned bottom-up (merge promotion), the same router reproduces word-level runs. `analyze_routed` reconstructs a byte-exact surface for replay (UTF-8 multi-byte characters are stored as raw bytes, not decoded per segment).

isPart as a Grammar Layer

IsPartLayer lifts parthood into the role-collapsed grammar as a binary operator (arity 2). Under a query model it dispatches to `queryPart` (`_dispatch_method_name_for_rule`), so `isPart` and `isEqual` are the relative operators recognized by `_relative_start_categories` (`_RELATIVE_OP_NAMES`). Its identity vector lives in the **WholeSpace** operator codebook and participates in the same soft superposition as the other operators.

Logic

Overview

This document defines the logic system at three levels:

1. **Subsymbolic (vector / field level)** — geometry in **ConceptualSpace**
2. **Symbolic (scalar level in [-1,1])** — order + polarity in **WholeSpace**
3. **Rationality** — propositional truth store built on both

Executable implementations of the subsymbolic and symbolic operators are the **Method** subclasses in `Language.md`.

Terminology (2026-06-21; see `doc/old/2026-06-21-terminology-percepts-concepts-symbols.md`).

One noun per-space: **percept** = a **PartSpace**/**WholeSpace** thing (dimensionally embedded, EXTENSIONAL; *part* = atom / bottom-up, *whole* = property/region / top-down are its two subtypes); **concept** = a **ConceptualSpace** relation tying one part-percept one whole-percept (the Concept codebook, diachronic); **symbol** = a **SymbolSpace** thing (0-D, INTENSIONAL, references a concept; the symbol codebook). The Section 2 scalar emission is the *symbol* that migrates to **SymbolSpace**, while **WholeSpace** itself holds whole-percepts. Code identifiers below are unchanged (a separate code pass renames them).

1. Subsymbolic Layer

Objects: vector sets (B, N, D) interpreted as RBF / luminosity fields.

Operators

- **Union:** `union(A, B) = concat(A, B)`. Combine sets.
- **Intersection:** Co-supported regions (RBF product / merge).
- **Negation (affirming):** `neg(x) = -x`. Antipodal opposition on hypersphere.

- **Non (non-affirming):** `Basis.non()` — bitonic returns zero (complete withdrawal); monotonic is `relu(x - threshold)`.

Meronomy reconciliation (2026-06-11; MeronomySpec §3 wins). The sign flip $-x$ is licensed only at the REFERENCE space-role — downstairs the sign is occupied (form content) or nonexistent (one-sided percepts); only at the reference space-role is it vacant of denotational duty and free to carry polarity ($a = -1$ on A's row = present, certain denial; never negative mass in ground space). `non()`'s withdrawal reading is the prasajya move — lowering $|a|$ toward 0 without crossing — gated, never spontaneous. RadMin/RadMax below remain signed-scalar kernels at their own space-role; the odds embedding is retired from the meronymic path (legacy PiLayer keeps it for non-meronymic consumers).

- **Parthood** (fundamental): `Basis.part()` — clipped cosine projection:

$$\text{part}(x, y) = \frac{\max(0, x \cdot y)}{\|x\| \cdot \|y\|}$$

Satisfies Boole's contrapositive trivially (dot product and norms sign-invariant under joint negation). The full suite (**whole**, **equal**, **overlap**, **underlap**, **boundary**) composes through **part**.

2. Symbolization

For $X \in (B, N, D)$:

$$s(X) = 2 \cdot \text{mean}(\|x_i\|) - 1 \in [-1, 1]$$

Interpretation: $+1$ = strong presence, 0 = neutral, -1 = absence.

3. Symbolic Layer (Scalars in [-1,1])

Basis supports **monotonic** (plain min/max over positive-cone values, including user-supplied truth-set bivectors) and **bitonic** (sign-aware). Monotonic forms here; bitonic forms (RadMin, RadMax) in Section 7. See Spaces.md “Monotonicity of the lift / lower chain”.

Let $a, b \in [-1, 1]$.

Operator	Bitonic	Monotonic
Negation (affirming)	$\text{neg}(a) = -a$	—
Non (non-affirming)	$\text{non}(a) = 0$	$\text{non}(a) = \text{relu}(a - \tau)$ (learnable τ)
Union	RadMax (Section 7)	$a \cup b = \max(a, b)$
Intersection	RadMin (Section 7)	$a \cap b = \min(a, b)$

Parthood as Projection

Parthood is the **fundamental mereological operation**. For $A, B \in \mathbb{R}^D$:

$$\text{part}(A, B) = \frac{\max(0, A \cdot B)}{\|A\| \cdot \|B\|} \in [0, 1]$$

Boole's contrapositive $\text{part}(A, B) = \text{part}(-B, -A)$ holds trivially.

The full suite composes through **part**:

Method	Formula
whole (A, B)	part (B, A)
equal (A, B)	$\text{part}(A, B) \cdot \text{part}(B, A)$
overlap (A, B)	$0 < \text{equal}(A, B) < 1$
underlap (A, B)	$\text{equal}(A, B) = 0$
boundary (A, B)	\$

Under clipped cosine, **part** is symmetric, so **equal** reduces to part^2 . Asymmetric classical subsumption is recovered **relationally** via figure/ground: compare $\text{part}(A, B)$ against $\text{part}(A, \neg B)$. See Mereology.md.

For ternary scalar values (monotonic projection):

$\text{part}(a, b)$	+1	0	-1
+1	1	1	0
0	1	1	1
-1	0	1	1

Zero is vacuously part of everything (empty-set convention).

4. Key Insight

- Subsymbolic = geometry
- Symbolic = order + polarity
- Symbolization = norm projection

Cleanly separates representation (vectors) from logic (scalars).

5. Rationality

Rationality is the propositional logic layer, built on SS grammar rules **part**(S, S) and **equals**(S, S).

Truth Statements

A **truth statement** is any assertion processed through the full pipeline. The **TruthLayer** — owned by **SymbolSpace**, reached as `self.wordSpace.truth_layer` — stores activations **scaled by** DegreeOfTruth:

$$\text{stored} = \text{activation} \times \text{degree}$$

Degree	Stored Vector	Effect
+1	full activation	attractor
$0 < d < +1$	scaled activation	weak attractor
0	zero vector (inert)	prunable
$-1 < d < 0$	scaled, negated	weak disperser
-1	negated activation	disperser

TruthLayer (SymbolSpace.truth_layer)

TruthLayer is instantiated by `SymbolSpace.__init__`. `WholeSpace.forward` records activations into the TruthLayer governed by the single continuous `<truthCriterion>` bar (0 = record every activation, 1 = record none): a per-cell activation is recorded when its clamped magnitude clears `truthCriterion`. The same knob also gates learned-concept acceptance in `ConceptualSpace` (relations admitted into the Concept codebook), so one continuous knob governs all truth learning. The binary `<accumulateTruth>` / `<truthMinMagnitude>` switches are **retired**. Recording now fires wherever `WholeSpace.forward` runs — both normal training and the `store_truths` gold-ingestion epoch (which drops `truthCriterion` to 0 to capture every provided gold truth, then restores it). See STM.md Section 9 and Params.md:

- **record(activation, degree)** — store `activation * degree`.
- **query(activation)** — find the closest stored truth by cosine similarity.
- **field(concepts)** — project all stored truths into `ConceptualSpace` as a scalar field over concept vectors.

Integration with the meronymy (doc/old/merological-order-raising.md). The absolute truth set (propositions, e.g. `cat <= [animal() & orange() & object()]`) and the `RelativeTruthStore` (relations between two ideas, e.g. `[cats] <= [furry]`) are **the same two structures** the corpus callosum needs for the concept-relation meronymy: a two-code part whole LUT (absolute) plus relations over concept indices (relative) — both live in the Concept codebook. These are to be **integrated** — the part whole LUT *is* the absolute table; the concept-taxonomy relations *are* the relative table. Taxonomy relations are also learned **explicitly from trusted language** (“cats are furry” → the relation `[cats] <= [furry]`, admitted when trusted).

Truth Field

$$f(c) = \frac{1}{n} \sum_{i=1}^n c \cdot t_i$$

where $t_i = \text{activation}_i \times \text{degree}_i$. Two kinds of regions:

- **Attractors** ($f \rightarrow +1$): concept vectors near +DoT truths.
- **Dispersers** ($f \rightarrow -1$): concept vectors near -DoT truths.

Consonance and Dissonance

1. **Internal consonance** — stored truths should be mutually consistent.
2. **External consonance** — incoming statements are evaluated against the truth field.

Propositional Structure

Both parthood and equality are SS operations over symbolic activations:

- **part(S, S)** — containment by clipped-cosine projection. Grammar applies it as `score * B`, scaling the whole by parthood degree.
- **equals(S, S)** — identity as mutual parthood; delegates to `Basis.equal`.

Together they define a partial order over symbolic activations.

Truth Accumulation Pipeline

Truth entries arrive as (text, DoT) pairs. `store_truths()` stages all texts on `TheData` and runs a standard inference epoch via `runEpoch(split="runtime")`. During forward, `WholeSpace.forward()` records each raw symbolic activation (degree 1.0). After the epoch, each activation is scaled by its DoT:

$$t_i = \text{activation}_i \times \text{DoT}_i$$

Two-phase design (encode all, then scale) ensures all truths are encoded in the same pipeline context before DoT modulation.

6. Consistency and Verification

Internal Consistency (within a TruthSet)

Leave-one-out check: for each t_j , evaluate the field produced by the rest of the set at c_j :

$$d_j = f_{\setminus j}(c_j) = \frac{1}{n-1} \sum_{i \neq j} c_j \cdot t_i$$

- $d_j > 0$: consonant; $d_j \approx 0$: independent; $d_j < 0$: dissonant.

A disperser (DoT < 0) near an attractor (DoT > 0) naturally produces negative field values, flagging the contradiction.

Verification (incoming statement against stored truth)

`verify(statement)` processes the statement, then queries the field:

$$v = f(c_a) = \frac{1}{n} \sum_{i=1}^n c_a \cdot t_i$$

v	Interpretation
$v \rightarrow +1$	strong support
$v \approx 0$	no opinion
$v \rightarrow -1$	strong contradiction

`wordSpace.truth_layer.query()` returns the single closest truth and its degree for point lookups.

Logical Entailment (symbolic closure)

Subsymbolic field checks detect *geometric* consistency but not *logical* entailments. Example: `part(A, B) = +1` and `part(B, C) = +1` entail `part(A, C)`, but A and C could be geometrically distant. A symbolic closure layer would:

1. **Extract propositions** by decomposing stored SS activations into relation + two operands.
2. **Compute closure** via transitivity:
 - $\text{part}(A, B) \wedge \text{part}(B, C) \Rightarrow \text{part}(A, C)$
 - $\text{equals}(A, B) \wedge \text{equals}(B, C) \Rightarrow \text{equals}(A, C)$
 - $\text{equals}(A, B) \wedge \text{part}(A, C) \Rightarrow \text{part}(B, C)$

Each entailed proposition inherits DoT = min of the premise DoTs.

3. **Inject entailments** as new symbolic activations recorded in the TruthLayer.

Worth implementing for reasoning rather than retrieval. The geometric field alone is sufficient for soft consistency checks and nearest-truth lookups.

7. Radial Operators for Hypersphere Ternary Logic

Ternary Logic Operators			
+1 = true 0 = unknown -1 = false			
NOT (neg(a) = -a)		NON (non(a) = 0)	
a	¬a	a	non(a)
+1	-1	+1	0
0	0	0	0
-1	+1	-1	0

Intersection (agree → min a,b)			
a \ b	+1	0	-1
+1	+1	0	0
0	0	0	0
-1	0	0	-1

Union (agree → max a,b ; 0 absorbs)			
a \ b	+1	0	-1
+1	+1	+1	0
0	+1	0	-1
-1	0	-1	-1

PART (mereological, range [0, 1])			
a \ b	+1	0	-1
+1	+1	0	0
0	+1	+1	+1
-1	0	0	+1


| PART(a,b): 1 when a ⊆ b; 0 is part of everything; opposite signs → 0 | | | |

RadMin, RadMax, NOT, and NON operate over signed magnitude space [-1, +1]. Sign = direction of assertion; magnitude = strength; zero = **unknown**. +1 = True; 0 = Unknown (no assertion); -1 = False. Magnitude ordering: $x \subset y$ iff $|x| < |y|$.

Op	Behavior	Ternary truth table
RadMin (AND)	Sign disagree → 0; else $\mathrm{sign}(x) \cdot \min(x , y)$	x
RadMax (OR)	Sign disagree → 0; else $\mathrm{sign}(x) \cdot \max(x , y)$	x
NOT	-x. Involutive; preserves magnitude; NOT(0) = 0.	input → -input
NON	0 for all x. Absorbing; idempotent; withdraws assertion (vs. NOT which inverts).	input → 0

Both RadMin and RadMax are commutative and collapse to zero on sign disagreement.

Luminosity



Meronomy reconciliation (2026-06-11; MeronomySpec §3 rev b wins). There are TWO measures, split as follows. The **truth-set luminosity** — what `Mereology.Luminosity(truth_layer=...)` / `TruthLayer.luminosity` computes — is now the *catuşkoţi* coverage measure over the stored codes: per conceptual dimension the signed references split into true/false pole coverage (T_k , F_k) (elementwise max of $\text{relu}(\pm \text{row})$), and

$$\text{luminosity} = \text{mean}_k[(T_k - F_k) - \min(T_k, F_k)] \quad \text{in } [-1, 1]$$

— total area weighted by sign, minus the regions where the sign differs (contradictory evidence). It is computed **on the codes** (no decode pullback) and is order-independent. The meet/GLB form below survives only as the **orthogonalization criterion** (`TruthLayer._luminosity_without`) and **darkness()**'s mirror — it is no longer the truth-set measure. (The diagram shows the meet form.)

The meet-form coherence scalar (orthogonalization criterion):

$$||\text{relu}(\min(\text{truths}))||$$

Element-wise min computes the conjunction (point where all truths agree); relu removes negative dimensions; L2 norm gives brightness.

- **High luminosity:** truths are coherent and mutually reinforcing.
- **Low luminosity:** truths are sparse, contradictory, or incoherent.

Two roles in the model:

1. **Top-down bias:** concept input scaled by luminosity each conceptual iteration: `concept_input * (1 + truthBiasScale * luminosity)`.
2. **Loss modification:** low luminosity increases training loss.

Fusion

Fusion is the **mereological least upper bound** of the stored truth set:

```
fusion = max_i truths[i]
```

In bivector space (paired-index [p0, n0, p1, n1, ...]), fusion names the top-right corner of the smallest hyperrectangle dominating every stored truth componentwise. Geometric dual of luminosity:

- **Luminosity** = `||relu(min(truths))||` — GLB / meet, scalar coherence.
- **Fusion** = `max(truths)` — LUB / join, vector coverage.

Trust (DoT) is already baked into stored truths, so fusion is trust-weighted automatically.

Layout caveat. TruthLayer’s paired-index slicing ([..., 0::2] for positive poles, [..., 1::2] for negative) assumes a *repeated* bivector layout. The WholeSpace codebook uses a *leading* bivector plus positional trailers: [pos, neg, where..., when...]. Callers feeding WholeSpace activations into `luminosity()` must slice the leading 2 dims first (`acts[..., :2]`); see `truth_loss` call-site in `bin/Spaces.py`.

Supporting measures

- **Consistency** (`isConsistent()`): folds truths into a union via successive `Basis.disjunction()`. In bitonic mode, conflicting +/- assertions cancel.
- **Grounding** (`ground()`): finds the minimal subset entailing a query.
- **TruthLoss**: additive penalty for contradictions, measured by union norm reduction. See Reasoning.md.
- **Derive**: pairwise mereological inference via `part()`; generalized by `extrapolate()` to all two-argument methods.

Semantic Summary

Operator	Type	Role	Behavior on Contradiction
RadMin	Binary	Conjunction (AND)	Collapses to zero
RadMax	Binary	Disjunction (OR)	Collapses to zero
NOT	Unary	Sign-flipping negation	N/A
NON	Unary	Non-affirming negation	N/A

8. Cross-references

The level-crossing operations (`lift`, `lower`, `intersection`, `union`), `Pi/Sigma` layer ownership, witness-set / slab-lattice geometry, and `Ops.*` namespace inventory have moved out of this document.

- **Ops.* reference**: Language.md.
- **PiLayer / SigmaLayer ownership** (`forwardPi/reversePi/forwardSigma/ reverseSigma` aliases): Spaces.md.
- **Mereological suite** (`part`, `whole`, `equal`, `overlap`, `underlap`, `boundary`, `copart`): Mereology.md.
- **Tetralemma bivector layout** (TRUE/FALSE/BOTH/NEITHER corners): Spaces.md and Philosophy.md.

Reasoning System

2026-05-29 delta: the chart / signal-router reverse path now passes the space-role-local Basis (`subspace.what`) to binary GrammarLayer reverses as `basis=space_role_basis`. The mereology-guided recommender (`Ops._binary_op_recommend`) walks the Codebook's W rows to find operand pairs (x_1, x_2) such that $\text{op}(x_1, x_2) \approx \text{parent}$. Under `<codebook>none</codebook>` on WholeSpace the recommender has no rows to walk and falls back to the lossy (`parent, parent`) pseudo-inverse — degrading the reasoning loop's structural recovery on multi-stage chart parses. See `doc/old/2026-05-29-clean-stack-stm-basis-arg-radixlayer.md`.

Four methods on `BaseModel` for truth-aware inference, a bidirectional reasoning loop, and a grammar learning mode. Builds on the TruthLayer infrastructure (`Logic.md`) and grammar composition (`Language.md`).

Partitioned Symbol Space

Terminology (percept / concept / symbol). Throughout this doc “symbol”/“symbolic” denotes the genuine SymbolSpace space-role — the 0-D, non-dimensionally-embedded references emitted as `symbolSum` — not the ConceptualSpace part whole relation table (those are *concepts*) and not the dimensionally-embedded perceptual content of PartSpace/WholeSpace (those are *percepts*: part-percepts and whole-percepts). See `doc/old/2026-06-21-terminology-percepts-concepts-symbols.md`.

The symbol dimension is statically partitioned across conceptual orders using geometric decay. Each order writes only to its slice of `symbolSum`, while reading the full vector as feedback.

```
order 0:  [0,      D//2)      <- 1/2 of symbol_dim
order 1:  [D//2,   3D//4)     <- 1/4
order 2:  [3D//4,  7D//8)     <- 1/8
...
last order: remainder of D
```

Makes the symbol partition **self-describing**: position reveals conceptual order. Truth methods use `_activation_order()` to determine a query's order by finding the partition with the highest energy. Partition boundaries are precomputed once at model creation via `BasicModel._order_partitions`.

Reasoning Methods

`isConsistent()` -> dict

Analyzes the TruthSet for internal consistency by folding all stored truths into a single summary via successive `Basis.disjunction()`. In bitonic mode, conflicting +/- assertions on the same dimension cancel to zero. Returns `{'consistent': bool, 'score': float, 'sites': tensor, 'union_vector': tensor}`.

`ground(activation, threshold=0.6)` -> dict

Finds the minimal subset of the TruthSet entailing a query activation. Uses `_activation_order()` to filter truths by partition. Falls back to `TruthLayer.derive()` for indirect derivation. Returns `{'grounded': bool, 'basis': [indices], 'trace': [...], 'confidence': float}`.

`isTrue(activation)` -> float

Grounds a proposition and returns a scalar Degree of Truth in $[-1, 1]$. Positive = true, negative = false, zero = unknown. Delegates to `ground()`.

`extrapolate(seed_indices, max_new, attenuation)` -> dict

Generalizes `TruthLayer.derive()` to all two-argument grammar methods (union, intersection, equals, part). For each pair of stored truths, applies every eligible method and accepts results that preserve or

increase luminosity. Accepted truths recorded at `attenuation * min(DoT_i, DoT_j)`. Returns `{'added': [indices], 'rejected': [(i, j, rule, delta_lum), ...]}`.

Meronomy reconciliation (2026-06-11). The gate’s role is unchanged, but the gated quantity is now the MeronomySpec §3 rev-b measure: `TruthLayer.luminosity` = the *catuskoṭi* coverage measure over the codes, `mean_k[(T_k - F_k) - min(T_k, F_k)]` — signed area minus conflict — order-independent, no decode pullback. The same applies to the multiplicative luminosity modulation under “TruthLoss” below.

TruthLoss

Additive loss penalty for false propositions, via `<TruthLoss>` in `model.xml` (default 0.0 = disabled).

Measures the **union norm reduction** when a proposition is included in the `TruthSet` union via `Basis.disjunction()`:

```
truth_union = disjunction(all stored truths)
extended    = disjunction(truth_union, new_proposition)
penalty      = max(0, ||truth_union|| - ||extended||)
```

Case	Effect
Agreeing proposition	Preserves/extends union dims → no penalty
Unknown proposition (zero dims)	Passes through → no penalty
Contradicting proposition	Cancels conflicting dims → positive penalty

DoT weighting is implicit: stored vectors carry DoT in magnitude, so contradicting a high-DoT truth causes a larger norm drop.

TruthLoss is **additive** and coexists with **multiplicative** luminosity modulation (`totalLoss * (1 + lum_weight * (1 - luminosity))`).

Bidirectional Reasoning Loop

`BasicModel.reason(givens, target, direction, max_steps):`

- **Forward** (`givens` → `conclusion`): Encode `givens` into `TruthSet`, extrapolate new truths each step, check `isTrue(target)` until DoT exceeds threshold or `max_steps` is reached.
- **Reverse** (`target` → `grounding`): Encode `target`, call `ground()` to find minimal basis, extrapolate if insufficient.

Luminosity non-decrease is the validity certificate.

Grammar Learning

`BasicModel.grammar_learning_step()` learns grammar weights from a symbolic reconstruction objective:

1. Forward pass produces `symbolSum`.
2. Reverse pass reconstructs input.
3. Re-encode reconstruction to `symbolSum_hat`.
4. `Loss = ||symbolSum_hat - symbolSum||^2` (symbolic, not conceptual).
5. Optional luminosity validity penalty for rules decreasing luminosity.

Paraphrase-invariance holds because semantically similar sentences snap to nearby codebook entries.

Parser And Conceptual Order

Grammar mode is derived from the loaded grammar block. Default-only unary `pi` / `sigma` rules take the fast path; non-default operator rules enable grammar-directed parser dispatch.

`subsymbolicOrder` controls the number of $P \rightarrow C \rightarrow S$ stages and the symbol partition geometry. Higher-order symbols write to later partitions, so truth grounding, consistency, and extrapolation can respect conceptual order.

The parser backend is no longer selectable. Stage 3 of the substrate refactor (2026-05-27) retired the CKY chart and STM shift-reduce parsers; the signal router (`LanguageLayer`) is the single canonical parser. The former `SymbolSpace.parserBackend` / `routerKind` knobs (along with `chartTau`, `chartTopK`, `chartNoiseEps`) are **RETIRED** — setting any of them in a config raises a loud `ValueError` at load time (see `Language._assert_retired_chart_knobs_absent`).

Explicit ordered grammar is preferred:

```
S4 = lift(NP3, VP1)
S5 = lift(NP4, MP1)
```

Here all NPs share base category NP; the suffix gives the conceptual order. Lift and lower are the only syntactic operations that change argument/return order.

Configuration

Parameter	Location	Default	Description
<TruthLoss>	<training>	0.0	Additive truth-loss weight
<subsymbolicOrder>	<architecture>	1	Percept $\rightarrow$ Concept $\rightarrow$ Symbol it
<parserBackend>	<SymbolSpace>	—	RETIRED (Stage 3, 2026-05-27)
truthCriterion	<architecture> / <ConceptualSpace> / <WholeSpace>	1.0	Single continuous truth bar (0-1)
intraLossWeight	<training>	0.1	In-STM next-idea loss $\mathcal{L}_{\text{intra}}$
interLossWeight	<training>	0.1	Inter-sentence next-end-state
routerWireSerial	<architecture>	both	Per-word router-fire gating on
ltmCapacity	<SymbolSpace>	1024	LTM chain capacity (InterSe)

The relative-vs-absolute end-state machinery, the content-aware learn-score gate, and the tetralemma trust 4-tuple carried on accepted relative META edges are documented in `STM.md` Section 9.

Contemplative Awareness Methods

Four methods on `BaseModel` characterizing stages of contemplative awareness as spatial/computational properties. `Contiguous()`, `Continuous()`, and `Peaceful()` are implemented (each returns a measure in $[-1, +1]$; see `bin/Mereology.py`). `Peaceful()` reads the `TruthLayer` and returns `valence-symmetry` $\times$ `luminosity-uniformity` (balanced affirming/denying pole masses $\times$ uniformly-held per-proposition magnitude; 0.0 when no truths are stored). `Done()` remains a stub that raises `NotImplementedError` — it defines the target characterization, not the implementation.

Method	Stage	Characterization
<code>Contiguous()</code>	One-Pointedness (Shamatha / FA)	Single connected, convex region in PartSpace; contiguous span in WholeSpace
<code>Continuous()</code>	Simplicity (Continuity / OA)	Concept states flow continuously; Jacobian of forward map is bounded
<code>Peaceful()</code>	One Taste (Emotional Symmetry)	<code>TruthLayer</code> luminosity uniformly high across stored propositions
<code>Done()</code>	Buddhahood (Non-Meditation)	Model is a fixed point of forward-reverse; reconstruction loss zero

Shamatha Speech is the target grammar mode for `Contiguous()`: complete DNF object grammar plus spatiotemporal contiguity. Every `conjunction` / `disjunction` over object parts must keep `where()` support

connected and `when()` support continuous. Differs from serial mode — may reduce over all active percepts at once; rejects scattered object fields, not multi-percept fields. See Language.md Section Shamatha Speech Mode and plans/2026-04-28-shamatha-speech-contiguity-handoff.md.

Testing

Unit tests in `basicmodel/test/test_reasoning.py` cover all methods without requiring a trained model. English-level tests (syllogisms, contrapositives, semantic equivalence) are `@pytest.mark.xfail` until word identity is learned through training.

Training

2026-05-29 deltas:

- **Embedding unit-ball normalization.** The training loop calls `Lexicon.normalize()` (unit-ball projection on the per-row vectors) right after `optimizer.step()` in `bin/Models.py::runBatch`. Keeps embedding vectors from drifting off the unit ball under JOINT / BACKPROP modes.
- **Seeded retries for MM_xor tests.** `test_learns_xor_signal` and `test_convergence` in `test/test_mm_xor.py` use for `seed` in `(42, 123, 7): torch.manual_seed(seed);` ... (the previously-dead `seed` loop variable in `test_convergence` is now actually consumed). Pass-if-any-attempt-converges semantics.
- **Reconstruction is unconditional from concepts.** The `<reconstruct>` element (and `reconstructEnum`) was RETIRED (A1, 2026-06-09); there is no knob to set or override. Whenever reconstruction fires it is always concepts-seeded from the terminal `ConceptualSpace` STM snapshot.
- **Supervised output loss restored.** `bin/Models.py::runBatch` computes MSE on the supervised output in addition to the reconstruction loss (`reconstructionScale`-weighted).
- See `doc/old/2026-05-29-clean-stack-stm-basis-arg-radixlayer.md`.

Overview

Two phases: **embedding pretraining** and **network training**. Embedding pretraining builds word vectors from a large corpus. Network training uses those vectors as input representation and learns to predict and reconstruct sentences.

The two phases can overlap: `<trainEmbedding>` controls whether and how embeddings continue to evolve during network training.

The staged objective curriculum

Network training climbs five objectives, from substrate-building reconstruction to deliberate question answering. The order is not arbitrary — it is the architecture’s **two learning rules** in their natural developmental order:

- **EMA / occurrence** moves the codebooks: concepts form by being *seen*, with no gradient and no goal. This is *prediction as substrate* — unconscious, statistical, **System 1**. You cannot ask “why did the empire fall?” until *empire* and *fall* are codebook rows, and those are EMA-built by exposure.
- **Gradient on a task error** moves the attention readouts and the relation store: the model learns *where to look* and *what relates to what* by being *tested* — directed, credit-assigned, **System 2**.

So “bootstrap with prediction, then answer questions” is just the EMA substrate-builder running first and the gradient deliberate-layer on top (the §6c sentence protocol’s *prime subsymbolically from what you see, then pump symbolically*). Reconstruction (stages 1–2) is the **gate**: a question can only be *answered* once an

idea can be turned back into words, so the grammar-operation inverses those stages need are the load-bearing prerequisite for the whole curriculum.

stage	objective	trains
1	reconstruct, keep syntax , fill the missing words	the lexical/leaf inverse + the codebook (word slot)
2	reconstruct with no syntax / no words (from the idea alone)	the full generative inverse (deliverable C) + abstr
3	predict the next sentence	the inter-sentence (discourse AR) predictor
4	answer a question by reasoning (no search)	relations, modus ponens (consequents()), the ca
5	answer a question by LTM search	global attention (the typed .where) + the soft-re

Plain next-token/next-sentence prediction *under-trains this architecture's* retrieval and relation machinery (the local window usually suffices, so the separate global attention gets no reward); that is why prediction is the **substrate** here, not the goal — stages 4–5 train the deliberate layer, and the stochastic exploration finally earns its keep in stage 5 (search has only a distal reward, so it must explore to break symmetry).

Full design (with per-stage machinery, losses, and prerequisites): `doc/old/training-stages.md`.

The grammar-operation inverses stages 1–2 require are enumerated in `doc/old/2026-06-19-grammar-inverses-handof`

Phase 1: Embedding Pretraining (make `train` / `embed.py train`)

Produces a static embedding artifact (e.g. `BasicModel.kv`). Output path from `<embeddingPath>`.

Pipeline

FineWeb-EDU parquet shards

- > Pass 1: stream documents, lex + parse, count words, build vocabulary
- > Pass 2: stream documents again, train SBOW per sentence, discard examples
- > Save `WordVectors` to `sentence.pt`

SBOW (Sentence Bag of Words)

Sentence-level variant of CBOW. Given a sentence of N words:

Leave-one-out centroid for word i :

$$c_i = \frac{1}{N-1} \sum_{j \neq i} v_j$$

Logits (unnormalised, for every word in vocabulary V):

$$z_i = W_o c_i + b_o$$

Softmax probability:

$$P(w_i | c_i) = \frac{\exp(z_{i,w_i})}{\sum_{k=1}^V \exp(z_{i,k})}$$

Loss (cross-entropy, averaged over all N positions):

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N \log P(w_i | c_i)$$

Gradient of the loss has two terms:

- **Attractive** (from position i where w_i is the target): pulls v_i toward directions that increase its probability.
- **Repulsive** (from positions $j \neq i$ where v_i is part of the centroid): pushes v_i away from centroids whose targets are other words.

The softmax partition function ensures all V words compete: any word too close to a centroid it doesn't belong to increases the denominator and is pushed away with force proportional to its softmax probability. This adaptive repulsion distributes vectors across the embedding space.

```
vecs = embedding(idx)           # [N, dim]
total = vecs.sum(dim=0)         # [dim]
centroids = (total - vecs) / (N - 1) # [N, dim] leave-one-out
logits = linear(centroids)      # [N, vocab] full softmax
loss = cross_entropy(logits, idx)
```

CBOW (Continuous Bag of Words)

CBOW predicts a single target word from its context. Context mean (with padding for variable-length contexts):

$$\bar{c}_i = \frac{1}{|C_i|} \sum_{j \in C_i} v_j$$

Loss applies the same negative-sampling objective per-word to the padded context mean rather than the leave-one-out centroid:

$$\mathcal{L}_{\text{CBOW}} = -\log \sigma(\bar{c}_i^\top v_{w_i}) - \sum_{k=1}^K \log \sigma(-\bar{c}_i^\top v_{w_k^-})$$

where K is the number of negative samples and w_k^- are sampled words.

Two-Pass Architecture

- **Pass 1 (vocabulary):** Stream all documents, count every word. Words meeting `min_count` are promoted; model (embedding + linear head) allocated once at final vocab size.
- **Pass 2 (training):** Stream documents again. Per sentence, run one SBOW step and discard examples. Multiple epochs re-stream the data.

Configuration (Makefile variables)

Variable	Default	Description
BASIC_SHARDS	1	FineWeb-EDU parquet shards
BASIC_MAX_DOCS	10000	Max documents to process
BASIC_VEC_SIZE	100	Embedding dimensionality
BASIC_EPOCHS	10	Training epochs
BASIC_MIN_COUNT	10	Min word count for vocabulary inclusion
BASIC_BATCH_SIZE	8	Batch size (unused by SBOW; per-sentence)

Memory Considerations

Dominant cost is the linear output head: `vocab_size * vector_size` parameters plus optimizer state. With 200K vocab and 100 dims:

- Embedding: $200\text{K} \times 100 \times 4 \text{ bytes} = \sim 80\text{MB}$
- Linear head: $100 \times 200\text{K} \times 4 \text{ bytes} = \sim 80\text{MB}$
- Adam optimizer: $\sim 320\text{MB}$ total

Full softmax computes $(N, \text{vocab_size})$ logits per sentence. Training runs on CPU by default (`BASICMODEL_DEVICE=cpu`) to avoid MPS/GPU memory limits.

Phase 2: Network Training (`BasicModel.py`)

The network learns to predict and reconstruct sentences using pretrained embeddings.

Within-sentence training objective (IR-only)

Within-sentence training is **always IR** (masked-LM at the subsymbolic (PS)). `create_ir_mask` replaces a `mask_rate` fraction of WHAT positions with `NULL_PERCEPT` and snapshots the pre-mask event on `_ir_pre_mask_input`; `runBatch` computes `MSE(perceptualSpace at masked positions, _ir_pre_mask_input at masked positions)` — no head loss, no reverse pipeline. The legacy `<maskedPrediction>` knob and the AR / ARUS / ARIR modes were retired 2026-05-14; sentence-level AR moved to `InterSentenceLayer` (see `doc/Architecture.md` Section “Sentence-level AR (`InterSentenceLayer`)”).

Knob	Effect
<code>maskRate</code>	Bernoulli mask probability at the subsymbolic (PS) (BERT default 0.15)
<code>reconstructionScale</code> <code><reconstruct></code>	Blend weight between output and reconstruction loss; <code>total = (1 - r)*output + r*recon</code> . Legacy RETIRED (A1, 2026-06-09; <code>reconstructEnum</code> removed). There is no longer a target space-role knob.

`<trainEmbedding>`: Embedding Update Mode

Controls whether and how embeddings evolve during network training. When enabled (CBOW, SBOW, or BOTH), implements EM-like alternation:

- **E-step (network)**: Forward + masked prediction + backward + optimizer.
- **M-step (CBOW/SBOW)**: Run one embedding update on the same sentence.

Value	Embedding	Model layers	Description
NONE	Frozen	Trained	Only model layers train
CBOW	CBOW (padded context, own optimizer)	Trained	Model layers train separately
SBOW	SBOW (centroid, own optimizer)	Trained	Faster variant
BACKPROP	Backprop only	Trained	Codebook trained purely via model loss
BOTH	SBOW post-batch	Trained	Two optimizers
JOINT	Single backward	Trained	Single optimizer: combined model + SBOW loss

Gradient Flow Through Codebook

`<trainEmbedding>` determines whether the codebook weight is **detached** during forward/reverse and whether it appears in the main optimizer:

<code>trainEmbedding</code>	<code>detach()</code>	In main optimizer	Method
NONE	Yes	No	Frozen
CBOW	Yes	No	CBOW (own optimizer)
SBOW	Yes	No	SBOW (own optimizer)
BACKPROP	No	Yes	Model loss only
BOTH	No	Yes	SBOW + backprop
JOINT	No	Yes	Single combined loss

- **NONE** is recommended for constrained hardware (Apple MPS <16GB).
- **CBOW/SBOW** maintains clean EM separation.
- **BACKPROP** is the simplest trainable mode. Embedding shaped entirely by the task; best for small vocabularies.
- **BOTH risks gradient interference** — reconstruction pulls embeddings apart (distinguishability); SBOW pulls co-occurring words together. These forces can conflict because they use separate optimizers.

JOINT: Single Combined Loss

Avoids EM alternation by computing one combined loss before one backward:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{model}} + \lambda \cdot \mathcal{L}_{\text{SBOW}}$$

where λ is `trainEmbeddingRatio` (default 0.1).

SBOW loss uses the same negative-sampling objective:

$$\mathcal{L}_{\text{SBOW}} = -\frac{1}{N} \sum_{i=1}^N \left[\log \sigma(c_i^\top v_{w_i}) + \sum_{k=1}^K \log \sigma(-c_i^\top v_{w_k^-}) \right]$$

Advantages. One optimizer, one momentum buffer; gradient coherence; no post-batch step.

Trade-off. Both objectives share LR and Adam state. If loss surfaces have very different curvature, λ may need tuning.

Why MSE over Embeddings (not Cross-Entropy)

- “cat” prediction when target is “kitten” incurs less loss than “democracy” — embedding space captures semantic similarity.
- Output dim is the embedding dim (~100) vs. vocab size (~200K).

Training Loop

Data flows through `SentenceStreamDataset` wrapped in `DataLoader`. The ordered training list is split into $B = \text{batchSize}$ contiguous slabs of length $L = \text{len(split)} // B$; at step t , row b is item $b * L + t$. Temporal context is coherent across steps. No per-epoch global shuffle.

For each B-wide batch:

1. `BasicModel.Start(inputData)` cascades reset through every Space.
2. `InputSpace.forward()` lexes/embeds once into $[B, N, D]$ (left-aligned, right-padded to N). No K axis, no cursor unfold.
3. `create_ir_mask` replaces a `mask_rate` fraction of WHAT positions with `NULL_PERCEPT`; pre-mask event stashed on `_ir_pre_mask_input` and the position mask on `_ir_mask_positions`.
4. `_forward_body` runs T stages on B rows.

5. `_forward_head` produces `[B, N, predDim]` (a side channel — IR loss is computed at the subsymbolic (PS), not at the head).
6. `runBatch` computes loss via `TheError.add`:
 - `reconstruction` (subsymbolic (PS)): MSE between post-body `perceptualSpace` and `_ir_pre_mask_input` at masked positions.
 - `reconstruction_reverse` (concepts-seeded, unconditional): the reverse pass is always seeded from the terminal `ConceptualSpace` STM snapshot (the `<reconstruct>` enum was retired), weighted by `reconstructionScale`.
 - `arma` (sentence-level): `InterSentenceLayer.observe(s_t)` MSE between the ARMA(p, q) prediction and the current sentence rep, weighted by `armaScale`.
7. One `backward()` + `optimizer.step()` per `DataLoader` yield.
8. Embedding training (CBOW/SBOW/BOTH) runs once per batch.

SBOW vs CBOW

Property	CBOW	SBOW	Masked Prediction (Phase 2)
Targets per sentence	1 (pick one word)	N (every word)	N (one per masked pos)
Context	Other N-1 words	LOO centroid of N-1	All unmasked (+ future truncati
Positive updates	1 per step	N per step	N per step
Repulsive force	vocab-1 implicit via softmax	$N \times (\text{vocab}-1)$	Implicit via MSE
Signal density	Low	High	High
Loss	Cross-entropy over vocab	Cross-entropy over vocab	MSE over embeddings
Updates embeddings	Yes (own optimizer)	Yes (own optimizer)	Only for BACKPROP/BOTH/JOINT
Used in (<code><trainEmbedding></code>)	CBOW	SBOW, BOTH, JOINT	BACKPROP, BOTH, JOINT

SBOW provides richer signal — every word acts as both context (for others) and target (predicted from others). The full softmax pushes vectors to distribute across the hypersphere.

Three-File Architecture

Artifact	Example	Contents	Updated by
XML config	<code>data/BasicModel.xml</code>	Architecture, hyperparameters, paths	Hand-edited
Embedding	<code>data/BasicModel.kv</code>	Word vectors ($V \times d$)	Phase 1
Weights	<code>data/BasicModel.ckpt</code>	Model layer parameters	Phase 2

Checkpoint excludes embedding parameters (`wv._vectors`). Enables:

- Retrain embeddings without touching model weights (`--force-embeddings`)
 - Retrain model with frozen codebook (`<trainEmbedding>NONE`)
 - Swap codebooks between models sharing the same architecture
-

Embedding Space Geometry

The full softmax provides a distributional guarantee that random negative sampling cannot match. For predicting w_i from centroid c_i :

$$\mathcal{L}_i = -\log \frac{\exp(z_{i,w_i})}{\sum_{j=1}^V \exp(z_{i,j})}$$

The gradient with respect to the logit for any non-target word $k \neq w_i$:

$$\frac{\partial \mathcal{L}_i}{\partial z_{i,k}} = P(k \mid c_i)$$

Words with high softmax probability (too close to the centroid) receive the strongest repulsive gradient; words already far away contribute almost nothing. This naturally spreads vectors across the hypersphere.

At equilibrium, no word can move closer to any centroid without increasing the loss. Random negatives, by contrast, mostly sample words already far away on a high-dimensional hypersphere and contribute little useful gradient.

Profiling

--profile wraps Phase 2 in cProfile:

```
make train_micro_remote # add --profile via train.py
# Or:
python bin/train.py --model data/BasicModel.xml --profile --max-docs 500
```

Output: a .prof file in output/profiles/ plus a top-30 summary to stdout. View with `snakeviz`:

```
pip install snakeviz
snakeviz output/profiles/train_20260319_111441.prof
```

For live profiling, `py-spy` can attach by PID (requires root on macOS):

```
pip install py-spy
sudo py-spy record -o profile.svg --pid <PID> --duration 30
```

Ergodic Exploration via Adaptive Bias-Variance Control

BasicModel decouples **what** the network learns (weights W) from **how much it explores** (a scalar α governing the bias-variance tradeoff). Two algorithms operate in tandem:

Component	Algorithm	Role
Weights W	Standard Adam	Gradient descent on the loss landscape
Tradeoff α	Gradient Energy Sensor	Measures loss-surface curvature to set exploration

α is **not trained by gradient descent**. It is a *sensor* reading the gradient energy through the temperature parameter and converting it into an exploration policy.

1. The Ergodic Weights

Each layer uses an effective weight matrix sampled from a locally specified ergodic distribution:

$$W_{\text{eff}} = \alpha \cdot W + (1 - \alpha) \cdot \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, I)$$

W stores the learned structure that shapes the distribution as α rises. Derived quantities:

$$\text{bias} = \alpha, \quad \text{var} = 1 - \alpha$$

- $\alpha = 0$: Pure exploration. $W_{\text{eff}} = \varepsilon$. Random noise.
- $\alpha = 1$: Pure exploitation. $W_{\text{eff}} = W$. Learned weights only.
- $0 < \alpha < 1$: Interpolation.

Training begins at $\alpha = 0$ and transitions autonomously toward exploitation as the loss landscape flattens.

2. Why Standard Adam Fails on α

The temperature $T = 1 - \alpha$ receives gradients via:

$$\frac{\partial \mathcal{L}}{\partial T} = \sum_{i,j} \frac{\partial \mathcal{L}}{\partial (W_{\text{eff}})_{ij}} \cdot \varepsilon_{ij}$$

Because ε is re-sampled every forward pass, this gradient is **zero-mean**:

$$\mathbb{E}\left[\frac{\partial \mathcal{L}}{\partial T}\right] = \sum_{i,j} \frac{\partial \mathcal{L}}{\partial (W_{\text{eff}})_{ij}} \cdot \mathbb{E}[\varepsilon_{ij}] = 0$$

Standard Adam's first moment kills the signal: $m_t \approx 0 \Rightarrow m_t / \sqrt{v_t} \approx 0$. Adam produces no update.

3. The Gradient Energy Sensor

3.1 Modified Second-Moment Estimator

Discard the first moment; use only the second moment as **output** rather than normalizer:

$$v_t = \beta \cdot v_{t-1} + (1 - \beta) \cdot g_t^2$$

where $g_t = \partial \mathcal{L} / \partial T$ and $\beta = 0.999$.

3.2 Bias Correction

$$\hat{v}_t = \frac{v_t}{1 - \beta^t}$$

3.3 What $\hat{v}_t$ Measures

The second moment of a zero-mean random variable is its variance:

$$\mathbb{E}[g_t^2] = \text{Var}(g_t) = \sum_{i,j} \left(\frac{\partial \mathcal{L}}{\partial (W_{\text{eff}})_{ij}} \right)^2 = \left\| \frac{\partial \mathcal{L}}{\partial W_{\text{eff}}} \right\|_F^2$$

The **squared Frobenius norm** of the loss gradient — the **gradient energy**.

Gradient energy	Meaning	Desired behavior
High $\hat{v}_t$	Loss sensitive to weight changes	Keep exploring (low α)

Gradient energy	Meaning	Desired behavior
Low $\hat{v}_t$	Loss surface is flat	Exploit (high α)

3.4 The Alpha Update Rule

$$\alpha_t = \frac{1}{1 + \tau \cdot \sqrt{\hat{v}_t}}$$

where τ (`global_temp`) is a sensitivity knob.

Properties:

- $\sqrt{\hat{v}_t} \rightarrow 0 \Rightarrow \alpha_t \rightarrow 1$ (exploit)
- $\sqrt{\hat{v}_t} \rightarrow \infty \Rightarrow \alpha_t \rightarrow 0$ (explore)
- $\alpha_t \in (0, 1)$ always
- τ controls transition speed

After computing α_t : $\text{bias}_t = \alpha_t$, $\text{var}_t = 1 - \alpha_t$.

3.5 Layer-Local Certainty (Per-Neuron Sigma)

Each ergodic layer maintains per-neuron **sigma** — running variance of its gradient energy, tracked via Welford’s algorithm in `observe_sigma()`. Sigma drives per-neuron bias and var:

$$\text{var}_j = \frac{\sigma_j}{\sigma_j + \kappa}, \quad \text{bias}_j = 1 - \text{var}_j$$

Low sigma (consistent gradient, found a minimum) yields high bias, low var. High sigma (unstable gradient) yields low bias, high var.

`set_sigma(sigma)` controls κ :

- `sigma=1`: encourage exploration (low κ)
- `sigma=0`: suppress exploration (high κ , $\text{var} \approx 0$)

4. Comparison: Adam vs. Gradient Energy Sensor

Property	Adam (for W)	Sensor (for α)
Type	Optimizer	Sensor (measurement)
First moment m_t	Yes — tracks direction	No — direction is zero-mean
Second moment v_t	Yes — normalizes step size	Yes — is the output
Role of $\sqrt{v_t}$	Denominator (adaptive LR)	Numerator (energy estimate)
Update rule	$W \leftarrow W - \eta \cdot m / \sqrt{v}$	$\alpha \leftarrow 1 / (1 + \tau \sqrt{\hat{v}})$
Descent?	Yes	No — maps energy to policy
Zero-mean safe?	No — produces $0 / \sqrt{v} = 0$	Yes — by design

5. Derivation: Why Drop m_t

Let $g_t = \nabla_T \mathcal{L}$. Since noise ε is i.i.d. each step, $\mathbb{E}[g_t] = 0$. The first moment EMA $m_t = \beta_1 m_{t-1} + (1 - \beta_1)g_t$ gives $\mathbb{E}[m_t] = \beta_1 \mathbb{E}[m_{t-1}]$. By induction $\mathbb{E}[m_t] = \beta_1^t \mathbb{E}[m_0] = 0$ — exponentially to zero, carrying **no information**.

The second moment converges to a meaningful quantity:

$$\mathbb{E}[v_t] \rightarrow \mathbb{E}[g_t^2] = \|\nabla_{W_{\text{eff}}} \mathcal{L}\|_F^2$$

The gradient energy — exactly what we need.

6. Dropout via Bias

bias and var support dropout:

$$\text{bias}_t^{(\text{drop})} = \begin{cases} 0 & \text{with probability } p \\ \alpha_t & \text{with probability } 1 - p \end{cases}$$

When bias is dropped to zero, the layer acts as pure noise regardless of α . The temperature $T = 1 - \alpha$ remains unchanged so noise contribution is consistent.

Dropout is a schedule on bias: ordinary Bernoulli dropout is the binary case. Annealed dropout corresponds to increasing the expected bias over training. Nonzero temperature bakes regularization directly into the model — every forward pass perturbs the effective weights, discouraging brittle co-adaptation.

7. The `global_temp` Sensitivity Knob

τ (`global_temp`) scales responsiveness:

τ	Effect
$\tau = 0$	$\alpha = 1$ always — pure exploitation
$\tau \ll 1$	Quickly converges to $\alpha \approx 1$
$\tau = 1$	Balanced — α tracks gradient energy on natural scale
$\tau \gg 1$	Maintains high exploration even with moderate gradients

τ can be scheduled (e.g., annealed from high to low) for a coarse exploration-to-exploitation curriculum.

8. Initialization and Bootstrap

Problem. If initialized at $\alpha = 1$ (pure exploitation), the noise term vanishes and $\partial \mathcal{L} / \partial T = 0$. The sensor would read zero energy and keep $\alpha = 1$ forever — an **exploitation trap**.

Solution. Initialize at $\alpha = 0$.

At $\alpha = 0$:

- $W_{\text{eff}} = \varepsilon$ — pure noise carries gradients through the network
- Temperature gradient is nonzero: $g_t = \sum_{ij} (\partial \mathcal{L} / \partial \varepsilon_{ij}) \cdot \varepsilon_{ij} \neq 0$
- Sensor measures initial gradient energy and begins tuning α upward

- As W learns useful structure, energy decreases and α rises

Zero-initializing W is safe and often cleaner than random initialization. Symmetry is broken by the sampled noise; learned weights don't contribute at $\alpha = 0$ anyway. Starting from $W = 0$ makes the exploratory regime honest.

9. Layer Architecture

All ergodic layers follow the effective-weight pattern. When `ergodic=True`, noise buffers are registered and weights are zero-initialized.

ErgodicLayer Base Class

`ErgodicLayer` (`Model.py`) provides per-neuron `bias` and `var` buffers driven by gradient variance (`observe_sigma()` / `sigma_to_ergodic()`). Subclasses with child `ErgodicLayers` must forward `set_sigma()` and `paramUpdate()`.

InvertibleLinearLayer — factor-level noise injection

LDU factorisation $W = L \cdot D_{\text{embed}} \cdot U$ (see `Architecture.md`). Exact inverse via triangular solves; no SVD.

When `ergodic=True`, noise is injected at the factor level to preserve exact invertibility at all temperatures:

$$\begin{aligned} L_{\text{eff}} &= I + \text{strict_lower}(\text{raw_L} + t \cdot \text{noise_L}) \\ U_{\text{eff}} &= I + \text{strict_upper}(\text{raw_U} + t \cdot \text{noise_U}) \\ d_{\text{eff}} &= b \cdot d_{\text{clamped}} + t \cdot \text{noise_d} \end{aligned}$$

W_{eff} stays in LDU form, so its exact inverse is always available — no approximation regardless of noise level. Compare to matrix-level blending $W_{\text{eff}} = b \cdot W + t \cdot N$, which destroys LDU structure:

$$W_{\text{eff}} \cdot W_{\text{eff}}^{-1} \approx (b^2 + t^2)I + bt(WN^{-1} + NW^{-1}) \neq I$$

Error grows with temperature — exactly the regime where exploration is most active. Factor-level injection avoids this.

`noise_d` is sampled with magnitude in `[eps, 1]` and random sign, so it's always invertible; combined with the `stable=True` clamp on the deterministic diagonal, `d_eff` stays bounded away from zero.

`SigmaLayer(invertible=True)` and `PiLayer(invertible=True)` use `InvertibleLinearLayer(ergodic=True)` internally. `bias` and `var` from the parent `ErgodicLayer` forward as `bias=` / `temp=` arguments, so the gradient-energy sensor governs noise injection in the inner linear primitive.

10. Training Loop Integration

```
optimizer = torch.optim.Adam(model.getParameters(), lr=lr) # Adam for W
loss = model(x, y) # forward W_eff = bias*W + var*noise
loss.backward() # Adam: dL/dW; sensor: dL/dT
optimizer.step()
model.paramUpdate() # sensor updates alpha
```

11. Summary of the Full Algorithm

Initialize: $\alpha_0 = 0$, $v_0 = 0$, $t = 0$, $\beta = 0.999$.

Each step:

1. Sample $\varepsilon \sim \mathcal{N}(0, I)$.
2. $W_{\text{eff}} = \alpha_t W + (1 - \alpha_t)\varepsilon$.
3. Forward: $\hat{y} = f(x; W_{\text{eff}})$; loss $\mathcal{L}(\hat{y}, y)$.
4. Backward: $\nabla_W \mathcal{L}$ and $g_t = \nabla_T \mathcal{L}$.
5. **Adam** updates W .
6. **Sensor** updates α :

$$v_t \leftarrow \beta v_{t-1} + (1 - \beta)g_t^2, \quad \hat{v}_t = \frac{v_t}{1 - \beta^t}, \quad \text{var}_j = \frac{\hat{v}_j}{\hat{v}_j + \kappa}, \quad \text{bias}_j = 1 - \text{var}_j$$

7. Zero temperature gradient: $g_t \leftarrow 0$.

Machine Minds

Alec Rogers, September 24, 2025

Abstract

Problem

Society depends on machine minds, but we lack a good theory of them. While human and machine minds functionally overlap, their design goals differ, and treating a machine mind as a human mind has significant consequences. We should have a different theory of mind for machine minds, informed by how those minds are designed and trained.

Background

Anthropomorphizing machines helps us understand them relative to human minds. This essay focuses on three questions: what are machine minds, what do machine minds feel, and what do machine minds know.

Solution

Three measures make a machine mind more humanistic: weight ergodicity, network invertibility, and output certainty. Weight ergodicity treats weights as samples from a random process rather than fixed constants. Network invertibility makes the architecture bidirectional and partially invertible. Output certainty means knowing that it does not know, rather than only the probability of one answer versus another. All models tested use a 20-neuron, single hidden layer on MNIST.

Invertibility matters for meaningful semantic embedding: an LLM's symbols should correspond to something we can identify, not be abstract tokens inside Searle's Chinese translation engine.

Background

Scope: Large Language Models (reinforcement learning with known inputs/outputs and unknown weights — the transformer architecture).

Part I: What Machine Minds Are — Weight Ergodicity

Weights as Ergodic Samples

Measurement theory views weights as **samples from an ergodic distribution** rather than fixed constants. W at any moment is one realization of a stochastic process:

$$W = \mu M + \sigma V$$

where μ and σ are the bias and variance parameters; M has norm 1 and zero variance (analogous to a typical weight matrix), and V is sampled from a standard normal with unit variance.

This perspective comes from **Lagrangian Field-theoretic Gradient Control** (LFGC), treating the weight landscape as a field where boundary conditions shape the distribution from which weights are drawn. The ergodic hypothesis guarantees that time averages (training steps) equal ensemble averages (weight configurations).

See Rogers, A. (2026), *Local Freedom under Global Constraint*, Zenodo, doi: 10.5281/zenodo.18644302.

Temperature and Annealing

A single per-layer *temperature* parameter, updated via ADAM, characterizes mean and variance:

$$\mu = 1 - temp, \quad \sigma = temp$$

Temperature begins at 1 and decreases to 0.

Advantages of Ergodic Weights

1. **No learning rate required.** Biases find different weight-space locations via randomness and converge as temperature decreases.
2. **Zero initialization.** Weights can be initialized to zero — annealing bounces the model out of local minima more effectively than gradient steps.
3. **Theoretical convergence guarantee.** Simulated-annealing literature provides a guarantee of convergence to global optimum (given infinite iterations).

Standard *dropout* can be interpreted as per-neuron variance in the bias parameter. The same regularization occurs naturally from time-varying weight biases.

Part II: What Machine Minds Feel — Network Invertibility

Bidirectional Perception

Egocentric viewpoint sees perception as passive; physics suggests it is bidirectional. Approximated via invertibility of weight multiplications and activation functions.

Symbols, Worldlines, and Karmic Inertia

The symbols an LLM processes have meanings given by **worldlines** — the accumulated history of contexts in which a token has appeared, the gradients that have shaped its embedding, and the network states it has participated in. The meaning of a symbol is the integral of its worldline.

Weight adaptation (or inference) puts a **force** into the system. That force adapts either the system’s inputs or outputs depending on which has less **karmic inertia** — resistance to change. The bidirectional architecture means the force propagates in both directions:

- **Forward (encoding):** input representation adapts to match learned categories — perception shaped by expectation.
- **Reverse (decoding):** internal state adapts to reconstruct the input — expectation grounded in perception.

The invertible architecture makes both directions explicit and trainable.

Implementation

Even when not invertible, bidirectionality plays a role in backpropagation and autoencoders. We examine simultaneously training in both forward (predict output from input) and reverse (predict input from output of the *penultimate* layer) directions — the network is composed of a bijective front (*recognizer*) and a surjective back (*generalizer*). See also Rogers et al. 2001 for related work on sharing weights between policy and critic.

Part III: What Machine Minds Know — Output Certainty

Certainty vs. Probability

A network’s error function corresponds to its desire. Cross-entropy loss *requires* output classification — expressing the probability of one class *versus another*, not the certainty that the input belongs to a given class.

To remedy: blend cross-entropy with a product of cross-entropy and the prediction norm. Zero output is penalized by cross-entropy; incorrect predictions carry additional penalty proportional to magnitude.

Certainty-Weighted Loss

$$\begin{aligned} error &= -\sum_{c=1}^N t_c \log(p_c) \\ certainty &= |\hat{y}| \\ surprise &= (\alpha) \cdot certainty \cdot error + (1 - \alpha) \cdot error \end{aligned}$$

For amplitude certainty:

$$error = (\hat{y} - y)^2$$

Knowing What You Do Not Know

A network should **know that it does not know**, not merely the probability of one answer vs. another. Cross-entropy forces probability distribution across all known classes, even when “I don’t know” is correct. Certainty weighting separates “which class?” from “how sure?”

Methods

Python + PyTorch. MNIST dataset: 60K train, 10K test, 28×28 monochrome at 256 bit depth. Inputs are preprocessed (mean and variance removed) and shuffled.

Input: 28×28 array. Output: one-hot target class. Architecture: 20 hidden unit NN. Batch size 10, 9 epochs, 7 trials for error bars.

Standard benchmark: ReLU, LR 0.01, ADAM, cross-entropy loss.

Base ergodic: zero-init weights, single per-layer temperature, ADAM updates temperature, dropout (0.75) on middle layer.

Five models evaluated: base + four ergodic variants (one plain ergodic, one with input normalization, two reversible — one with separate forward/backward layers, one invertible). The invertible model represents W in SVD form (rotation, eigenvalues, rotation) using Givens rotations and diagonal eigenvalues to avoid direct matrix inversion.


Weight update equations for ergodic models avoid a learning rate. Compared to traditional gradient descent ($\Delta_W = \alpha(y - \hat{y})\nabla_F x$):

$$\begin{aligned} e &= y - \hat{y} \\ c &= \alpha \hat{y}^2 + (1 - \alpha) \\ \Delta_W &= (ce - \alpha \hat{y} e^2)x \end{aligned}$$


The network simultaneously minimizes error and “incorrect certainty”.

Results


Model comparison with error bars shows the simple model performs slightly better than all proposed



paradigms, with all proposed paradigms performing almost identically:



Reconstruction of input (separate weight matrix, MSE loss):



Discussion

Benefits:

- Random initialization unnecessary — only the simple model with random weight variation shows significant performance variation.
- Learning rate replaced by weight noise variance (dropout role replaced by bias variance).
- Error function encapsulates output certainty in addition to between-class certainty.

Limitations:

- Pearson correlation between certainty and accuracy is insignificant — the implementation used $certainty = p_c$ (softmax output) rather than $|\hat{y}|$, creating contradictory constraints since softmax outputs sum to 1 while the certainty norm imposes another constraint. Corrected code should use norm-weighted MSE.
- PyTorch autograd tunes weights directly, interfering with simultaneous temperature tuning in ways not captured.

Conclusion

Strong philosophical grounding for several paradigms, but implementation issues limited testing. Effectiveness in multilayer NNs and challenging tasks remains unclear.

References

Rogers, Shannon, Lendaris; 2001: *A comparison of DHP based antecedent parameter tuning strategies for fuzzy control*, IFSA/NAFIPS, IEEE.

Rogers, A; 2003: *Analysis of the Instantaneous Estimate of Autocorrelation*, unpublished.

Rogers, A; 2026: *Local Freedom under Global Constraint*, Zenodo, doi: 10.5281/zenodo.18644302.

XML Configuration Parameters

Model configuration is specified in XML files (e.g. `data/XOR_exact.xml`). Default values are loaded from `data/model.xml` and overridden by model-specific configs. The schema is in `data/model.xsd`.

Overlay merge order: `data/model.xml` $\rightarrow$ `<config>.xml` $\rightarrow$ runtime overrides.

Note on naming: the schema and the XML configs use `nInput` / `nOutput` (sequence-length sentinels), `nDim` (per-vector dim), `nVectors` (codebook size), and `nInputDim` / `nOutputDim` (raw shape hints for the `LiftingLayer`). Older docs sometimes call `nOutput` / `nInput` `nActive`; treat that as the same thing as `nOutput`.

2026-05-28: `<nWhere>` / `<nWhen>` are retired from per-file configs. `data/model.xml` carries per-Space defaults (`<nWhere>2</nWhere>` on InputSpace / PartSpace / WholeSpace, 0 elsewhere). The `.where` field is the canonical positional identifier; see `doc/old/2026-05-28-where-keyed-taxonomy.md`.

2026-06-16: `.where` and `.when` are now **endpoint-sum brackets** `[start, end]` (angle = span center, magnitude = extent), the invertible `EndpointSumWhere` form. An instant (`start==end`) is byte-identical to the prior single-quadrature point, so `whereScale` / `whenScale` are unchanged. `.when` tense is the interval-vs-now relation (the magnitude-D tense scheme is retired); see `doc/Spaces.md`.

<architecture>

Core model settings. Training and data parameters are in nested sub-elements `<training>` and `<data>` (see below).

Parameter	Type	Default	Description
<code>data/dDataType</code>	string	"numeric"	The data space-role (was the retired architecture-level <code>modelType</code>), s
<code>reconstruct</code>	—	—	RETIRED (A1, 2026-06-09): the <code>reconstructEnum</code> / <code><reconstruct</code>
<code>subsSymbolicOrder</code>	int	1	Iteration depth of the P→C→S pipeline. <code>order > 1</code> partitions the s
<code>processSymbols</code>	bool	false	Apply extra symbolic processing after Sigma.
<code>monotonic</code>	bool	false	Constrain invertible Sigma / Pi to $W \geq 0$ so the lift / lower chain is
<code>ergodic</code>	bool	false	Ergodic exploration: every Layer uses <code>W_eff = bias * W + temp *</code>
<code>naive</code>	bool	false	Materialise <code>W_eff</code> densely in <code>InvertibleLinearLayer</code> . Slower; debu
<code>serial</code>	bool	derived	Forward-dispatch mode. <code>true</code> = serial / grammatical (per-word [B,
<code>symbolicOrder</code>	int	0	Symbolic / relational loop budget: the RAMSIFIED DEPTH of the
<code>subsSymbolicStack</code>	bool	false	P4 sigma/pi stacks: DISTINCT per-pass subsymbolic layers — PS s
<code>subsSymbolicNoop</code>	string	(empty)	Pass indices whose stack layer is the IDENTITY pass-through, e.g.
<code>sparseReplace</code>	—	—	RETIRED (2026-07-02 P3 two-phase forward): phase separation ma
<code>routerWireSerial</code>	string	"both"	Per-word router-fire gating on the serial path: per-word (fire per wo
<code>learning</code>	bool	false	Two-pass soft-superposition training for the grammar chooser. Whe
<code>exploreTemperature</code>	decimal	0.5	Superposition temperature $t \in [0, 1]$ for pass B of learning . 0 = the
<code>transformChooser</code>	string	"anchordot"	Placement scorer for the structured grammar layers. <code>anchordot</code> = s
<code>reconstructFromIdea</code>	bool	false	Reverse from the idea rather than replaying the forward parse. Whe
<code>categoryCodebook</code>	bool	true	MetaSymbol participation-category codebook for the role-collapsed g
<code>adverbEigEdit</code>	bool	false	Legacy/direct <code>LiftLayer</code> helper flag for the adverb sparse eigenvalu
<code>mereologyRaise</code>	bool	false	Mereological order-raising: perception's autobind hook builds a mere
<code>embeddingPath</code>	string	(empty)	gensim <code>KeyedVectors</code> path. Empty disables embedding load.
<code>weightsPath</code>	string	(empty)	Model weights checkpoint path. Empty falls back to <code>output/<name></code>
<code>maxResponseLength</code>	int	4096	Inference token budget (characters / bytes / tokens). Caps output a
<code>amp</code>	string	"off"	AMP autocast mode: <code>off</code> , <code>fp16</code> , <code>bf16</code> .
<code>truthBiasScale</code>	decimal	0.1	Strength of luminosity-based bias on concept input during forward.
<code>LuminosityWeight</code>	decimal	0.1	Multiplicative loss penalty for incoherent truths.
<code>UniversalityWeight</code>	decimal	0.1	Multiplicative loss penalty for unkind propositions.
<code>allowExcludedMiddle</code>	int	1	Catuskoti policy: 1 permits NEITHER, -1 enforces classical LEM.
<code>allowContradiction</code>	int	0	Catuskoti policy: 0 forbids BOTH (non-contradiction), 1 permits pa
<code>gateL1Lambda</code>	float	0.0	L1 sparsity penalty on lift / lower gates. 0 disables.
<code>loadBalanceWeight</code>	float	0.0	Sparse-MoE load-balance loss weight. Currently inert / no-op : the
<code>l1Lambda</code>	decimal	0.0	Architecture-wide L1 penalty hook. Most configs leave this at 0 and
<code>discontinuityLambda</code>	decimal	0.0	Architecture-wide discontinuity penalty hook (legacy).
<code>conceptualWidth</code>	string	"tapered"	Symbol-partition geometry for <code>subsSymbolicOrder > 1</code> . <code>tapered</code> =
<code>maxActivePerLayer</code>	int	8	Maximum active nodes per layer for reverse / mereological walks.
<code>codebookRetire</code>	bool	false	Retire codebook path where supported. Mostly a migration hook.

Parameter	Type	Default	Description
<code>symbolLearning.enabled</code>	bool	false	Enables runtime symbol-learning hooks. Disabled by default.
<code>mereologicalTreeSize</code>	int	–	Retired compatibility knob. Values are accepted but ignored; the co
<code>writeSyntax</code>	bool	false	Dump POS-labelled syntax tree at end of every <code>MentalModel.forward</code>
<code>syntaxOutPath</code>	string	–	Output path for <code>writeSyntax</code> .

Luminosity and universality are **multiplicative** — they scale the total loss by

$$(1 + \textit{LuminosityWeight} \cdot (1 - \text{luminosity}) + \textit{UniversalityWeight} \cdot (1 - \text{universality})).$$

`TruthLoss` (in `<training>`, below) is **additive** — it directly penalises propositions that contradict the `TruthSet`. Both coexist. See `Ethics.md` Section Universality and `Reasoning.md` Section `TruthLoss` for details.

`<architecture><data>`

Data loading and filtering.

Parameter	Type	Default	Description
<code>dataset</code>	string	"xor"	Dataset key: <code>xor</code> , <code>mnist</code> , <code>tomatoes</code> , <code>text</code> , <code>inline</code> . Omit for inference-only mode.
<code>numShards</code>	int	1	Shard count (streaming text datasets).
<code>maxDocs</code>	int	10000	Per-shard document cap.
<code>shardDir</code>	string	(empty)	Text shard directory (e.g. <code>data/fineweb</code>). Only used when <code>dataset=text</code> .
<code>minFrequency</code>	float	0.0	Vocabulary admission threshold (fraction of corpus). Words below it are held in a
<code>classificationMin</code>	float	–	Lower bound for classification-accuracy reporting.
<code>classificationMax</code>	float	–	Upper bound for classification-accuracy reporting.

Inline-mode configs additionally accept `<input use="train|test|validation">...|...|...</input>` and matching `<output>` elements directly inside `<data>`.

`<architecture><training>`

Training loop and I/O.

Parameter	Type	Default	Description
<code>numTrials</code>	int	1	Independent training runs per invocation.
<code>numEpochs</code>	int	3	Epochs per trial.
<code>batchSize</code>	int	64	Contiguous streams through the dataset. Each batch row <code>b</code> receives the next
<code>numWorkers</code>	int	0	<code> DataLoader </code> prefetch workers. 0 = synchronous in-process batch assembly.
<code>learningRate</code>	float	0.001	Adam learning rate.
<code>reconstructionScale</code>	float	0.5	Weight of reconstruction loss vs prediction loss in $[0, 1]$: $\mathcal{L}_{\text{total}} = (1 - r) \mathcal{L}_{\text{ou}}$
<code>whatScale</code>	float	0.7	Loss weight on the <code>.what</code> (content) channel.
<code>whereScale</code>	float	0.2	Loss weight on the <code>.where</code> (positional) channel.
<code>whenScale</code>	float	0.1	Loss weight on the <code>.when</code> (temporal) channel.
<code>negSamples</code>	int	64	Negative samples per positive for CBOW / SBOW. Memory cost: $O(\text{batch} \cdot$
<code>TruthLoss</code>	float	0.0	Additive penalty for propositions that contradict the <code>TruthSet</code> (union-norm
<code>maskRate</code>	float	0.15	Within-sentence IR Bernoulli mask rate. BERT default 0.15.
<code>trainEmbedding</code>	string	"NONE"	Embedding update mode: NONE, CBOW, SBOW, BACKPROP, BOTH, JOINT. See ta

Parameter	Type	Default	Description
embeddingScale	float	0.25	JOINT-mode embedding loss weight λ : $\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{model}} + \lambda \cdot \mathcal{L}_{\text{emb}}$. Ignored
autoload	bool	true	Load weights from <code>weightsPath</code> at construction. Stale incompatible checkp
autosave	bool	false	Save weights after training completes.
checkpointEveryBatches	int	0	Save mid-training checkpoint every N optimizer steps. 0 disables periodic ch
profile	bool	false	cProfile the training loop.
certainty	bool	false	Per-neuron certainty tracking in ergodic layers. Allows individual neurons to
sentencePrediction	bool	false	Enables <code>InterSentenceLayer</code> , an ARMA(p, q) predictor over sentence repr
armaScale	float	0.1	Loss weight for the ARMA sentence-prediction MSE.
sentencePrimingScale	float	0.05	AR prediction cast into <code>concept_dim</code> and added as a bias to <code>concept_inpu</code>
intraLossWeight	float	0.1	Loss weight on the in-STM next-idea term $\mathcal{L}_{\text{intra}} = \text{MSE}(\hat{c}_t, c_t)$ from <code>IntraS</code>
interLossWeight	float	0.1	Loss weight on the inter-sentence next-end-state term $\mathcal{L}_{\text{inter}} = \text{MSE}(\hat{p}, p)$ fro

<trainEmbedding> — Embedding Update Modes

Value	Embedding method	Model layers	Gradients through codebook
NONE	Frozen	Trained	No
CBOW	CBOW (padded context)	Trained	No
SBOW	SBOW (leave-one-out centroid)	Trained	No
BACKPROP	Backprop only	Trained	Yes
BOTH	SBOW post-batch	Trained	Yes (two optimizers)
JOINT	Single backward	Trained	Yes ($\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{model}} + \lambda \cdot \mathcal{L}_{\text{SBOW}}$)

<architecture><server>

HTTP server settings (used by `serve.py`).

Parameter	Type	Default	Description
host	string	"127.0.0.1"	Bind address.
port	int	8001	Port. Override per model to run several in parallel.
timeout	int	120	Per-request timeout (s).

BasicModel / MentalModel Quick Reference

The two configurations that train the full $P \rightarrow C \rightarrow S$ pipeline:

```

<architecture>
  <truthBiasScale>0.1</truthBiasScale>
  <LuminosityWeight>0.1</LuminosityWeight>
  <UniversalityWeight>0.1</UniversalityWeight>
  <subsymbolicOrder>1</subsymbolicOrder>
  <monotonic>false</monotonic>
</architecture>

<training>
  <TruthLoss>0.0</TruthLoss>
</training>

```

The grammar section of `MentalModel.xml` declares `lift` as an SS binary rule:

```
<S>lift(S, S)</S>      <!-- lift is binary; SVO is assembled by
                        the chart-compose path S -> S VO,
                        VO -> V O, not a ternary lift -->
```

Transitive SVO is produced compositionally via typed grammar rules such as `S -> S VO / VO -> V O`, or via STM `ParseState` extraction for a subject lift over a transitive verb-phrase derivation. Grammar mode is derived from the loaded grammar.

Space sections

Every space block uses the same sentinel convention:

- `nInput` = 0 — `InputSpace`: derive from `TheData`; other spaces: derive from the previous space's `outputShape[0]`.
- `nOutput` = 0 — same as `nInput` for this space (passthrough count).
- `nVectors` = 0 — same as `nOutput` (used for non-codebook spaces).

Per-Space `<nWhere>` / `<nWhen>` declarations were retired on 2026-05-28 (see top-of-file note); the per-Space defaults live in `data/model.xml` and don't need re-declaration in per-experiment configs.

<InputSpace>

Lifts raw data into the model's internal representation.

Parameter	Type	Default	Description
<code>nOutput</code>	int	sentinel	Sequence length: max tokens per input. For XOR: 2. For text models, OOV
<code>nDim</code>	int	1	Per-vector content dim.
<code>nInputDim</code> / <code>nOutputDim</code>	int	0	Raw shape hints. 0 defers to the loader's native dim.
<code>nVectors</code>	int	sentinel	Codebook size. Defaults to <code>nOutput</code> .
<code>codebook</code>	mode	none	none, quantize, or project. Legacy true/false are accepted as quantize/

<lexer> moved to <WholeSpace> (Phase 4b of the analysis/synthesis dual-input plan, rev. 2026-06-09): lexing is analytic CUTTING, owned by the analysis side. An `InputSpace`-side `<lexer>` now fails schema validation and the runtime reader. `InputSpace` emits the dual view — `forward(x) -> (percepts_in, concepts_in)`: the atom view (content `[B, N, 1]`) for the perceptual branch and the unity view (`[B, 1, N]`) for the symbolic branch.

`inputShape` uses the data's native dim (e.g. 784 for MNIST, `inputLength` for text); `outputShape` uses `nDim` from `TheObjectEncoding`. `LiftingLayer` bridges the two.

<PartSpace>

Transforms lifted input into perceptual features via its synthesis fold — a `SigmaLayer` (additive/union; the Pi/Sigma swap of the analysis/synthesis plan, rev. 2026-06-09: PS is bottom-up SYNTHESIS over atoms; the multiplicative Pi fold moved to `WholeSpace` as top-down analysis).

Parameter	Type	Default	Description
<code>nOutput</code>	int	sentinel	Number of active perceptual feature vectors. Should be $\geq$ <code>InputSpace.nOutput</code> . For
<code>nDim</code>	int	1	Per-vector content dim.
<code>nVectors</code>	int	sentinel	Codebook size. When $>$ <code>nOutput</code> , enables vector quantization with top-k selection o
<code>invertible</code>	bool	false	true: one <code>PiLayer(invertible=True)</code> handles both directions; false: separate pi1
<code>hasAttention</code>	bool	true	DEPRECATED and INERT: the legacy boolean no longer constructs a QKV attent

Parameter	Type	Default	Description
<code>nonlinear</code>	bool	<code>true</code>	Tanh-bound output to $[-1, 1]$.
<code>synthesis</code>	string	<code>"lexicon"</code>	Bottom-up union strategy (renamed from <code><chunking></code> , Phase 4a — the legacy spelling).
<code>butterfly</code>	bool	<code>false</code>	FFT-style element-pair butterfly cascade on the PS fold (<code>self.sigma</code> post-swap) for
<code>wordLearning</code>	int	0	Active lexicon-growth mode. 0 = frozen codebook (training default). ≥ 1 = on first

Sigma layer math: $y_j = \tanh(Wx + b)$ (the additive/union fold). The PS `<codebook>` element was retired (the percept prototypes live on the `.what` basis; radix's PerceptStore IS the surface codebook). See Architecture.md and Philosophy.md for the analysis/synthesis orientation.

<ConceptualSpace>

Transforms perceptual features into abstract concepts via Sigma layers (additive / linear).

Parameter	Type	Default	Description
<code>nOutput</code>	int	sentinel	Active concept vectors. For XOR: 3.
<code>nDim</code>	int	1	Per-vector content dim.
<code>nVectors</code>	int	sentinel	Codebook size.
<code>invertible</code>	bool	<code>false</code>	<code>true</code> : <code>SigmaLayer(invertible=True)</code> (exact inversion via $\text{atanh} + W^{-1}$); <code>false</code> : separable
<code>hasAttention</code>	bool	<code>false</code>	DEPRECATED and INERT: no longer constructs an attention pass in conceptual processing
<code>nonlinear</code>	bool	<code>true</code>	Tanh-bound output to $[-1, 1]$.
<code>codebook</code>	mode	<code>none</code>	<code>none</code> , <code>quantize</code> , or <code>project</code> .
<code>butterfly</code>	bool	<code>false</code>	Butterfly cascade on CS's PiLayer (same machinery as PartSpace). Required by <code>MM_xor</code>
<code>stmCapacity</code>	int	8	STM ring depth (within Miller's 7 ± 2 band). Per-batch buffer [<code>B</code> , <code>stmCapacity</code> , <code>nDim</code>]

Sigma layer math: $y_j = \tanh(Wx + b)$. See Architecture.md. `InvertibleSigmaLayer` has been merged into `SigmaLayer(invertible=True)`.

2026-05-29 (clean-stack STM): `ConceptualSpace.forward` no longer applies `sigma_in / sigma_cs` on the forward path. The Stage-10 additive composition `sigma_in(combined) + sigma_cs(prev)` is replaced with per-stage space-role attribution (`folded = primary` at stage 0, `folded = sym` at $k > 0$). The sigma layers remain allocated and are still invoked on the reverse path, but their parameters get no gradient on forward — they're dead weight under the current experiment. See doc/old/2026-05-29-clean-stack-stm-basis-arg-radixlayer.md.

<WholeSpace>

Discrete symbolic representation — the information bottleneck between perception and output, and (post the analysis/synthesis plan, rev. 2026-06-09) the home of top-down ANALYSIS: SS owns the Pi fold (multiplicative/intersection), the `<analysis>` division knob, and the `<lexer>` intake knob. See Language.md and Philosophy.md for the design.

Parameter	Type	Default	Description
<code>nOutput</code>	int	sentinel	Total symbols. When reconstruction symbols are enabled, <code>nOutput</code>
<code>nDim</code>	int	1	Per-vector content dim.
<code>nVectors</code>	int	sentinel	Codebook size. When <code>codebook=quantize</code> , equals the number of
<code>codebook</code>	mode	<code>none</code>	<code>none</code> , <code>quantize</code> , or <code>project</code> . Required for the full symbolic pipeli
<code>analysis</code>	string	<code>"byte"</code>	Top-down division strategy over the unity view (Phase 4b): <code>byte</code>
<code>lexer</code>	string	—	Intake granularity InputSpace executes (<code>word</code> , <code>sentence</code> , <code>byte</code> , <code>ra</code>)
<code>nonlinear</code>	bool	<code>true</code>	Tanh-bound the activation.
<code>sortNetwork</code>	bool	<code>false</code>	Bitonic-sort regularizer on the codebook ordering.

Parameter	Type	Default	Description
<code>sortPasses</code>	int	0	Sort-network depth. 0 disables.
<code>l1Lambda</code>	decimal	0.01	L1 sparsity penalty on the symbolic activation.
<code>symbolResidualScale</code>	decimal	1.0	Symbol-residual injection scale (legacy mereology hook).
<code>outputSymbolResidualScale</code>	decimal	0.0	Output-side symbol-residual injection scale.
<code>commitmentBeta</code>	float	0.25	Space-level VQ-VAE commitment-loss β . Asymmetric dispatch
<code>semanticArrangement</code>	float	0	Task 5 (C-13): post-sentence semantic arrangement over SS heat -
<code>decorrelationWeight</code>	decimal	0.0	ImpenetrableLayer decorrelation regularizer on codebook rows.
<code>spectralFlatnessWeight</code>	decimal	0.0	ImpenetrableLayer spectral-flatness regularizer.
<code>truthCriterion</code>	unitInterval	1.0	Single continuous truth bar governing BOTH (a) WholeSpace
<code>trust</code>	unitInterval	1.0	Model-level multiplier for incoming assertions/testimony. Runtime

WholeSpace owns one square invertible `PiLayer` at `self.pi` — the top-down analysis (product/intersection) fold — bridging the $C \leftrightarrow S$ boundary in both directions via its exact inverse (the `Pi/Sigma` swap, rev. 2026-06-09; SS owned a `SigmaLayer` at `self.sigma` before the corrected orientation). Consumers use `model.wholeSpace.pi` directly; the grammar rule binding registers it under both the `pi` rule name and the legacy `sigma` alias.

<OutputSpace>

Maps symbols to final predictions via linear layers.

Parameter	Type	Default	Description
<code>nOutput</code>	int	sentinel	Output values. For XOR: 1.
<code>nDim</code>	int	1	Per-output dim.
<code>nVectors</code>	int	sentinel	Codebook size. Defaults to <code>nOutput</code> .
<code>invertible</code>	bool	false	Rarely set on <code>OutputSpace</code> .
<code>codebook</code>	mode	none	none , quantize , or project .
<code>nonlinear</code>	bool	false	<code>OutputSpace</code> is linear by default; rescales via <code>Data.denormalize</code> .
<i>(readout — retired 2026-06-19)</i>	—	—	The <code>OutputSpace</code> regression head (backlog #11): with <code>nVectors == 1</code>

`LinearLayer` with (`bias`, `temp`) support for ergodic mode.

<SymbolSpace>

Grammar infrastructure (`SyntacticLayers`, `TruthLayer`). Lives outside <architecture> as a top-level sibling.

Parameter	Type	Default	Description
<code>syntacticHiddenDim</code>	int	64	<code>SyntacticLayer</code> MLP hidden size.
<code>truthMaxEntries</code>	int	1024	<code>TruthLayer</code> SS-codebook capacity.
<code>ltmCapacity</code>	int	1024	Long-term-memory (LTM) chain capacity on InterSentenceLayer : the pe
<code>parserBackend</code>	—	—	RETIRED (Stage 3 cleanup, 2026-05-27): the chart backend is gone and th
<code>routerKind</code>	—	—	RETIRED (Stage 3 cleanup, 2026-05-27): there is no longer a selectable ro
<code>chartCollapse</code>	string	"root"	Soft-superposition output collapse mode, preserved for the signal router's r
<code>chartTau</code>	—	—	RETIRED (Stage 3 cleanup, 2026-05-27) alongside the chart. Rejected wit
<code>wMax</code>	—	—	RETIRED: the legacy STM-capacity alias is no longer read. STM depth co
<code>chartTopK</code>	—	—	RETIRED (Stage 3 cleanup, 2026-05-27) alongside the chart. Rejected wit
<code>chartNoiseEps</code>	—	—	RETIRED (Stage 3 cleanup, 2026-05-27) alongside the chart. Rejected wit
<code>iterationsPerWord</code>	int	1	P->C iterations per word in the per-word stem path.

Parameter	Type	Default	Description
downwardGeneration	bool	false	Emit a codebook head via <code>SymbolSpace.reconstruct()</code> and stash on <code>self</code> .
language.interpretation	float	0.5	Soft-interpolation weight between forward and generation directions.

The `<language><grammar>` block contains space-role-scoped rules in `<compose>` and `<generate>` sections, each holding `<percepts>` / `<concepts>` / `<symbols>` sub-blocks of `<rule>` elements. Per the 2026-05-07 rollback, the grammar XML is the sole source of truth for which rule fires at each space-role (no code-level `default_rule` fallback). Legacy bare-space-role elements like `<P>P = sigma(P)</P>` directly under `<grammar>` are still accepted.

InvertibleLinearLayer (programmatic)

The LDU-factorised invertible linear primitive used internally by `SigmaLayer(invertible=True)` and `PiLayer(invertible=True)`. Not set via XML directly.

Parameter	Type	Default	Description
naive	bool	false	false : apply L, D, U sequentially via triangular solves (no W materialisation). true : materialise W .
stable	bool	false	Clamps each diagonal entry d_i to magnitude $[\epsilon, 1]$ with sign preserved before the ergodic blend.
ergodic	bool	false	Factor-level noise injection. Registers <code>noise_raw_L</code> , <code>noise_raw_U</code> , <code>noise_d</code> buffers; resamples on <code>forward</code> .

Class renames and removals.

Old name	New name / status
LULayer	Renamed to <code>InvertibleLinearLayer</code>
InvertibleLinearLayer (SVD-based)	Removed; replaced by LDU factorisation
InvertibleSigmaLayer	Merged into <code>SigmaLayer(invertible=True)</code>
InvertiblePiLayer	Merged into <code>PiLayer(invertible=True)</code>

See `Architecture.md` for the LDU factorisation and factor-level noise injection.

Example: XOR Configuration

```

<model>
  <architecture>
    <data><dataType>embedding</dataType></data>

    <training>
      <numEpochs>400</numEpochs>
      <learningRate>0.01</learningRate>
      <autoload>false</autoload>
      <trainEmbedding>JOINT</trainEmbedding>
      <embeddingScale>0.05</embeddingScale>
    </training>
  </architecture>

  <InputSpace>
    <nVectors>8</nVectors>

```

```

    <nDim>10</nDim>
    <nOutput>8</nOutput>
    <nOutputDim>10</nOutputDim>
  </InputSpace>

  <PartSpace>
    <nInput>8</nInput>
    <nVectors>1000</nVectors>
    <nDim>10</nDim>
    <nOutput>8</nOutput>
    <hasAttention>false</hasAttention>
    <invertible>true</invertible>
    <codebook>quantize</codebook>
  </PartSpace>

  <ConceptualSpace>
    <nInput>8</nInput>
    <nVectors>8</nVectors>
    <nDim>10</nDim>
    <nOutput>8</nOutput>
    <invertible>true</invertible>
    <codebook>quantize</codebook>
  </ConceptualSpace>

  <WholeSpace>
    <nInput>8</nInput>
    <nVectors>8</nVectors>
    <nDim>2</nDim>
    <nOutput>8</nOutput>
    <codebook>quantize</codebook>
  </WholeSpace>

  <OutputSpace>
    <nOutput>1</nOutput>
    <nVectors>1</nVectors>
  </OutputSpace>
</model>

```

Everything not listed inherits from `data/model.xml`.

Example: Embedding / Chatbot Configuration

```

<model>
  <architecture>
    <data><dataType>embedding</dataType></data>
    <weightsPath>BasicModel.ckpt</weightsPath>
    <embeddingPath>BasicModel.kv</embeddingPath>

    <data>
      <shardDir>data/fineweb</shardDir>
      <minFrequency>0.00001</minFrequency>
    </data>
  </architecture>
</model>

```



```

<training>
  <numEpochs>2</numEpochs>
  <batchSize>128</batchSize>
  <reconstructionScale>0.1</reconstructionScale>
  <trainEmbedding>BACKPROP</trainEmbedding>
  <autosave>true</autosave>
  <checkpointEveryBatches>500</checkpointEveryBatches>
  <sentencePrediction>true</sentencePrediction>
  <armaScale>0.1</armaScale>
</training>
</architecture>

<!-- space definitions inherit from model.xml -->
</model>

```

Key points:

- `<data><dataType>embedding</dataType>` activates the embedding-based input pipeline alongside the neural model.
- `trainEmbedding=BACKPROP` trains model layers with gradients flowing through the codebook.
- XML config defines the architecture. `weightsPath` stores the neural model checkpoint.
- `minFrequency` gates vocabulary admission; words below it are buffered until they exceed the threshold.